

H3

The Relational Languages

Fundamentals of Databases

Alvaro A A Fernandes, SCS, UoM

[COMP23111 Handout 03 of 12]

Acknowledgements

- These slides are adaptations (mostly minor, but some major) of material authored and made available to instructors by **A. Silberschatz, H. F. Korth, and S. Sudarshan** to accompany their textbook **Database System Concepts, 6th Edition, McGraw-Hill, 2006, 978-0-07-352332-3** and **H. Garcia-Molina, J.D. Ullman, and J. Widom** to accompany their textbook **Database Systems: The Complete Book, 2nd Edition, Pearson Prentice Hall, 2009, 978-0-13-135428-9**
- Copyright remains with them, whom I thank.
- All errors are my responsibility.

In Previous Handouts

- We learned that data is an enterprise asset and that DBMSs are crucial to manage it well.
- We learned the importance of adopting different levels of abstraction in designing and implementing databases.
- We learned how data models lead to schemas and instances that enable a logical view of the data.
- We learned about the main components of DBMSs and the various architectures used to deploy them.
- We started learning about the relational approach to data modelling.
- We started learning SQL, focussing on its DDL and DML capabilities.

In This Handout

- What is an algebra?
- What is a relational algebra?
- What are the basic operations of relational algebra?
- What are the most important derived operations?
- What are the most important operations to have in an extended relational algebra?
- How do we relate SQL to relational algebra?
- What are the main constructs and query types in SQL?

What is an algebra?

What is an Algebra?

- An algebra is a mathematical structure consisting of:
 - ▶ Operands: variables or values, typically drawn from a single sort (i.e., one type only), from which new values can be constructed
 - ▶ Operators: symbols denoting procedures that construct new values from given values and with closure
- Equivalences between expressions follow from operator semantics (e.g., commutativity, associativity)
- By closure we mean that evaluating any well-formed expression yields an element of the same type as the inputs, i.e., from the set of operands

A Familiar Example

7

- Take as operands, the real numbers and an infinite set of letters used to denote them.

$$2+2$$

- Take as operators: addition (+), subtraction (-), multiplication (*), division (/), all binary.

$$4*(3/2)$$

$$(4*3)/2$$

$$x+1$$

- We can construct some example expressions.

$$2/(4*x)$$

- Addition and multiplication are commutative and associative.

$$x+y = y+x$$

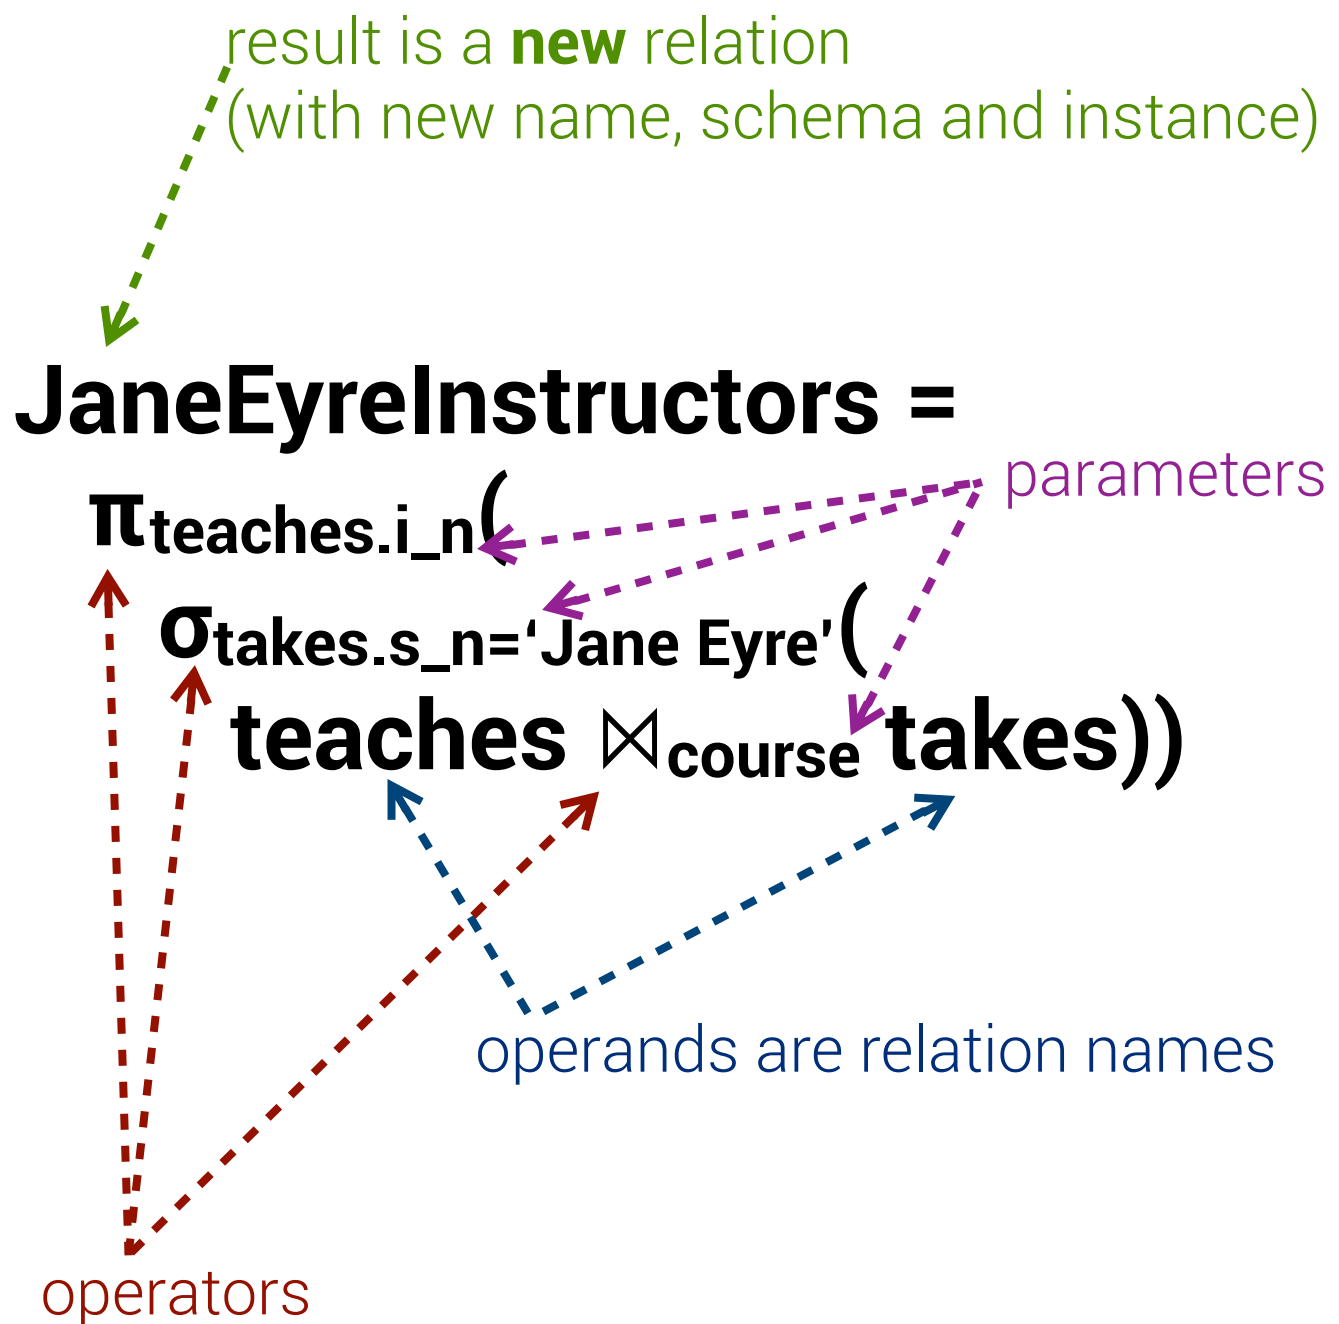
$$(x*y)*z = x*(y*z)$$

- Closure holds.

Relational Algebras

8


- A relational algebra (RA) is an algebra whose operands are relation names (i.e., variables that denote relation instances).
- Relational-algebraic operators are designed to do the most common things that we need to do with relations in a database.
- They can take parameters (e.g., a predicate, or a list of attribute names).
- The result is an algebra that can be used as a query language over relational databases.



Relational Algebras

9

- When evaluated against a database instance, a relational-algebraic expression yields a relation that is the result of the query.
- The evaluation is grounded on the relation instances denoted by variables in the expression.
- One special characteristic is that, because attributes and relations have names (i.e., there is the notion of a relation schema), we need to:
 - ▶ Take into account the schemas of the input relations, and
 - ▶ Infer (and perhaps explicitly act to specify) the schema of the output relation.



Some Terminology

10

- A relation R with schema $R = \{a_1:t_1, \dots, a_n:t_n\}$ has arity (or degree) n
- An instance of R is a set of tuples conforming to the schema of R
- By conformance to a schema we mean that
 - ▶ every tuple must have the schema-defined arity (i.e., number of columns) and
 - ▶ the value of each schema-defined attribute a_i in every tuple must have the schema-defined type t_i
- The cardinality (of a given instance of R) is the number of tuples in the instance, which we denote by $|R|$

Core Relational Algebra

Core Relational Algebra

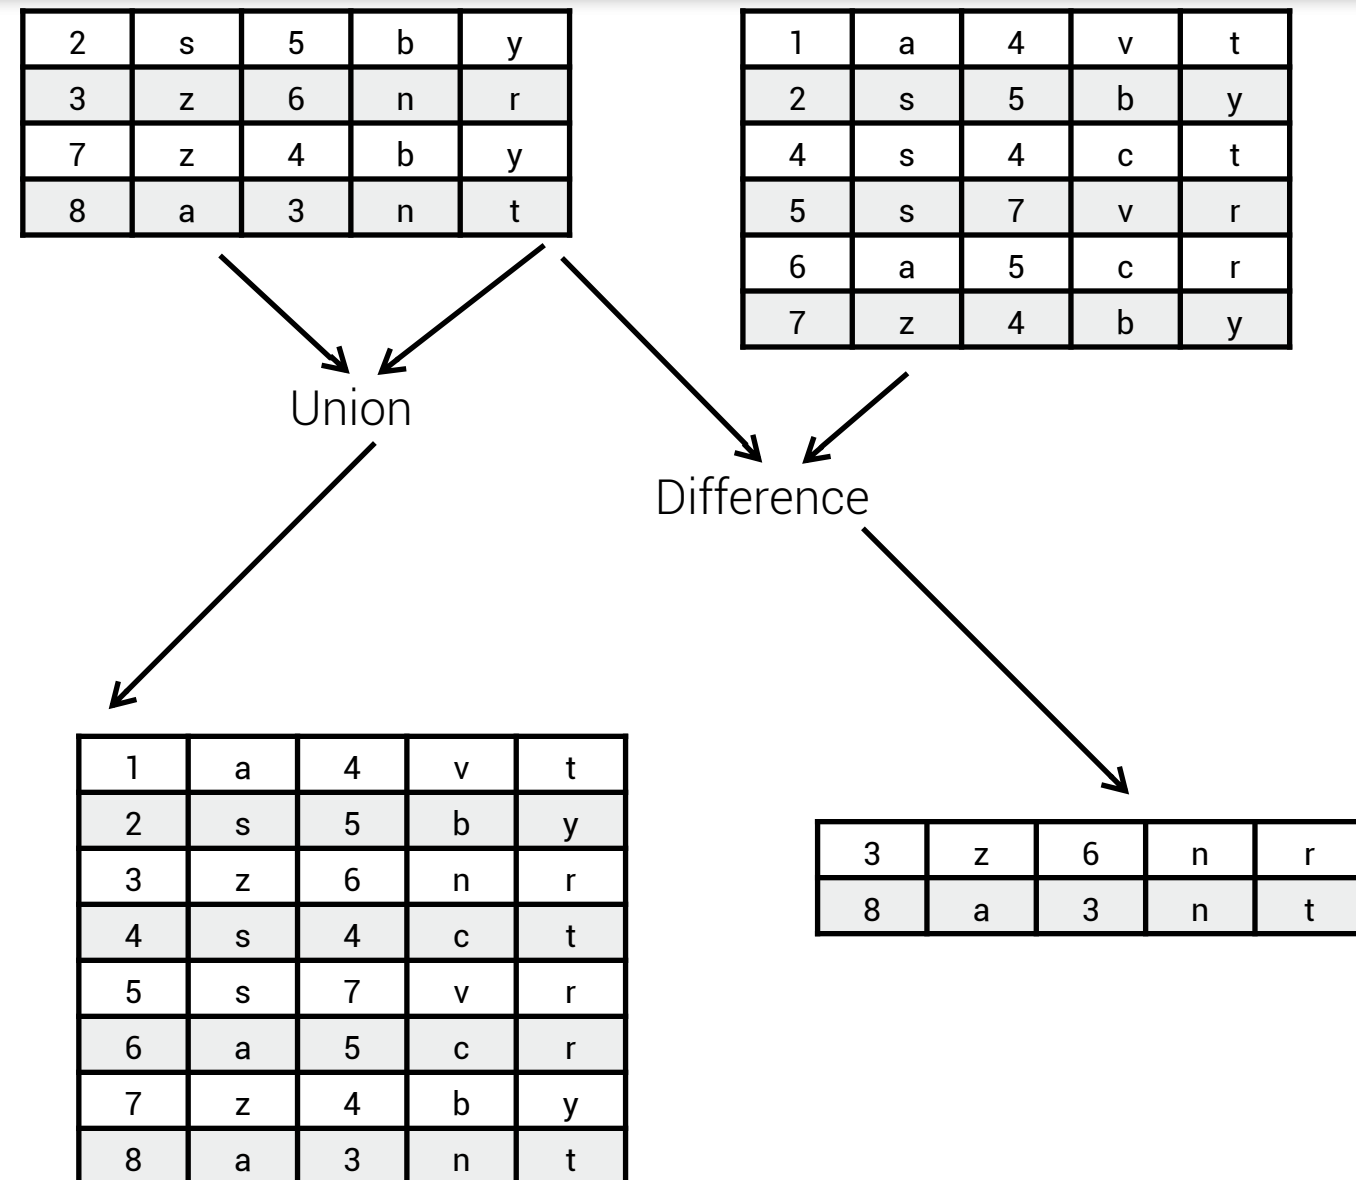
12

- Some relational-algebraic operators are primitive and some are derivable.
- Consider, by analogy, addition and multiplication.
- One can derive multiplication from addition, whereas addition is not derivable from the usual arithmetic operators.
- Addition is, in some sense, primitive, whereas multiplication is not.
- We define a core RA by taking a set of primitive operators only
- Then, we specify the derivation of additional ones from our choice of primitives.

Core Relational Algebra: Union, (Intersection) and Difference

13

- Because relations are sets of tuples, the binary set-theoretic operators **Union** and **Difference** are part of the core RA.
- **Intersection** is derivable from union and difference: $A \cap B = ((A \cup B) \setminus (A \setminus B)) \setminus (B \setminus A)$.
- In the case of RA, for set-theoretic operations to be well-formed, both operands must have compatible relation schemas.
- The schemas of two relations R and R' are *compatible* if they have the same arity and for every attribute of type t_i in R there is a corresponding attribute of type t_i in R' .

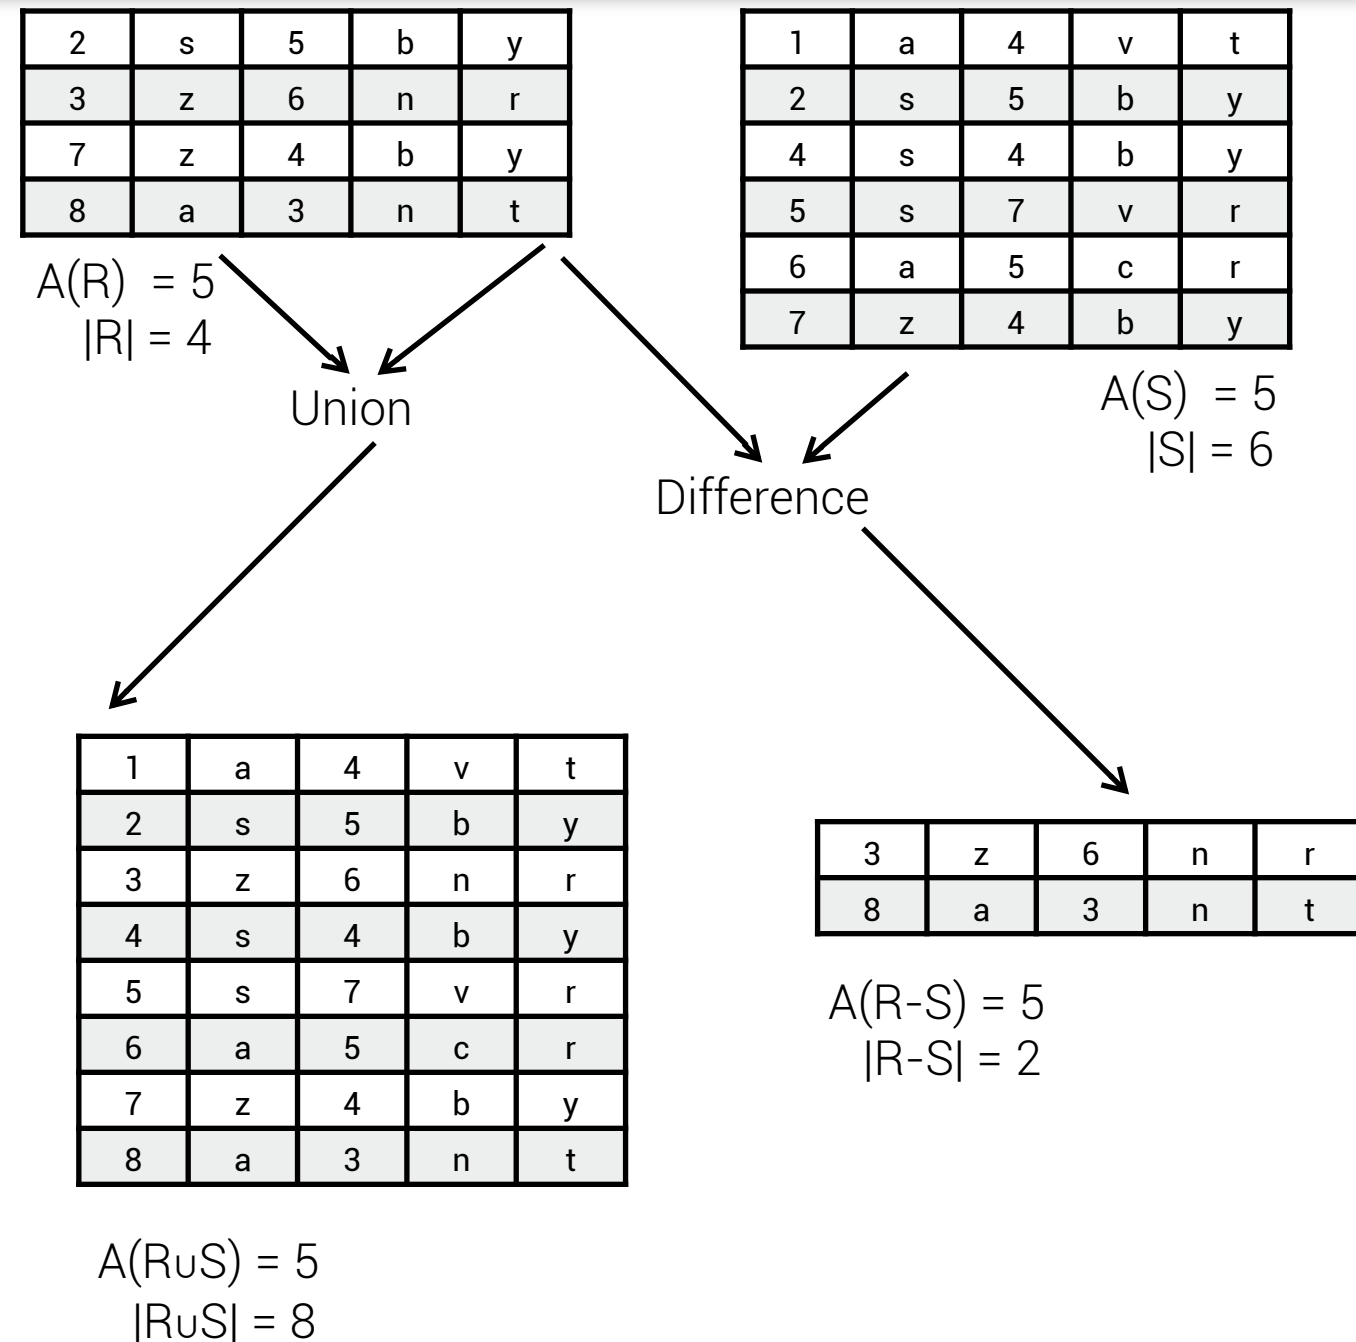


What about Intersection?

Core Relational Algebra: Union, (Intersection) and Difference

14

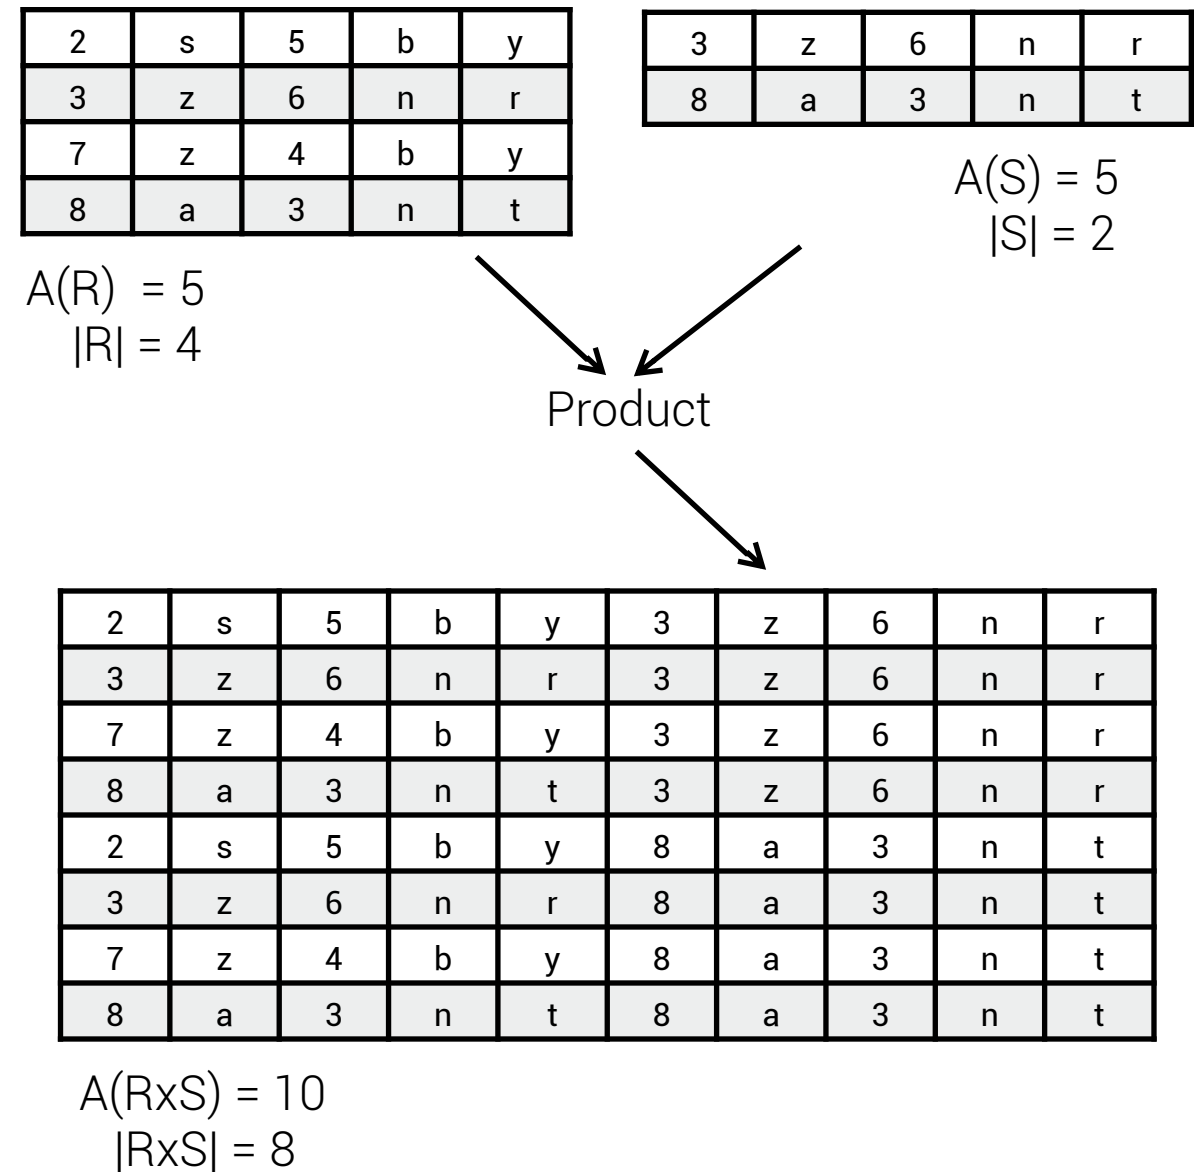
- Notice that these operations have the potential to yield a result with different cardinality than either operand, but the arity of the result does not differ from that of the operands.
- The cardinality may be
 - ▶ larger (in the case of **Union**)
 - ▶ smaller (in the case of **Intersection** and **Difference**)



Core Relational Algebra: (Cartesian) Product

15

- The binary set-theoretic (**Cartesian**) **product** operator is also in the core RA.
- It concatenates every tuple of one operand with every tuple of the other to produce tuples in the result.
- It creates compositions/associations of relations.
- It normally yields a result with larger cardinality and larger arity than either operand.



Core Relational Algebra: Selection

16

- **Selection** is a unary operator that picks only certain rows from its operand.
- It has the potential to yield a result with smaller cardinality than either operand, but the arity of the result does not differ.

A	B	C	D	E
1	a	4	v	t
2	s	5	b	y
3	z	6	n	r
4	s	4	b	y
5	s	7	v	r
6	a	5	c	r
7	z	4	b	y
8	a	3	n	t

$$A(R) = 5$$
$$|R| = 8$$

$S = \text{Selection}[B \neq 's'](R)$

A	B	C	D	E
1	a	4	v	t
3	z	6	n	r
6	a	5	c	r
7	z	4	b	y
8	a	3	n	t

$$A(S) = 5$$
$$|S| = 5$$

Core Relational Algebra: Projection

17

- **Projection** is an operator that picks only certain columns from its operand.
- It has the potential to yield a result with smaller arity than either operand, and may also make the cardinality smaller (why?)

A	B	C	D	E
1	a	4	v	t
2	s	5	b	y
3	z	6	n	r
4	s	4	b	y
5	s	7	v	r
6	a	5	c	r
7	z	4	b	y
8	a	3	n	t

$A(R) = 5$
 $|R| = 8$

$S = \text{Projection}[B,D,E](R)$

B	D	E
a	v	t
s	b	y
z	n	r
s	v	r
a	c	r
z	b	y
a	n	t

$A(S) = 3$
 $|S| = 7$

Core Relational Algebra: Renaming

18

- **Renaming** is a unary operator that determines the name that one or more attributes in the operand will have the result relation.
- It changes neither the arity nor the cardinality of the result relation.
- Its use is a formal necessity and it is often convenient in practice.

A	B	C	D	E
1	a	4	v	t
2	s	5	b	y
3	z	6	n	r
4	s	4	b	y
5	s	7	v	r
6	a	5	c	r
7	z	4	b	y
8	a	3	n	t

$A(R) = 5$
 $|R| = 8$

$S = \text{Renaming}[B \rightarrow X, E \rightarrow Y](R)$

A	X	C	D	Y
1	a	4	v	t
2	s	5	b	y
3	z	6	n	r
4	s	4	b	y
5	s	7	v	r
6	a	5	c	r
7	z	4	b	y
8	a	3	n	t

$A(S) = 5$
 $|S| = 8$

Extending the Core Relational Algebra

Derivable RA Operators

20

- Besides **Intersection**, the most useful derivable operators are:
 - ▶ **Joins**, in many different forms
 - ▶ **Division**
- The core RA extended with a set of useful derived operators is referred to as an extended RA.

RA Notation and Examples

Selection: Notation and Example

$$S := \sigma_c(R)$$

- c is a condition (as in 'if' statements) that refers to attributes of R .
- S contains all those tuples of R that satisfy c .

Selection: Notation and Example

23

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

$\text{JoeMenu} := \sigma_{\text{bar} = \text{'Joe's'}}(\text{Sells})$

JoeMenu		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75

Projection: Notation and Example

24

$$S := \pi_L(R)$$

- L is a list of attributes from the schema of R .
- S is constructed by
 - ▶ looking at each tuple of R ,
 - ▶ extracting the attributes on list L , in the order specified, and
 - ▶ creating from those components a new tuple for S if it is not a duplicate of a tuple already placed in S .

Projection: Notation and Example

25

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

$\text{BeerPrices} := \pi_{\text{beer}, \text{price}}(\text{Sells})$

BeerPrices	
beer	price
Bud	2.50
Miller	2.75
Miller	3.00

(Cartesian) Product: Notation and Example

26

$$R_3 := R_1 \times R_2$$

- Pair each tuple t of R_1 with each tuple t' of R_2 .
- Then, the concatenation $t+t'$ is a tuple of R_3 .
- The schema of R_3 is the concatenation of the attribute declarations of R_1 and R_2 .
- But beware, if an attribute A exists in R_1 and R_2 , use $R_1.A$ and $R_2.A$ to disambiguate
- If projecting out, use renaming.

(Cartesian) Product: Notation and Example

27

R ₁	
A	B
1	2
3	4

$$R_3 := R_1 \times R_2$$

R ₂	
C	D
5	6
7	8
9	10

R ₃			
A	B	C	D
1	2	5	6
1	2	7	8
1	2	9	10
3	4	5	6
3	4	7	8
3	4	9	10

Renaming: Notation and Example

28

$$R_2 := \rho_{R_2(A_1, \dots, A_n)}(R_1)$$

- Assigns to the resulting relation the name R_2 with attributes named A_1, \dots, A_n and the same tuples as R_1 .
- A simplified notation is

$$R_2(A_1, \dots, A_n) := R_1$$

Renaming: Notation and Example

29

Bars	
name	addr
Joe's	Maple St.
Sue's	River Rd.

$\text{Barz} := \rho_{\text{Bars}}(\text{bar}, \text{addr}) (\text{Bars})$

Barz	
bar	addr
Joe's	Maple St.
Sue's	River Rd.

θ -Join: Notation and Example

30

$$R_3 := R_1 \bowtie_{\theta} R_2$$

- Equivalent to taking the product

$$R := R_1 \times R_2$$

then applying a selection

$$R_3 := \sigma_{\theta}(R)$$

where θ can be any Boolean-valued condition

- If the only predicate that occurs in θ is '=', we refer to the θ -join as an equijoin.

θ -Join: Notation and Example

31

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Coors	3.00

Bars	
name	addr
Joe's	Maple St.
Sue's	River Rd.

$\text{BarInfo} := \text{Sells} \bowtie_{\text{Sells.bar}=\text{Bars.name}} \text{Bars}$

BarInfo				
bar	beer	price	name	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Joe's	Maple St.
Sue's	Bud	2.50	Sue's	River Rd.
Sue's	Coors	3.00	Sue's	River Rd.

(Natural) Join: Notation and Example

32

$$R_3 := R_1 \bowtie R_2$$

- A useful join variant that connects two relations by:
 - ▶ Checking for equality of attributes of the same name
 - ▶ Projecting out one copy only of each pair of equated attributes.
- So, equivalent to an equijoin to which we apply a projection that eliminates duplicate columns.

(Natural) Join: Notation and Example

33

Sells		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Coors	3.00

Barz	
bar	addr
Joe's	Maple St.
Sue's	River Rd.

BarInfo := Sells \bowtie Barz

BarInfo			
bar	beer	price	addr
Joe's	Bud	2.50	Maple St.
Joe's	Miller	2.75	Maple St.
Sue's	Bud	2.50	River Rd.
Sue's	Coors	3.00	River Rd.

Extended Projection: Notation and Example

34

- Using the same projection operator, we allow the list L to contain arbitrary expressions involving attributes:
 - ▶ Deriving attributes values, e.g., using arithmetic with name assignment, e.g.,

$A+B*C \rightarrow D$

- ▶ Introducing duplicate occurrences of the same attribute (typically with name disambiguation)

Extended Projection: Notation and Example

35

R	
A	B
1	2
3	4
5	6
7	8

$$S := \pi_{A+B \rightarrow C, A, A \rightarrow A2}(R)$$

S		
C	A	A2
3	1	1
7	3	3
11	5	5
15	7	7


Building Complex RA Expressions

Building Complex RA Expressions

37

$v := u - 1$
 $x := v^2$
 $y := x + 1$
 $z := y / 3$

$$z = ((u - 1)^2 + 1) / 3$$



- By combining operators with parentheses and precedence rules, we can build arbitrarily complex expressions.
- Three notations, just as in arithmetic:
 - ▶ Sequences of assignment statements with (typically simple) expressions in the right-hand side.
 - ▶ Single (typically complex) expressions that are equivalent to entire sequences of assignments.
 - ▶ Operator (or expression) trees where nodes are operators or values (variables or literals) and edges denote operator-operand relationships.

As Sequences of Assignment Statements

38

$$R_1 \bowtie_c R_2$$

\equiv

$$S := R_1 \times R_2$$

$$R_3 := \sigma_c(S)$$

- Creates temporary relation names for later reference
- Renaming can be implied by giving left-hand side relations a list of attributes.
- For example, the derivation of θ -join from product and selection can be expressed as a sequence of assignments.

As Single Expressions

39

$$R_1 \bowtie_C R_2$$

\equiv

$$R_3 := \sigma_C(R_1 \times R_2)$$

- Observe precedence and use parentheses, as follows (from highest to lowest):

$$[\sigma, \pi, \rho]$$

$$[\times, \bowtie]$$


$$n$$

$$[u, \setminus]$$

- Set difference is also denoted by $-$
- Renaming may need to be explicit

As Operator Trees

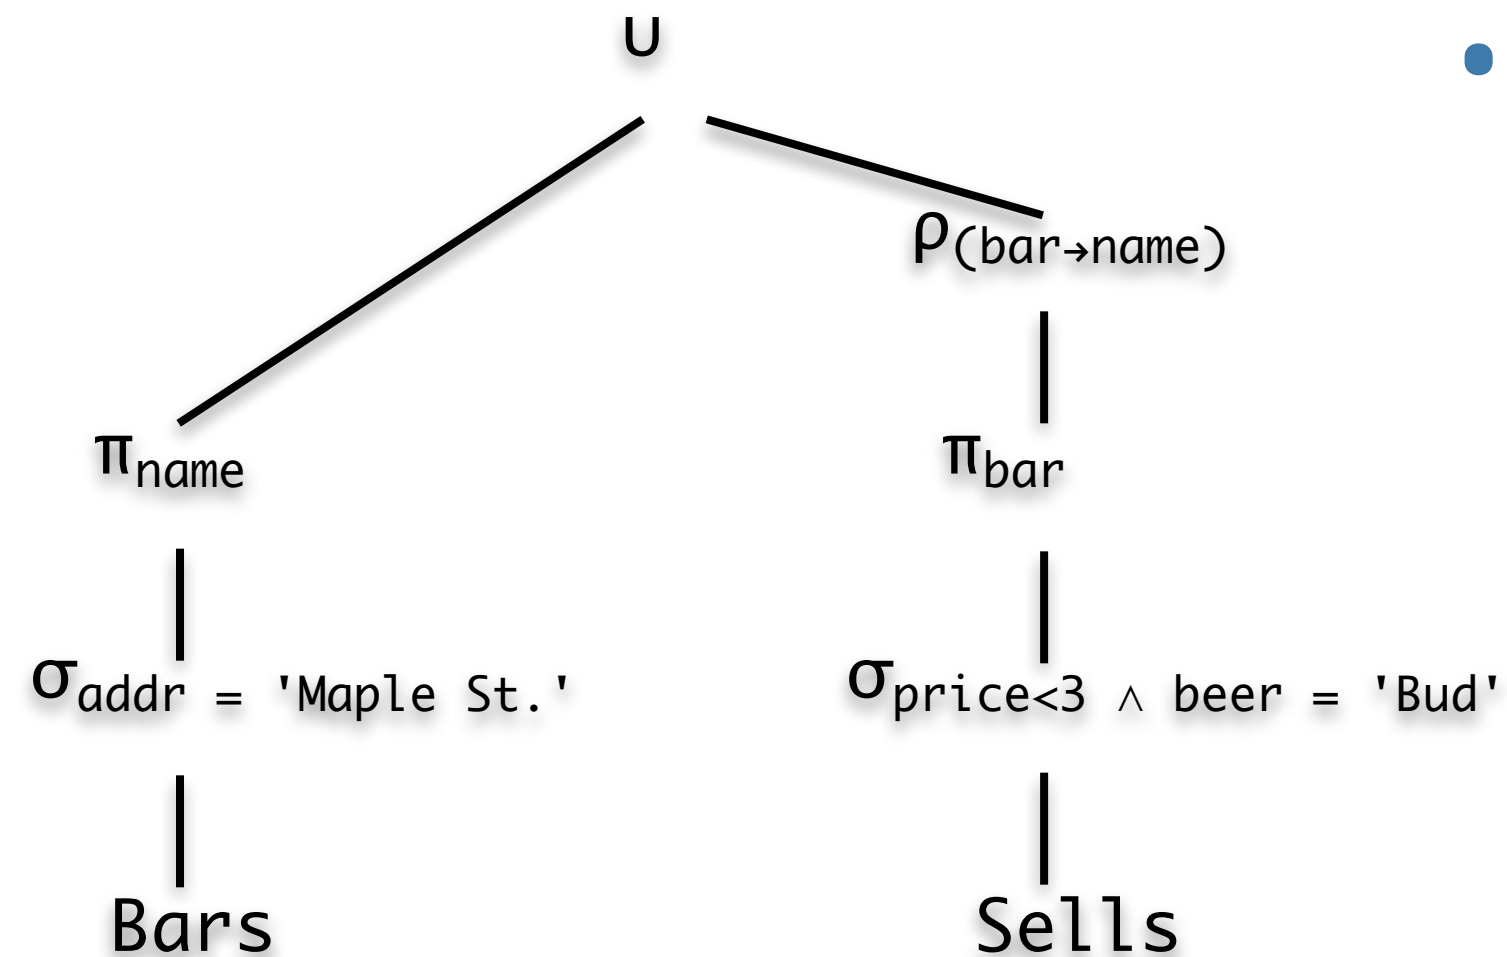
40



- Leaves are operands: either variables standing for relations or particular, constant relations.
- Interior nodes are operators, applied to their child or children.

More Example Operator Trees

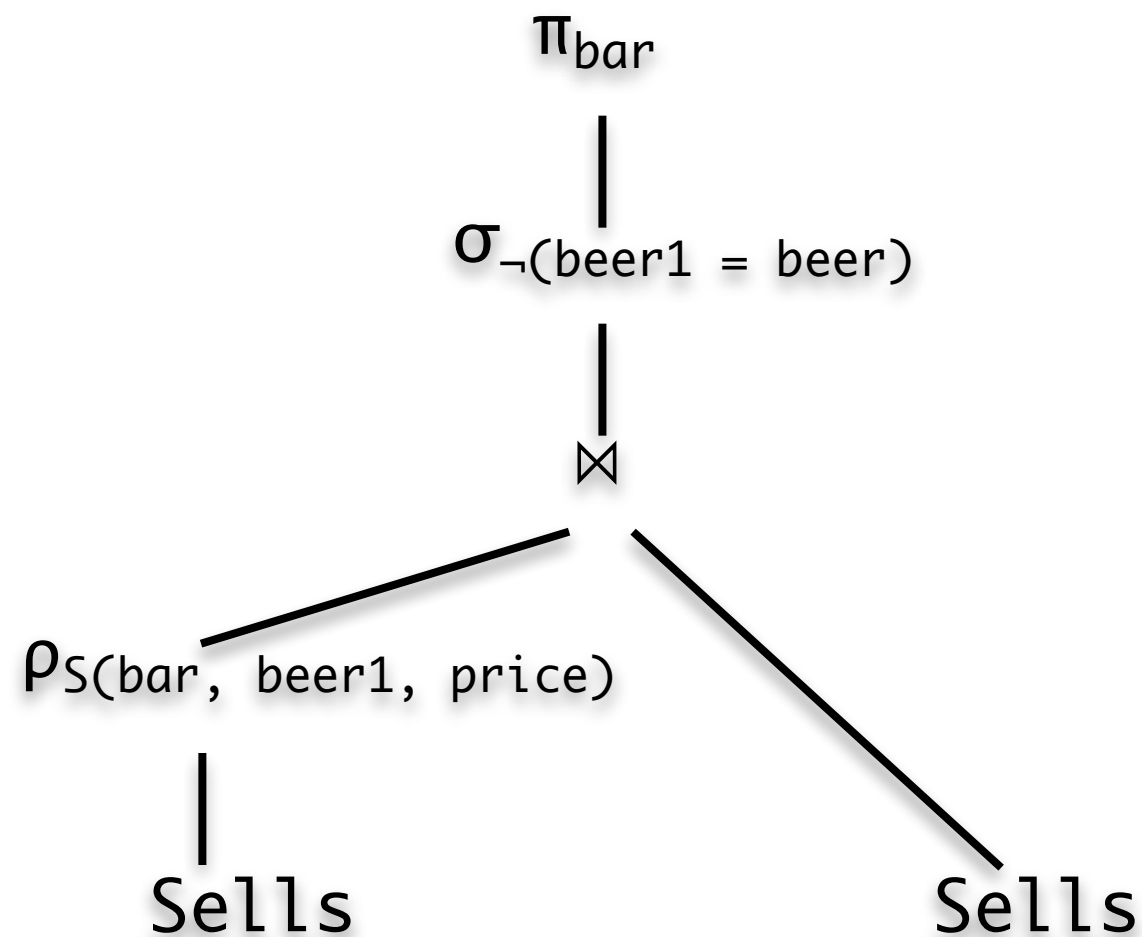
41



- Using the relations **Bars(name, addr)** and **Sells(bar, beer, price)**, find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.

More Example Operator Trees

42



- Using **Sells(bar, beer, price)**, find the bars that sell two different beers at the same price.
- Here is one strategy:
 - ▶ By renaming, define a copy of **Sells**, called **S(bar, beer1, price)**.
 - ▶ Then, the natural join of **Sells** and **S** consists of quadruples **(bar, beer, beer1, price)** such that the bar sells these two beers at this price.
 - ▶ Finally, we pick only pairs of beers that are different.

RA Type Inference

Inferring the Schema of the Result Relation

44

- **Union, Intersection, Difference**

- ▶ The schemas of the two operands must be the same, so it is also the schema of the result

- **Selection**

- ▶ the schema of the result is the same as the schema of the single operand

- **Projection**

- ▶ the list of attributes from the single operand tells us the (presumably different) schema of the result

- **Renaming**

- ▶ The operator is designed to define the result schema.

Inferring the Schema of the Result Relation

45

- **Product**

- ▶ The result schema is the list union of the attributes of the two relations.
- ▶ Use $R.A$, etc., to distinguish two attributes named A .

- **Theta-join**

- ▶ The same as for product

- **Natural join**

- ▶ The result schema is the set union of the attributes of the two relations.

RA on Bags

Relational Algebra on Bags

47

$$\{\{1,2,1,3\}\} \neq \{\{1,2,3\}\}$$

$$\{\{1,2,1,3\}\} = \{\{1,2,3,1\}\}$$

$$\{\{1,2,3\}\} = \{1,2,3\}$$

- Classical relational algebra defines relations as sets.
- But, it is often implemented (e.g., as the basis for SQL) with relations as bags.
- A bag (or multiset) is like a set, but
 - ▶ may contain duplicates, i.e., any element may appear more than once.
 - ▶ order remains immaterial.
- Sometimes written with double curly brackets, e.g., $\{\{1,2,1,3\}\}$.

Why Bags?

- SQL, the most important query language for relational databases, is actually a bag-based language.
- Some operations, like projection, are more efficient to implement over bags than over sets.

Why Bags?

49

- This is because duplicate elimination is an expensive operation to compute.
- This, in turn, is because one needs to process all the input tuples before one can produce any result tuple.
- Operations for which this is true (such as projection on sets) are referred to as blocking operations.

Operations on Bags

50

- Selection applies to each tuple, so its effect on bags is like its effect on sets, but duplicates are kept.
- Projection also applies to each tuple, but as a bag operator, we do not need to eliminate duplicates.
- Products and joins are done on each pair of tuples, so duplicates in bags have no effect on how we operate.

Selection on Bags

51

R	
A	B
1	2
5	6
1	2

$\sigma_{A+B<5}(R)$

<result>	
A	B
1	2
1	2

Projection on Bags

52

R	
A	B
1	2
5	6
1	2

$\pi_A (R)$

<result>
A
1
5
1

Product on Bags

53

R	
A	B
1	2
5	6
1	2

S	
B	C
3	4
7	8

$R \times S$

<result>			
A	R.B	S.B	C
1	2	3	4
1	2	7	8
5	6	3	4
5	6	7	8
1	2	3	4
1	2	7	8

θ -Join on Bags

54

R	
A	B
1	2
5	6
1	2

S	
B	C
3	4
7	8

 $R \bowtie_{R.B < S.B} S$

<result>			
A	R.B	S.B	C
1	2	3	4
1	2	7	8
5	6	7	8
1	2	3	4
1	2	7	8

Union, Intersection and Difference on Bags

55

- An element appears in the union of two bags the sum of the number of times it appears in each bag.

$$\{\{1,2,1\}\} \cup \{\{1,1,2,3,1\}\} = \{\{1,1,1,1,1,2,2,3\}\}$$

- An element appears in the intersection of two bags the minimum of the number of times it appears in either.

$$\{\{1,2,1,1\}\} \cap \{\{1,2,1,3\}\} = \{\{1,1,2\}\}$$

- An element appears in the difference of two bags A and B as many times as it appears in A, minus the number of times it appears in B.

$$\{\{1,2,1,1\}\} - \{\{1,2,3\}\} = \{\{1,1\}\}$$

Some Laws that Hold for Sets *Do* Hold for Bags

56

- Some, but not all algebraic laws that hold for sets also hold for bags.
- The commutative law for union
- Since addition is commutative, adding the number of times that a tuple appears in R and S does not depend on the order of R and S .

$$R \cup S = S \cup R$$

does hold for bags.

Some Laws that Hold for Sets *Do Not* Hold for Bags

57

- Set union is idempotent, meaning that, for sets:

$$S \cup S = S$$

- However, for bags, if x appears n times in S , then it appears $2n$ times in the result of bag union

- Thus, in general, for bags:

$$S \cup S \neq S$$

- For example

$$\{\{1\}\} \cup \{\{1\}\} = \{\{1,1\}\} \neq \{\{1\}\}.$$

A Powerful Extended RA

Extended Relational Algebras

59

- The relational algebra admits of various extensions.
- The main kinds of extended operations are:
 - ▶ duplicate elimination
 - ▶ sorting
 - ▶ several versions of join
 - ▶ aggregation and grouping
 - ▶ division

Extended RA: Notation

60

- δ : duplicate elimination
- τ : sorting
- γ : aggregation/grouping
- \Join , \Join_L , \Join_R : outer join, left outer join, right outer join
- \Join_L , \Join_R : left semijoin, right semijoin
- \Join_{\neg} : antijoin
- \div : division
- Note that the notations for extended RA operators is much less agreed upon than for core RA operators.

Extended RA: Duplicate Elimination

61

$$R_2 := \delta(R_1)$$

- R_2 consists of one copy only of each tuple that appears in R_1 .
- Note that this may be computationally expensive to implement.

R_1	
A	B
1	2
5	6
1	2

$$R_2 := \delta(R_1)$$

R_2	
A	B
1	2
5	6

Extended RA: Sorting

62

$$R_2 := \tau_L(R_1)$$

- where $L = (a_1, d_1), \dots, (a_n, d_n)$ is a list of attribute/direction pairs such that each a_i is an attribute in R_1 and each d_i is either 'A' (for ascending order) or 'D' (for descending order).
- R_2 is the list of tuples of R_1 sorted first on the value of the first attribute in L , then on the second attribute in L , and so on, with ties broken arbitrarily
- τ is the only operator whose result is more constrained than a set.

R_1	
X	Y
1	2
5	6
1	5

$$S := \tau_{(X,A),(Y,D)}(R)$$

R_2	
X	Y
1	5
1	2
5	6

Extended RA: Aggregation

63

- Aggregation functions, strictly, are not operations in the relational algebra: their operands are not relations.
- Rather, they apply to (entire) columns (i.e., sets, bags, or lists) of a table and produce a single result.
- They map a bulk/collection value into a scalar value.
- The most commonly-used aggregation functions are:
 - ▶ COUNT
 - ▶ SUM
 - ▶ AVG
 - ▶ MIN
 - ▶ MAX

R	
A	B
1	3
3	4
3	2

$$\text{COUNT}(R.A) = 3$$

$$\text{SUM}(R.A) = 7$$

$$\text{AVG}(R.B) = 3$$

$$\text{MIN}(R.A) = 1$$

$$\text{MAX}(R.B) = 4$$

Extended RA: Grouping

64

$$R_2 := \gamma_L(R_1)$$

- where L is a list of elements that are:
 - ▶ either individual (so called, grouping) attributes
 - ▶ or the application of an aggregation function on an attribute A of R_1 , with (re)naming, as in extended projection.

Extended RA: Grouping

65

- To obtain the result, we rearrange the input tuples according to all the grouping attributes on the list L .
- That is, we form one group for each distinct list of values in the input tuples, taking into account the grouping attributes in L .
- Over each such group, we compute the aggregation function for each such function in L .
- The result has one tuple for each group and has columns as follows:
 - ▶ One for each grouping attribute in L
 - ▶ One for each aggregation function in L

Extended RA: Grouping

66

R		
A	B	C
1	2	3
4	5	6
1	2	5



<tmp>		
A	B	C
1	2	3
1	2	5
4	5	6



$S := \gamma_{A,B, \text{AVG}(C) \rightarrow X} (R)$

S		
A	B	X
1	2	4
4	5	6

Extended RA: Outer Join, Left and Right

67

- When we refer to join, without qualification, we mean, more precisely, an inner join.
- If we join $R \bowtie_c S$, a tuple of R that has no tuple of S with which it joins is said to be dangling and does not contribute to the result.
- Similarly for a tuple of S that has no tuple of R with which it joins.

Extended RA: Outer Join, Left and Right

68

- A join that emits such dangling tuples into the result relation by padding them with NULL values is called an outer join.
- If a join condition is given, we write $R \bowtie_{\theta} S$; otherwise, we assume a natural outer join; both of these are referred to as full outer joins.
- An outer join in which only the tuples from the left (resp., right) operand are padded is called a left (resp., right) outer join, and we write $R \ltimes_{\theta} S$ (only pad R tuples) and $R \rtimes_{\theta} S$ (only pad S tuples), resp..

R	
A	B
1	2
4	5

S	
B	C
2	3
6	7

$$T := R \bowtie S$$

T		
A	B	C
1	2	3
4	5	NULL
NULL	6	7

Extended RA: Left/Right Semijoin

69

- A left semijoin is defined as:

$$R \ltimes S \equiv \pi_L(R \bowtie S)$$

where L is the list of attributes in R

- A right semijoin is defined as:

$$R \ltimes' S \equiv \pi_{L'}(R \bowtie S)$$

where L' is the list of attributes in S

- So a left (resp., right) semijoin is a join that only retains in the result relation the columns of the left (resp., right) operand.

R	
A	B
1	2
5	6
1	2

$$T := R \ltimes S$$

T	
A	B
1	2
5	6

S	
B	C
2	4
6	8

$$U := R \ltimes' S$$

U	
B	C
2	4

Extended RA: Antijoin

70

- An antijoin is the complement of a left semijoin, i.e., it is defined as:

R	
A	B
1	2
5	6
3	4

S	
B	C
2	4
6	8

$$T := R \triangleright S$$

$$R \triangleright S \equiv R \setminus (R \bowtie S)$$

T	
A	B
3	4

Extended RA: Division

71

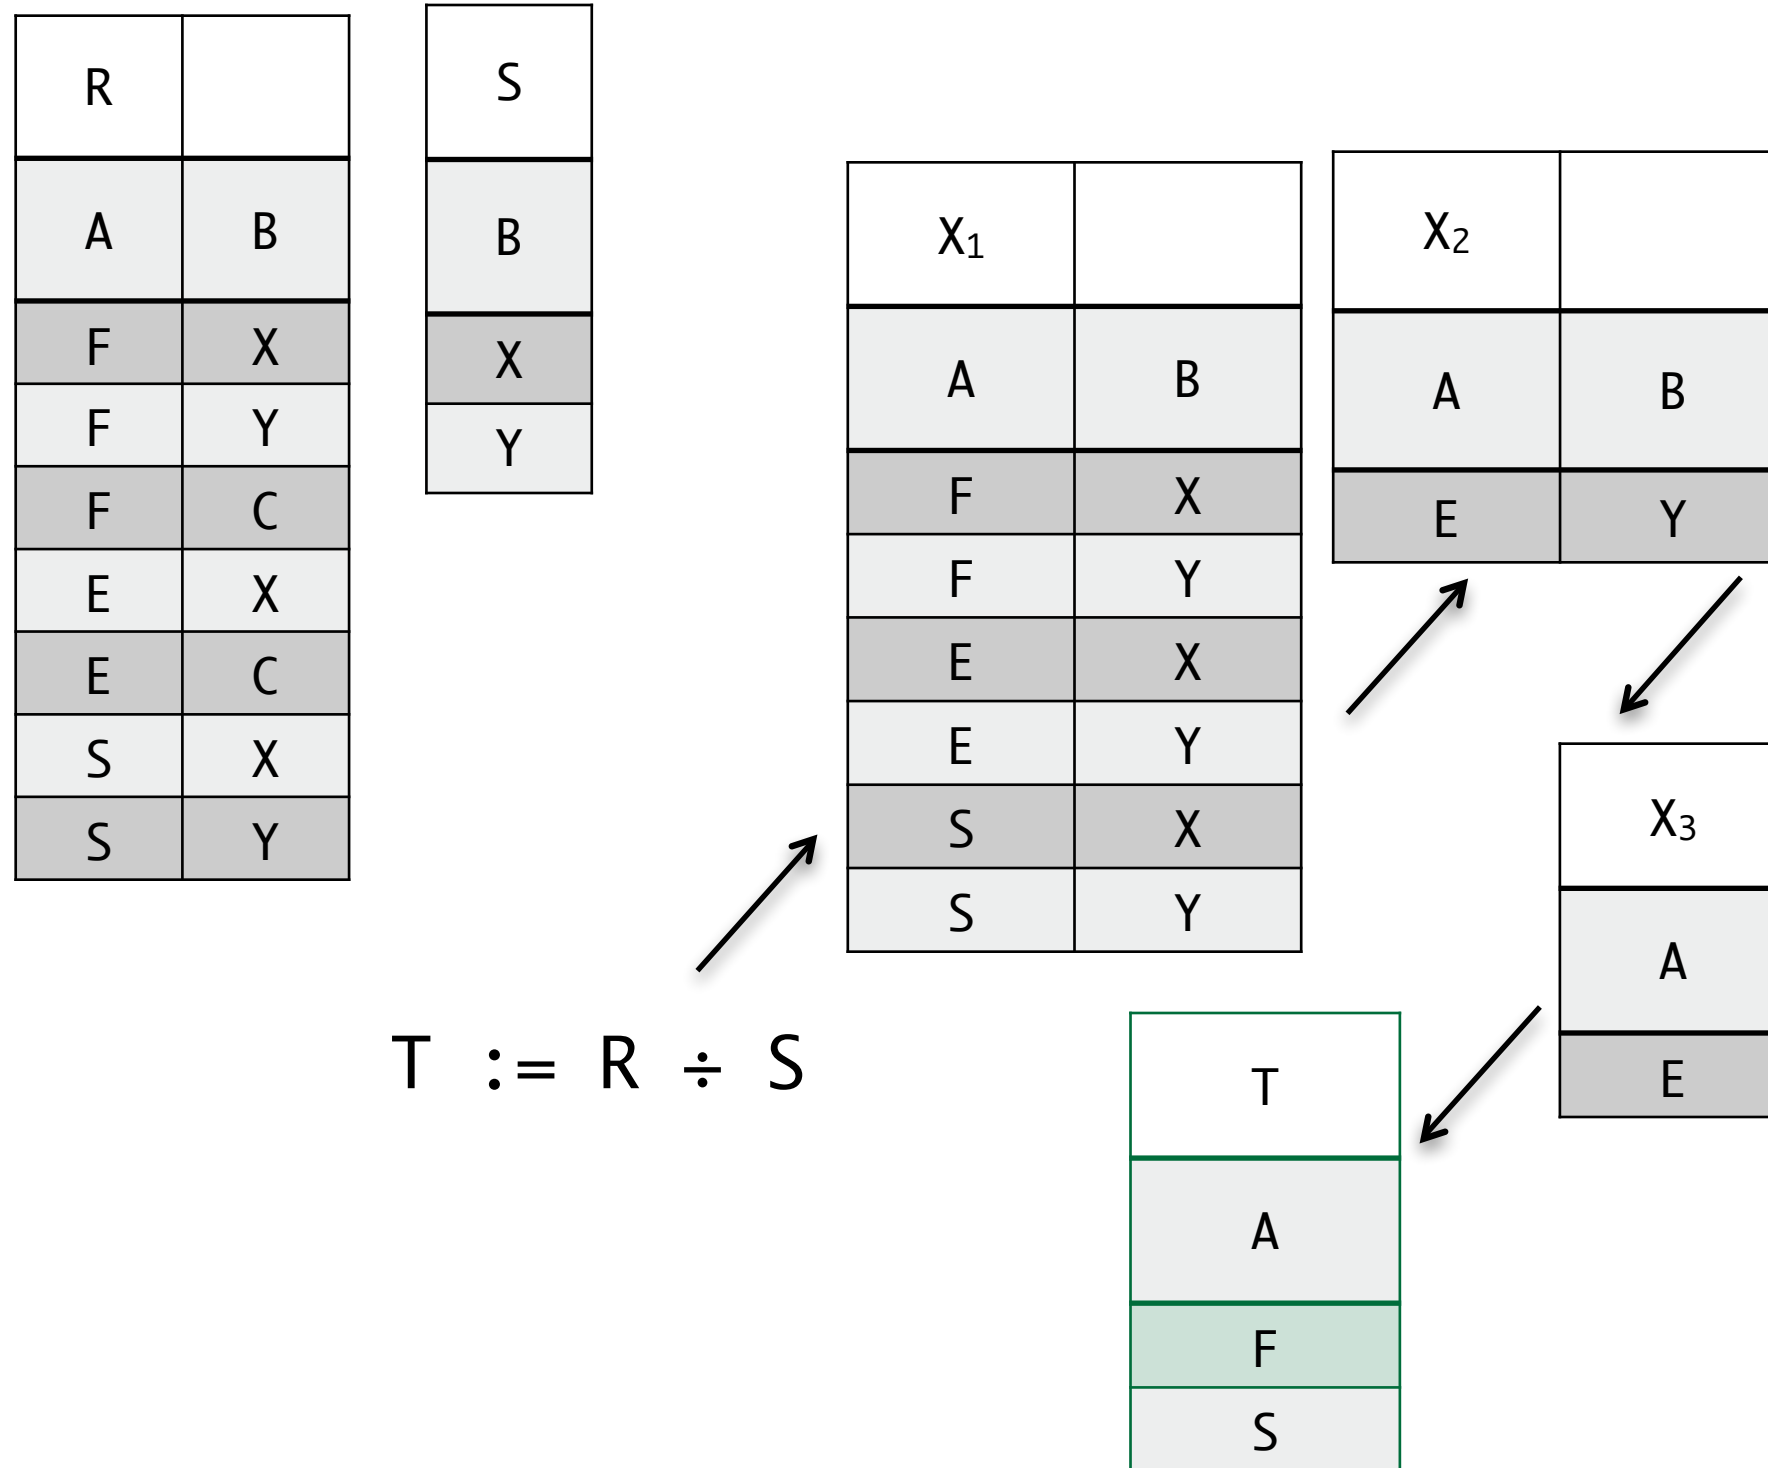
- Let R have two attributes $R.x$ and $R.y$, and S one attribute $S.y$ with the same domain as $R.y$.
- The division of R by S is the set of all $R.x$ values (as unary tuples) such that for every $S.y$ value there is a tuple $(R.x, R.y)$ in R .

$$\begin{aligned} T &:= R \div S \\ &\equiv \pi_L(R) \setminus (\pi_L((\pi_L(R) \times S) \setminus R)) \\ &\equiv \begin{aligned} X_1 &:= \pi_L(R) \times S \\ X_2 &:= X_1 \setminus R \\ X_3 &:= \pi_L(X_2) \\ T &:= \pi_L(R) \setminus X_3 \end{aligned} \end{aligned}$$

where L is the list of attribute names unique to R

Extended RA: Division

72



Translating SQL onto Relational Algebra

Direct Translation of Core SQL to RA

74

- There is a useful way of computing in your mind what the result of a SQL query would be.
- It is a high-level procedure for obtaining a direct translation from SQL to RA.
- The basic SQL query blocks map to RA as follows:
 - ▶ Form the (cascaded binary) product of the relations in the **FROM** clause
 - ▶ To this result, apply a selection operator whose predicate is the one in the **WHERE** clause
 - ▶ To this result, apply a projection whose attribute list is that in the **SELECT** clause
 - ▶ If **DISTINCT** is requested, apply a duplicate removal operator

R	
A	B
1	2
5	6
1	2

S	
B	C
3	4
7	8

<result>	
A	C
1	4
1	8
5	8

SELECT
DISTINCT R.A, S.C
FROM R, S
WHERE R.B < S.B
≡
 $\delta(\pi_{R.A, S.C}(\sigma_{R.B < S.B}(R \times S)))$

SQL with Grouping and Aggregation

75

- With grouping and aggregation, we extend the previous semantics as follows:
 - ▶ Before applying the projection, tuples are grouped by LG , where LG is the list of grouping attributes in the **GROUP BY** clause
 - ▶ When the projection is applied, the aggregation functions are applied per group, so that one tuple is generated for each group

R	
A	B
1	2
5	6
1	2

S	
B	C
3	4
7	8

<result>	
A	SUM_C
1	12
5	8

```
SELECT    R.A, SUM(S.C) AS SUM_C
FROM      R,S
WHERE     R.B < S.B
GROUP BY  R.A
```

≡

$\gamma_{R.A, \text{SUM}(S.C) \rightarrow \text{SUM_C}}(\sigma_{R.B < S.B}(R \times S))$

Importance of SQL to RA Translation

76

- Note that, while SQL has a declarative core, RA is wholly procedural.
- The order in which operations are applied is strictly specified by the structure of the relational-algebraic expression (i.e., parentheses and operator precedence).
- The ability to translate SQL to relational algebra is, therefore, a bridge from a declarative to a procedural language.
- Since RA is an algebra, one can use properties such as idempotence, commutativity, associativity as well as more specific equivalence laws to transform an expression E into another expression E' .
- The result of evaluating E' is the same as the result of evaluating E , but E' could be more efficient to compute.

Importance of SQL to RA Translation

77

- Consider calculating by hand

1982736458+
1982736458+
1982736458+
1982736458+
1982736458+
1982736458

compared with

$6 * 1982736458$

- One of the goals of an optimizer is to apply such laws to the direct translation of an SQL query to an RA expression and search (by rewriting) for a more efficient, equivalent RA expression.
- This search is often guided by the insight that we should reduce the size of a relation (in bytes, i.e., the average width of a tuple times the average number of tuples) of intermediate results as early as possible.

SQL as a Query Language

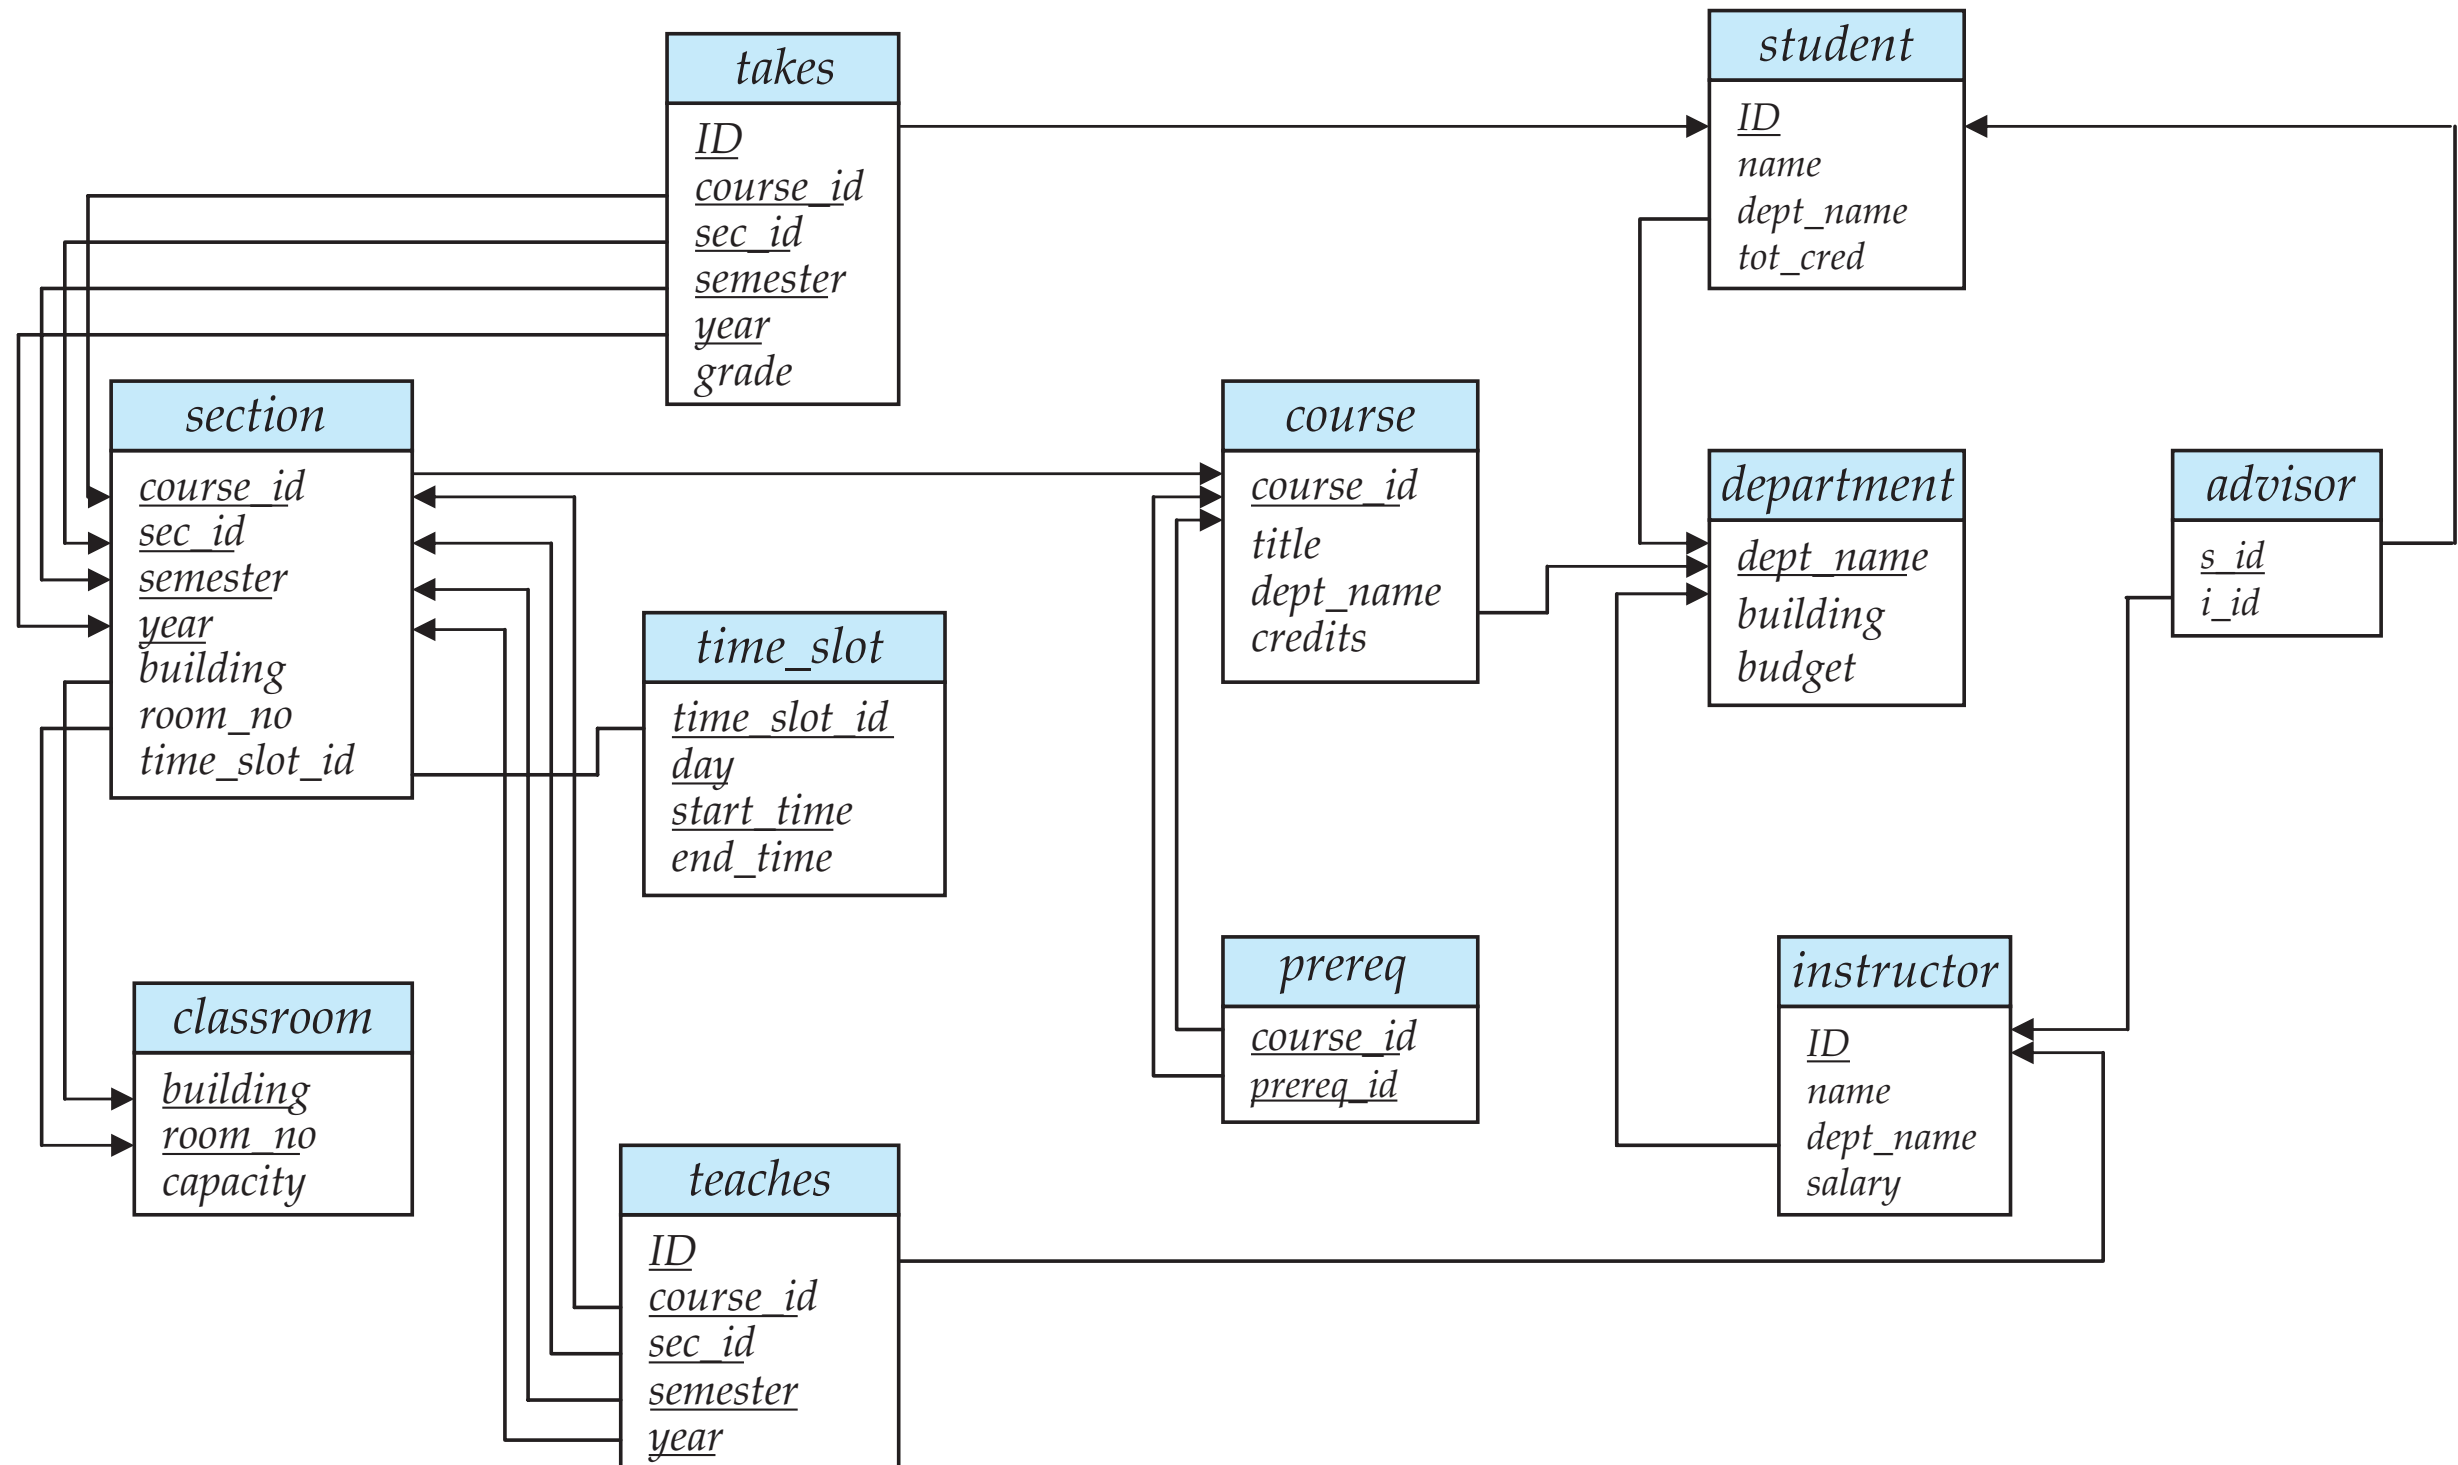
The History of SQL

79

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory in the 1970s
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - ▶ SQL-86, SQL-89, SQL-92
 - ▶ SQL:1999, SQL:2003, SQL:2008
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
 - ▶ Some examples here may not work on some DBMSs.

An Example Database

80



Some (Not All) Relation Schemas

81

```
create table department (  
    dept_name    varchar(20),  
    building     varchar(15),  
    budget       numeric(12,2)  
                check (budget > 0),  
    primary key (dept_name)  
);  
  
create table course (  
    course_id    varchar(8),  
    title        varchar(50),  
    dept_name    varchar(20),  
    credits      numeric(2,0)  
                check (credits > 0),  
    primary key (course_id),  
    foreign key (dept_name)  
        references department  
        on delete set null  
);  
  
create table instructor (  
    ID           varchar(5),  
    name         varchar(20) not null,  
    dept_name    varchar(20),  
    salary       numeric(8,2)  
                check (salary > 29000),  
    primary key (ID),  
    foreign key (dept_name)  
        references department  
        on delete set null  
);  
  
create table section (  
    course_id    varchar(8),  
    sec_id       varchar(8),  
    semester     varchar(6)  
                check (semester in ('Fall', 'Winter', 'Spring', 'Summer')),  
    year         numeric(4,0)  
                check (year > 1701 and year < 2100),  
    building     varchar(15),  
    room_number  varchar(7),  
    time_slot_id varchar(4),  
    primary key (course_id, sec_id, semester, year),  
    foreign key (course_id)  
        references course  
        on delete cascade,  
    foreign key (building, room_number)  
        references classroom  
        on delete set null  
);  
  
create table takes (  
    ID           varchar(5),  
    course_id    varchar(8),  
    sec_id       varchar(8),  
    semester     varchar(6),  
    year         numeric(4,0),  
    grade        varchar(2),  
    primary key (ID, course_id, sec_id, semester, year),  
    foreign key (course_id, sec_id, semester, year)  
        references section  
        on delete cascade,  
    foreign key (ID)  
        references student  
        on delete cascade  
);
```

Some (Not All) Relation Instances

82

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

course

course_id	title	dept_name	credits
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3
HIS-351	World History	History	3
MU-199	Music Video Production	Music	3
PHY-101	Physical Principles	Physics	4

department

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

prereq

course_id	prereq_id
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

section

course_id	sec_id	semester	year	building	room_number	time_slot_id
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	E
CS-190	2	Spring	2009	Taylor	3128	A
CS-315	1	Spring	2010	Watson	120	D
CS-319	1	Spring	2010	Watson	100	B
CS-319	2	Spring	2010	Taylor	3128	C
CS-347	1	Fall	2009	Taylor	3128	A
EE-181	1	Spring	2009	Taylor	3128	C
FIN-201	1	Spring	2010	Packard	101	B
HIS-351	1	Spring	2010	Painter	514	C
MU-199	1	Spring	2010	Packard	101	D
PHY-101	1	Fall	2009	Watson	100	A

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009
32343	HIS-351	1	Spring	2010
45565	CS-101	1	Spring	2010
45565	CS-319	1	Spring	2010
76766	BIO-101	1	Summer	2009
76766	BIO-301	1	Summer	2010
83821	CS-190	1	Spring	2009
83821	CS-190	2	Spring	2009
83821	CS-319	2	Spring	2010
98345	EE-181	1	Spring	2009

Core SQL

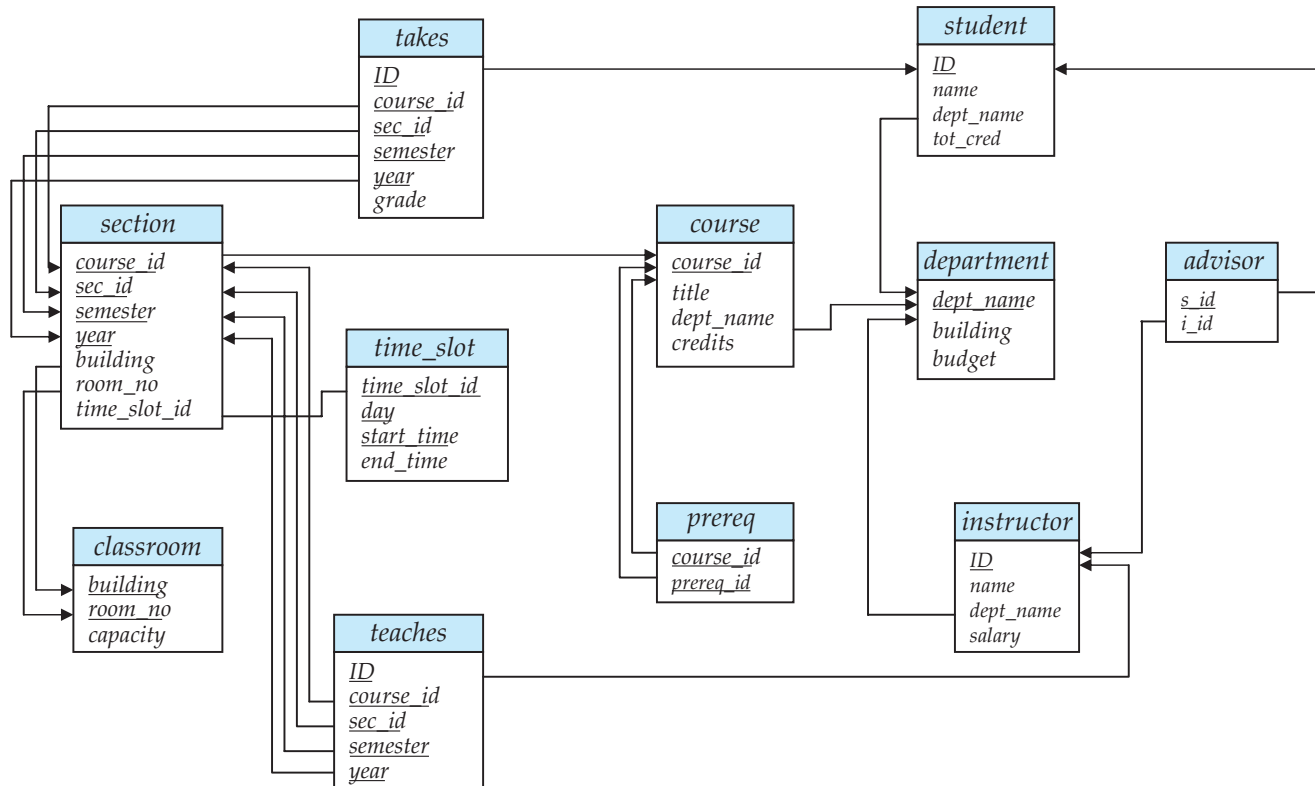
- A typical SQL query has the form:

```
SELECT  A1, A2, . . . , An  
FROM    r1, r2, . . . , rm  
WHERE   P;
```

- ▶ A_i denotes an attribute
 - ▶ R_i denotes a relation
 - ▶ P is a predicate.
- The result of an SQL query is a relation.

SQL: The SELECT Clause

85




- The SELECT clause lists the attributes desired in the result of a query.
- It corresponds to the projection operation of the relational algebra.
- For example, to find the names of all instructors:

```
SELECT name  
FROM instructor;
```

- SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
 - ▶ E.g. Name ≡ NAME ≡ name
- Some people use bold font wherever we use upper case.

SQL: The SELECT Clause

86



- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword DISTINCT after select.
- To find the names of all departments with instructor, and remove duplicates:


```
SELECT DISTINCT name  
FROM instructor;
```

- The keyword ALL specifies that duplicates must not be removed.

```
SELECT ALL name  
FROM instructor;
```

SQL: The SELECT Clause

87



- An asterisk in the select clause denotes "all attributes"

```
SELECT *  
FROM instructor;
```


- The select clause can contain arithmetic expressions ranging over the operations +, −, *, and /, and in which constants or attributes of tuples may occur as operands.
- The query

```
SELECT ID, name, salary/12  
FROM instructor;
```

would return a relation that is the same as the instructor relation, except that the value of the attribute salary is divided by 12.

SQL: The WHERE Clause

88




- The WHERE clause specifies conditions that every tuple in the result must satisfy.
- It corresponds to the selection operation of the relational algebra.
- To find all instructors in the Computer Science department with salary > 80000:

```
SELECT name
FROM   instructor
WHERE  dept_name = 'Comp. Sci'
AND    salary > 80000;
```

- Comparison results can be combined using the logical connectives AND, OR, and NOT.
- Comparisons can be applied to results of arithmetic expressions.

SQL: The FROM Clause

89



- The FROM clause lists the relations involved in the query.
- It corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product instructor X teaches:


```
SELECT *  
FROM    instructor, teaches;
```

generates every possible (instructor, teaches) pair, with all attributes from both relations.

- Cartesian product not is very useful directly, but useful combined with the where-clause condition (selection operation in relational algebra).

SQL: Set Operations: Union

90




- To find courses that ran in Fall 2009 or in Spring 2010:

```
(SELECT course_id
FROM section
WHERE sem = 'Fall' AND year = '2009')
UNION
(SELECT course_id
FROM section
WHERE sem = 'Spring' AND year = '2010')
;
```

SQL: Set Operations: Intersection

91

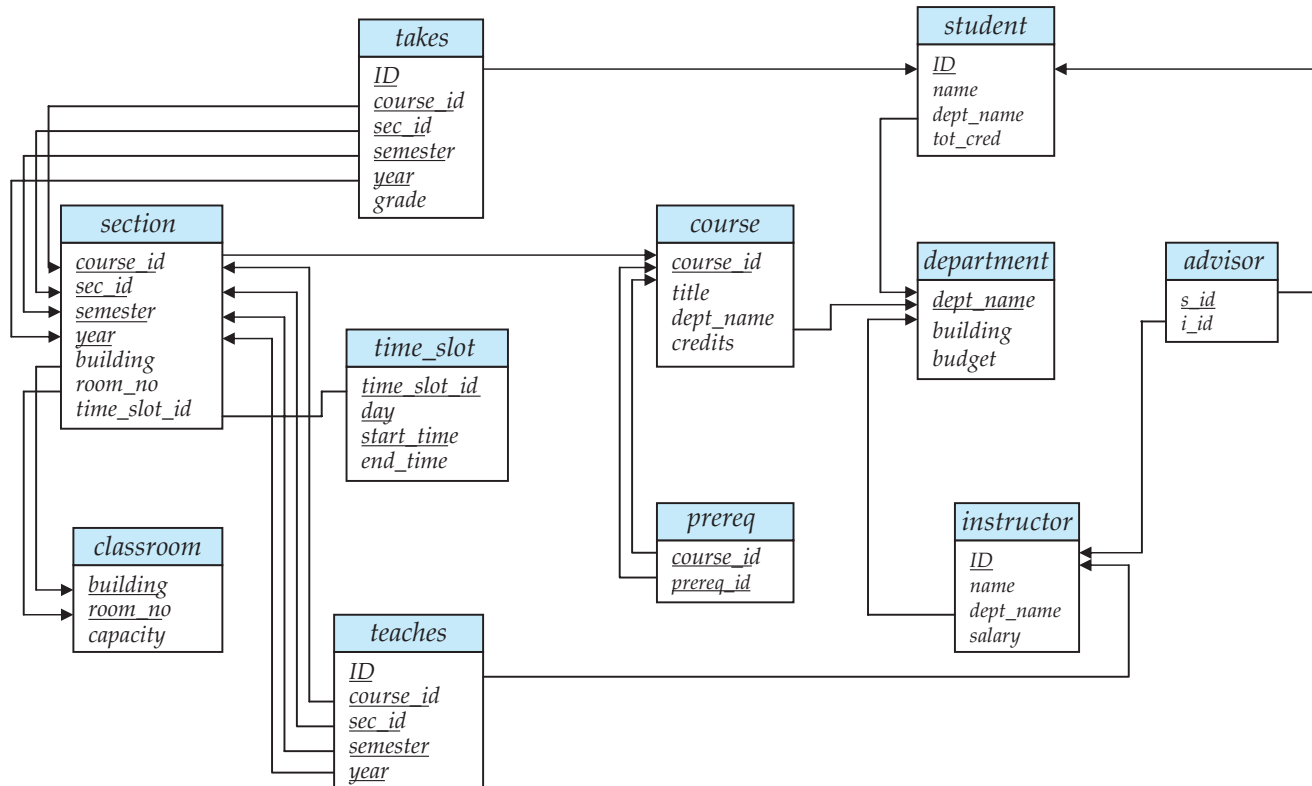


- To find courses that ran in Fall 2009 and in Spring 2010:

```
(SELECT course_id
FROM section
WHERE sem = 'Fall' AND year = '2009')
INTERSECT
(SELECT course_id
FROM section
WHERE sem = 'Spring' AND year = '2010')
;
```

SQL: Set Operations: Difference

92



- To find courses that ran in Fall 2009 but not in Spring 2010:

```
(SELECT course_id
FROM section
WHERE sem = 'Fall' AND year = '2009')
EXCEPT
(SELECT course_id
FROM section
WHERE sem = 'Spring' AND year = '2010')
;
```

- Oracle uses MINUS rather than EXCEPT.


SQL: Retaining Duplicates in Set Operations

93

- UNION, INTERSECT, and EXCEPT automatically eliminate duplicates
- To retain duplicates, we use the corresponding multiset/bag versions UNION ALL, INTERSECT ALL and EXCEPT ALL.
- Suppose a tuple occurs m times in R and n times in S , then, it occurs:
 - ▶ $m+n$ times in $R \text{ UNION ALL } S$
 - ▶ $\min(m,n)$ times in $R \text{ INTERSECT ALL } S$
 - ▶ $\max(0, m-n)$ times in $R \text{ EXCEPT ALL } S$

SQL: Joins

94



- For all instructors who teach some course, find their names and the course ID of the courses they teach:


```
SELECT instructor.name, teaches.course_id
FROM   instructor, teaches
WHERE  instructor.ID = teaches.ID;
```

- Find the course ID, semester, year and title of each course offered by the Comp. Sci. department:

```
SELECT s.course_id, s.semester, s.year, c.title
FROM   section s, course c
WHERE  s.course_id = c.course_id
AND    c.dept_name = 'Comp. Sci.'
```

SQL: Joins

95




- We can also use the JOIN/ USING keywords explicitly.
- For all instructors who teach some course, find their names and the course ID of the courses they teach (same as before):

```
SELECT name, course_id
FROM   instructor
JOIN   teaches
USING (ID);
```

SQL: Natural Join

96

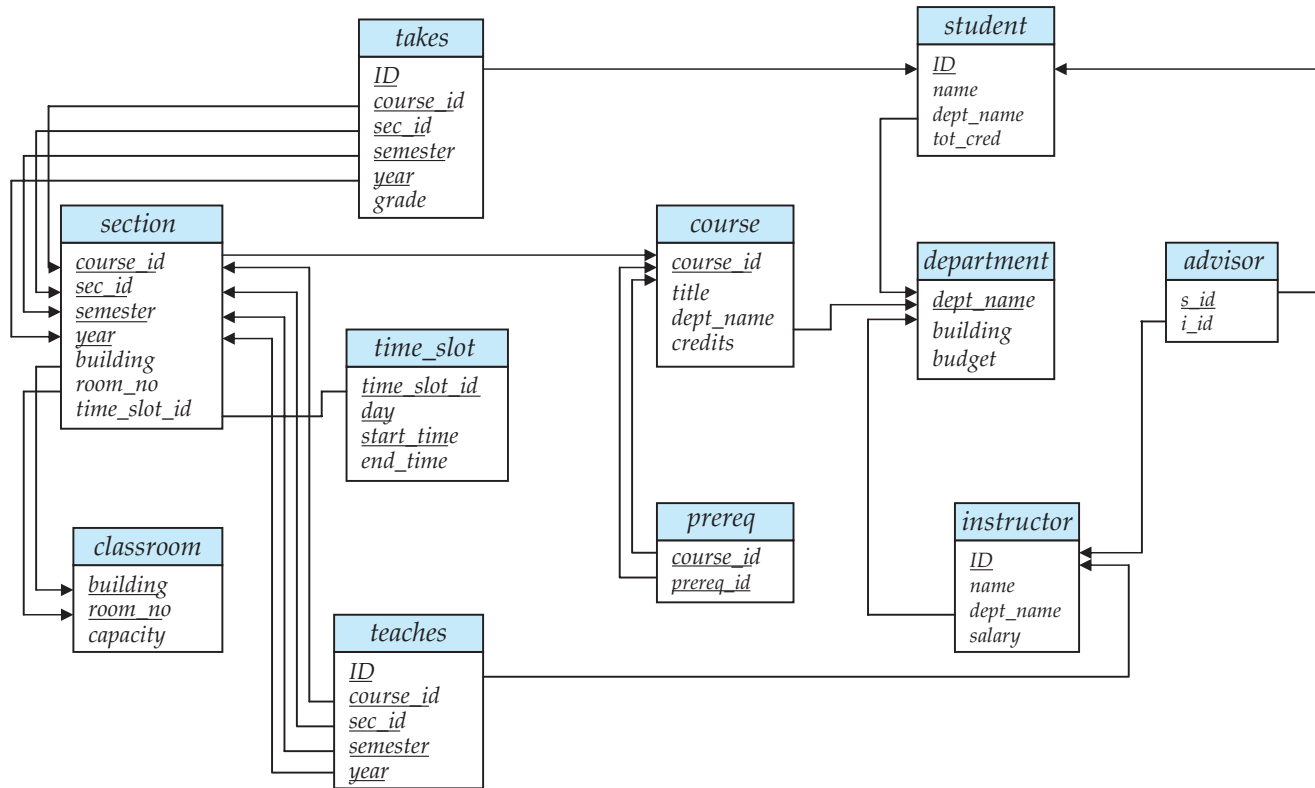


- Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column:

```
SELECT *
      FROM instructor
NATURAL JOIN teaches;
```


SQL: Natural Join


97



- There is danger in natural join!
- Beware of unrelated attributes with same name that get equated incorrectly!

SQL: Natural Join

98



- List the names of instructors along with the titles of courses that they teach:

- Incorrect version (erroneously makes course.dept_name = instructor.dept_name)

```
SELECT name, title
FROM   instructor NATURAL JOIN
       teaches NATURAL JOIN
       course;
```

- Correct version (an additional, specific join condition)


```
SELECT name, title
FROM   instructor NATURAL JOIN teaches, course
WHERE  teaches.course_id = course.course_id;
```

- Another correct version (different joins)

```
SELECT name, title
FROM   (instructor NATURAL JOIN teaches)
JOIN   course USING (course_id);
```

SQL: Renaming

99



- SQL allows renaming relations and attributes using the AS clause of the form:

<old-name> AS <new-name>

- E.g.

```
SELECT ID, name, salary/12 AS monthly_salary
FROM   instructor;
```

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci':


```
SELECT DISTINCT T.name
FROM   instructor AS T, instructor AS S
WHERE  T.salary > S.salary
AND    S.dept_name = 'Comp. Sci.';
```

- Keyword AS is optional and may be omitted
- Keyword AS **must** be omitted in Oracle SQL

SQL: Strings

SQL: String Operations

101




- SQL includes a string-matching operator for comparisons on character strings.
- The operator LIKE uses patterns that are specified using three special characters:
 - ▶ percent (%), which matches any substring
 - ▶ underscore (_), which matches any character
 - ▶ use backslash (\) as the escape character if you want to use % or _ as actual characters.
- Find the names of all instructors whose name includes the substring "dar":

```
SELECT name
FROM   instructor
WHERE  name LIKE '%dar%';
```

SQL: String Operations

102



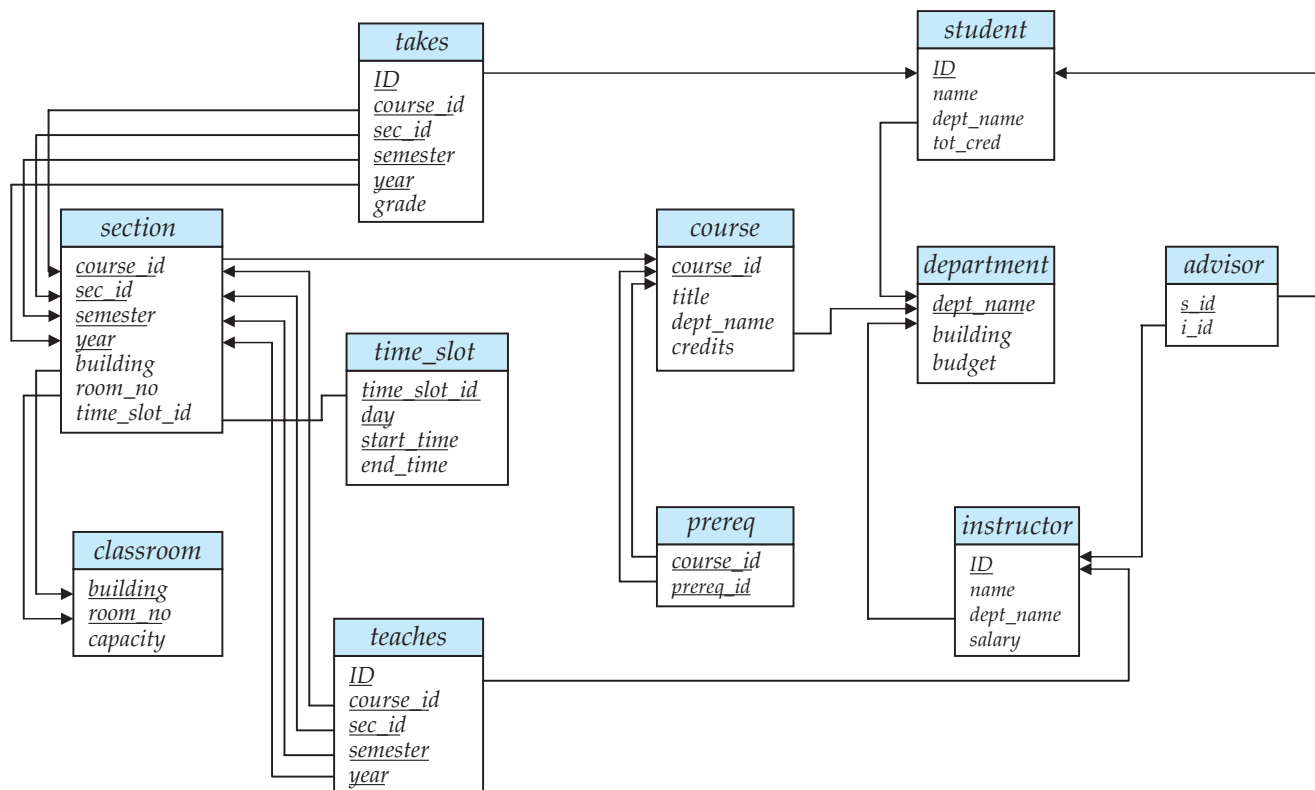
- Patterns are case-sensitive (as are literals).
- Some examples:
 - ▶ '100\%' matches the string '100%'
 - ▶ 'Intro%' matches any string beginning with "Intro".
 - ▶ '%Comp%' matches any string containing "Comp" as a substring.
 - ▶ '___' matches any string of exactly three characters.
 - ▶ '___ %' matches any string of at least three characters.

- SQL supports a variety of string operations such as
 - ▶ concatenation (using '||')
 - ▶ converting from upper to lower case (and vice versa)
 - ▶ finding string length, extracting substrings, etc.
- Consult the DBMS manual for the details.

SQL: Extended Features in the WHERE Clause

SQL: WHERE Clause Extended Features: Interval Membership

105




- SQL includes a BETWEEN comparison operator that can appear in a WHERE clause
- For example, find the names of all instructors such that $\$90,000 \leq \text{salary} \leq \$100,000$):

```
SELECT name
FROM   instructor
WHERE  salary BETWEEN 90000 AND 100000;
```

SQL: WHERE Clause Extended Features: Tuple Comparison

106

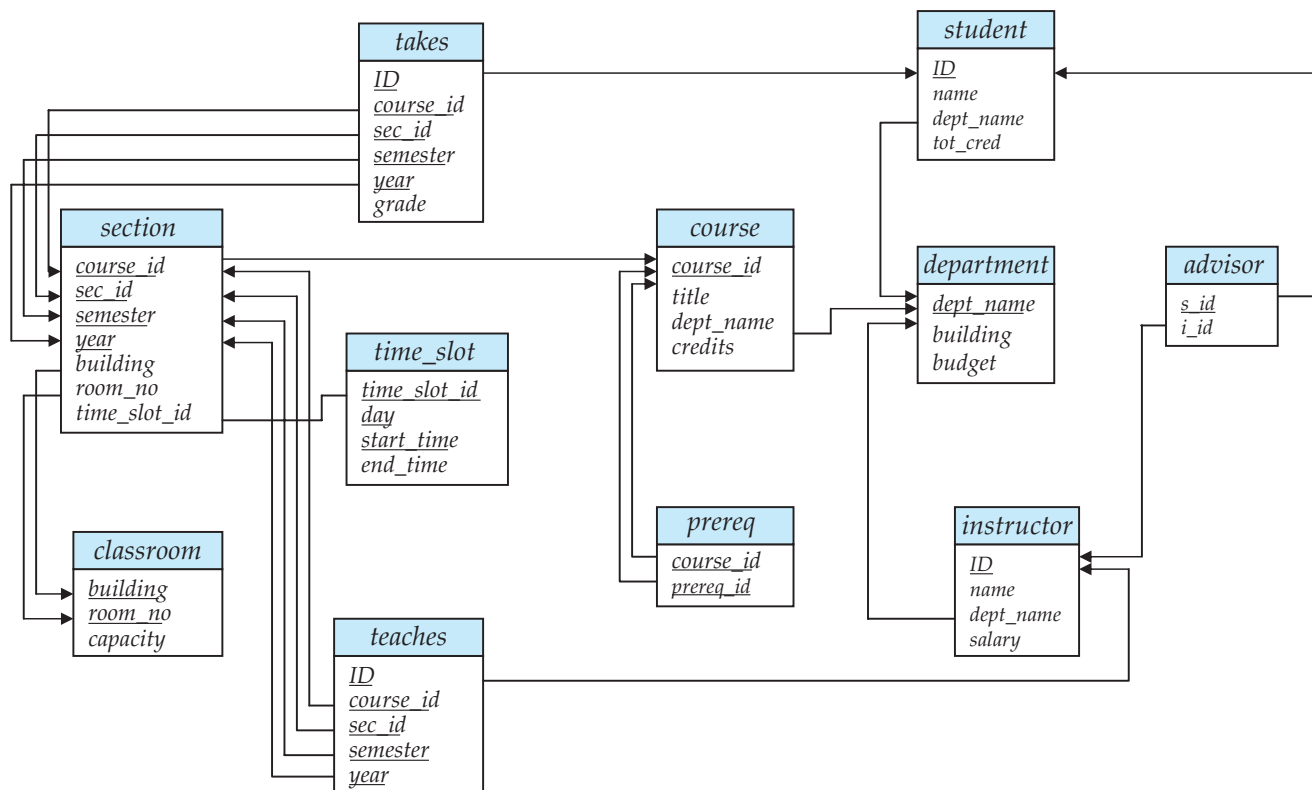


- SQL can do pairwise comparison of a list of elements:

```
SELECT name, course_id
FROM   instructor i, teaches t
WHERE  (i.ID,t.dept_name) = (t.ID,'Biology');
```

SQL: WHERE Clause Extended Features: NULL values

107



- It is possible for tuples to have a null value, denoted by NULL, for some of their attributes
- NULL denotes a value that may exist but hasn't been inserted into the tuple or a value that does not exist.
- The result of any arithmetic expression involving null is null (e.g., 5 + NULL returns NULL)
- The predicate IS NULL can be used to check for null values.
- For example, find all instructors whose salary is null:

```
SELECT name
FROM   instructor
WHERE  salary IS NULL;
```

SQL: WHERE Clause Extended Features:

Tests and Comparisons on NULL and UNKNOWN

108

- Any comparison with NULL returns UNKNOWN, e.g., $5 < \text{NULL}$ or $\text{NULL} <> \text{NULL}$ or $\text{NULL} = \text{NULL}$
- The three-valued logic using the truth value UNKNOWN is:
 - ▶ $(\text{UNKNOWN OR TRUE}) = \text{TRUE}$
 $(\text{UNKNOWN OR FALSE}) = \text{UNKNOWN}$
 $(\text{UNKNOWN OR UNKNOWN}) = \text{UNKNOWN}$
 - ▶ $(\text{TRUE AND UNKNOWN}) = \text{UNKNOWN}$
 $(\text{FALSE AND UNKNOWN}) = \text{FALSE}$
 $(\text{UNKNOWN and UNKNOWN}) = \text{UNKNOWN}$
 - ▶ $(\text{NOT UNKNOWN}) = \text{UNKNOWN}$
- "P IS UNKNOWN" evaluates to TRUE if predicate P evaluates to UNKNOWN
- The result of a WHERE clause predicate is treated as FALSE if it evaluates to UNKNOWN

SQL: Aggregation

- These functions operate on the multiset of values of a column of a relation, and return a value

AVG: average over

MIN: minimum in


MAX: maximum in

SUM: sum of

COUNT: number of

SQL: Aggregation

111



- Find the average salary of instructors in the Computer Science department:

```
SELECT AVG(salary)
FROM   instructor
WHERE  dept_name = 'Comp. Sci.';
```

- Find the total number of instructors who teach a course in the Spring 2010 semester:


```
SELECT COUNT (DISTINCT ID)
FROM   teaches
WHERE  semester = 'Spring'
AND    year = 2010;
```

- Find the number of tuples in the course relation:

```
SELECT COUNT(*)
FROM   course;
```

SQL: Aggregation: Treatment of NULL Values

112




- Find the sum of all instructor salaries:

```
SELECT SUM(salary)
FROM   instructor;
```

- The evaluation ignores any NULL amount.
- The result is NULL if there is no non-NULL amount.
- All aggregate operations except COUNT(*) ignore tuples with NULL values on the aggregated attributes.
- What if the collection only has NULL values?
 - ▶ COUNT returns 0
 - ▶ all other aggregates return NULL

SQL: GROUP BY Aggregation

113




- Find the average salary of instructors in each department:

```
SELECT dept_name, AVG(salary)
FROM   instructor
GROUP BY dept_name;
```

- Note that departments with no instructors do not appear in result.

SQL: GROUP BY Aggregation

114



- Attributes in the SELECT clause that appear outside (i.e., are not arguments) of aggregate functions must appear in the GROUP BY list.
- An **erroneous** query:

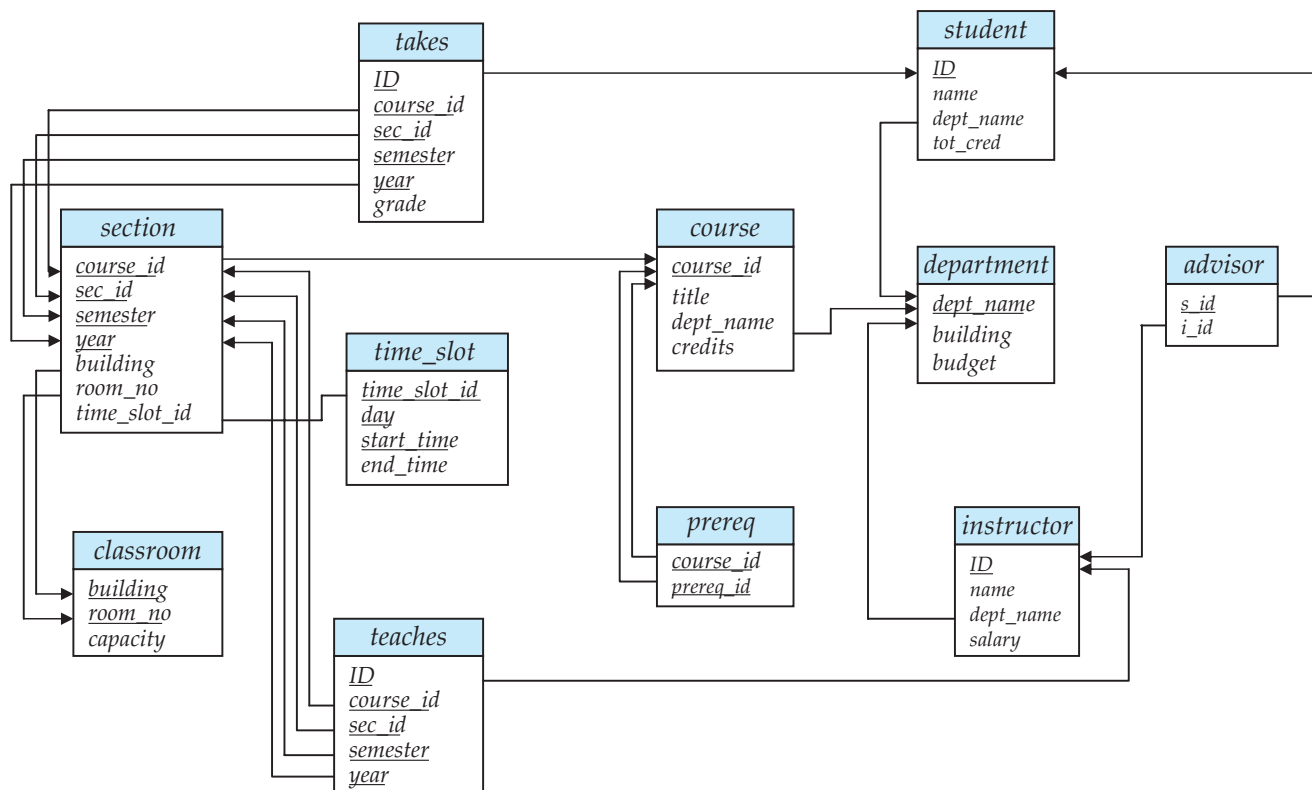
```
SELECT dept_name, name, AVG(salary)
FROM   instructor
GROUP BY dept_name;
```

- A **correct** query (but probably not what you want!):

```
SELECT dept_name, name, AVG(salary)
FROM   instructor
GROUP BY dept_name, name;
```

SQL: GROUP BY Aggregation: The HAVING Clause

115



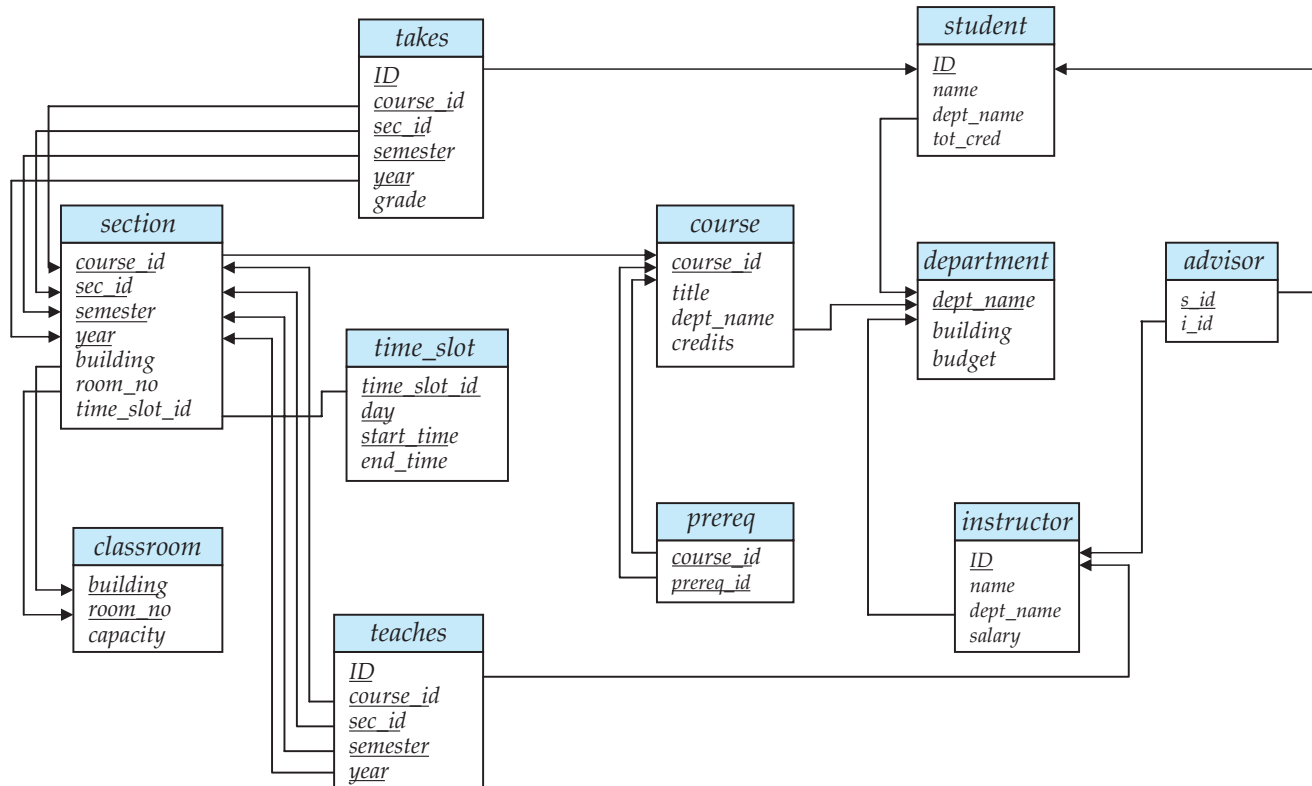
- Find the names and average salaries of all departments whose average salary is greater than 42000:

```
SELECT dept_name, AVG(salary)
FROM   instructor
GROUP BY dept_name
HAVING AVG(salary) > 42000;
```

- Whereas predicates in the WHERE clause are applied before forming groups, predicates in the HAVING clause are applied after the formation of groups.

SQL: Displaying Tuples in a Given Order

116



- List in alphabetic order the names of all instructors:

```
SELECT DISTINCT name
FROM   instructor
ORDER BY name;
```

- We may specify more than one attribute to order by
- For each attribute, we may specify DESC for descending order or ASC (which is the default) for ascending order.

SQL: Nested Subqueries


SQL: Nested Subqueries

118

- SQL queries are compositional, i.e., they allow for the nesting of subqueries.
- A subquery is a SELECT-FROM-WHERE expression that is nested within another query.
- Common uses of subqueries are to perform:
 - ▶ set membership tests
 - ▶ set comparisons
 - ▶ empty set tests

SQL: Nested Subqueries: The IN Clause for Testing Set Membership

119



- Find courses offered in Fall 2009 and in Spring 2010:


```
SELECT DISTINCT course_id
FROM section
WHERE semester = 'Fall'
AND year = 2009
AND course_id IN (SELECT course_id
                  FROM section
                  WHERE semester = 'Spring'
                  AND year = 2010);
```

- Find courses offered in Fall 2009 but not in Spring 2010:

```
SELECT DISTINCT course_id
FROM section
WHERE semester = 'Fall'
AND year = 2009
AND course_id NOT IN (SELECT course_id
                     FROM section
                     WHERE semester = 'Spring'
                     AND year = 2010);
```

SQL: Nested Subqueries: The IN Clause for Testing Set Membership

120



- The example query below can be written in a much simpler manner; this formulation is simply to illustrate SQL features.
- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101:

```
SELECT COUNT (DISTINCT ID)
FROM takes
WHERE (course_id, sec_id, semester, year)
IN (SELECT course_id, sec_id, semester, year
    FROM teaches
    WHERE teaches.ID = 10101);
```


SQL: Nested Subqueries:

Semantics of the SOME Clause

122

$$e \phi \text{ SOME } r \equiv \exists t \in r \text{ s.t. } (e \phi t)$$

where e is an expression that yields a value in the domain of r , t is a tuple in r and $\phi \in \{<, \leq, >, \geq, =, \neq\}$.

- So:
 - ▶ if $e = 5$ and $r = (0, 5, 6)$, then
' $e < \text{SOME } r$ ' evaluates to true
 - ▶ if $e = 5$ and $r = (0, 5)$, then
' $e < \text{SOME } r$ ' evaluates to false
 - ▶ if $e = 5$ and $r = (0, 5)$, then
' $e = \text{SOME } r$ ' evaluates to true
 - ▶ if $e = 5$ and $r = (0, 5)$, then
' $e \neq \text{SOME } r$ ' evaluates to false
- Note that ' $= \text{SOME}$ ' \equiv ' IN '
but ' $\neq \text{SOME}$ ' $\not\equiv$ ' NOT IN '

SQL: Nested Subqueries: Semantics of the ALL Clause

124

$$e \phi \text{ ALL } r \equiv \exists t \in r \text{ s.t. } (e \phi t)$$

where e is an expression that yields a value in the domain of r , t is a tuple in r and $\phi \in \{<, \leq, >, \geq, =, \neq\}$.

- So:
 - ▶ if $e = 5$ and $r = (0, 5, 6)$, then ' $e < \text{ALL } r$ ' evaluates to false
 - ▶ if $e = 5$ and $r = (6, 10)$, then ' $e < \text{ALL } r$ ' evaluates to true
 - ▶ if $e = 5$ and $r = (4, 5)$, then ' $e = \text{ALL } r$ ' evaluates to false
 - ▶ if $e = 5$ and $r = (4, 6)$, then ' $e \neq \text{ALL } r$ ' evaluates to true
- Note that ' $= \text{ALL}$ ' \neq ' IN '
but ' $\neq \text{ALL}$ ' \equiv ' NOT IN '

SQL: Nested Subqueries: Testing for the Empty Set Using the EXISTS Clause


125

- The EXISTS clause is used to return true if the argument subquery is nonempty and false otherwise:

$\text{EXISTS } r \equiv r \neq \emptyset$

$\text{NOT EXISTS } r \equiv r = \emptyset$

SQL: Nested Subqueries: Testing for the Empty Set Using the EXISTS Clause




- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”:

```
SELECT s1.course_id
FROM   section s1
WHERE  s1.semester = 'Fall' AND s1.year = 2009
AND    EXISTS (SELECT *
               FROM   section s2
               WHERE  s2.semester = 'Spring' AND year = 2010
               AND    s1.course_id = s2.course_id);
```

- This is an example of a **correlated subquery**, i.e., a subquery in which a range variable from the outer query (i.e., s1) is used in the inner query (it is compared with s2).

SQL: Nested Subqueries: Testing for the Empty Set Using the EXISTS Clause

127




- Find all students who have taken all courses offered in the Biology department.

```
SELECT DISTINCT s.ID, s.name
FROM   student s
WHERE
NOT EXISTS ((SELECT course_id
              FROM   course
              WHERE  dept_name = 'Biology')
            EXCEPT (SELECT t.course_id
                     FROM   takes t
                     WHERE  s.ID = t.ID));
```

- Note that $X \setminus Y = \emptyset \equiv X \subseteq Y$

SQL: Nested Subqueries: Testing for the Absence of Duplicates Using the UNIQUE Clause

128




- The UNIQUE clause is used to return true if the argument subquery does not return duplicate tuples and false otherwise.
- It returns true on the empty set.
- Find all courses that were offered at most once in 2009:

```
SELECT c.course_id
FROM   course c
WHERE  UNIQUE (SELECT s.course_id
               FROM   section s
               WHERE  c.ID = s.ID AND s.year = 2009);
```


SQL: Nested Subqueries: Subqueries in the FROM Clause

129




- SQL allows a subquery expression to be used in the FROM clause.
- Find the average instructor salary in those departments where the average salary is greater than \$42,000:

```
SELECT dept_name, avg_salary
FROM   (SELECT dept_name, AVG(salary) AS avg_salary
        FROM   instructor
        GROUP BY dept_name)
WHERE  avg_salary > 42000;
```

- Note that this acts as a HAVING clause.
- Note also that avg_salary in the result is the renaming of the attribute resulting from computing AVG(salary) in the inner query.

SQL: Nested Subqueries: Subqueries in the FROM Clause

130



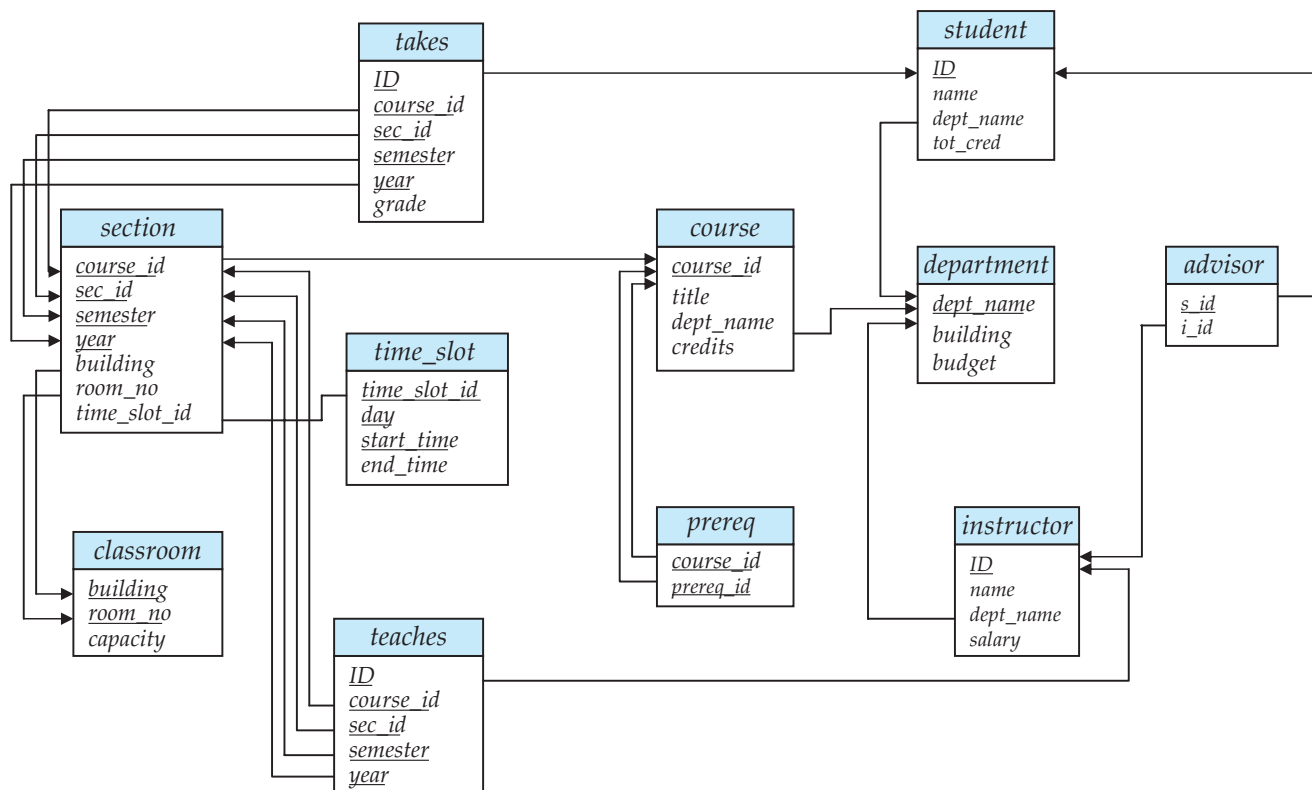
- Another way to write the previous query:

```
SELECT dept_name, avg_salary
FROM (SELECT dept_name, AVG(salary)
      FROM instructor
      GROUP BY dept_name)
      AS (dept_name, avg_salary)
WHERE avg_salary > 42000;
```

- In the second case, the renaming of the attributes in the nested subquery is done using a tuple of attribute names.

SQL: Nested Subqueries: Temporary Views Using the WITH Clause

131



- The WITH clause provides a way of defining a temporary view whose definition is available only to the query in which the WITH clause occurs.


- Find all departments with the maximum budget:

```
WITH max_budget(value)
AS (SELECT max(budget)
     FROM department)
SELECT budget
FROM department, max_budget
WHERE department.budget = max_budget.value;
```

- Note that the WITH clause specifies a schema for the temporary relation, so that its name (i.e., max_budget) can appear in the FROM clause of the main query and its attribute (i.e., value) can be referred to in the WHERE clause of the main query.

SQL: Nested Subqueries: Temporary Views Using the WITH Clause

132



- The WITH clause is very useful for writing complex queries
- Supported by most database systems, with minor syntactic variations
- We can use two temporary views to find all departments where the total salary is greater than the average of the total salary at all departments:

```
WITH
    dept_total(dept_name, value)
AS (SELECT dept_name, SUM(salary)
     FROM instructor
     GROUP BY dept_name),
    dept_total_avg(value)
AS (SELECT AVG(value)
     FROM dept_total)
SELECT dept_name
FROM dept_total, dept_total_avg
WHERE dept_total.value >= dept_total_avg.value;
```

SQL: Nested Subqueries: Scalar Subqueries

133

- A **scalar subquery** is one that returns a single value and therefore can be used wherever a single value is expected.
- One example, in the SELECT clause:

```
SELECT dept_name,  
       (SELECT COUNT(*)  
        FROM instructor  
        WHERE department.dept_name = instructor.dept_name) AS num_instructors  
FROM   department;
```

- Another example, in the WHERE clause:

```
SELECT name  
FROM   instructor  
WHERE  salary*10 > (SELECT budget  
                   FROM   department  
                   WHERE  department.dept_name = instructor.dept_name);
```

- A runtime error is raised if the subquery returns more than one result tuple.

Summary

The Relational Languages

135

- An algebra is a mathematical structure made of a set of values/variables and a set of operations on those values.
- A relational algebra (RA) has relations as values and operations on relations.
- The basic operations of RA are selection, projection, product, renaming, as well as union, intersection and difference.
- The most useful derived operations are various forms of join (e.g., inner, outer, semi- and antijoins).
- The most useful operations in an extended RA are extended projection, duplicate removal, sorting, grouping and aggregation.
- SQL has a direct translation to (extended) RA, which allows optimizers to manipulate RA expressions using equivalence laws.
- SQL is a powerful, practical language for interacting with relational DBMSs.

In The Next Handout

136

- We'll begin our exploration of conceptual modelling.
- We'll see that conceptual modelling is aimed at providing an intuitive approach to identifying information needs.
- We'll observe that a good conceptual model facilitates the generation of well-formed logical models.
- In particular, we'll learn about the formalism for conceptual modelling known as Entity-Relationship (ER) Modelling.