

H6

Designing Databases: Mapping a Conceptual into a Logical Design

Fundamentals of Databases

Alvaro A A Fernandes, SCS, UoM

[COMP23111 Handout 06 of 12]

Acknowledgements

- These slides are adaptations (mostly minor, but some major) of material authored and made available to instructors by **Ramez Elmasri and Shamkant B. Navathe** to accompany their textbook **Database Systems: Models, Languages, Design, and Application Programming, 6th (Global) Edition, Addison-Wesley Pearson, 2011, 978-0-13-214498-8**
- Copyright remains with them and the publishers, whom I thank.
- Some slides had input from Sandra Sampaio, whom I thank.
- All errors are my responsibility.

In Previous Handouts

- We learned how data models lead to a distinction between schemas and instances that enables a logical view of the data.
- We learned about the relational approach to logical data modelling.
- We learned about the relational algebra and SQL, both its DDL and DML capabilities and its querying constructs.
- We learned of the practical benefits of performing conceptual modelling before logical design and why the EER approach serves well this purpose.

In This Handout

- We'll learn how a conceptual model can be mapped into a logical model that is capable of being implemented in a relational DBMS.

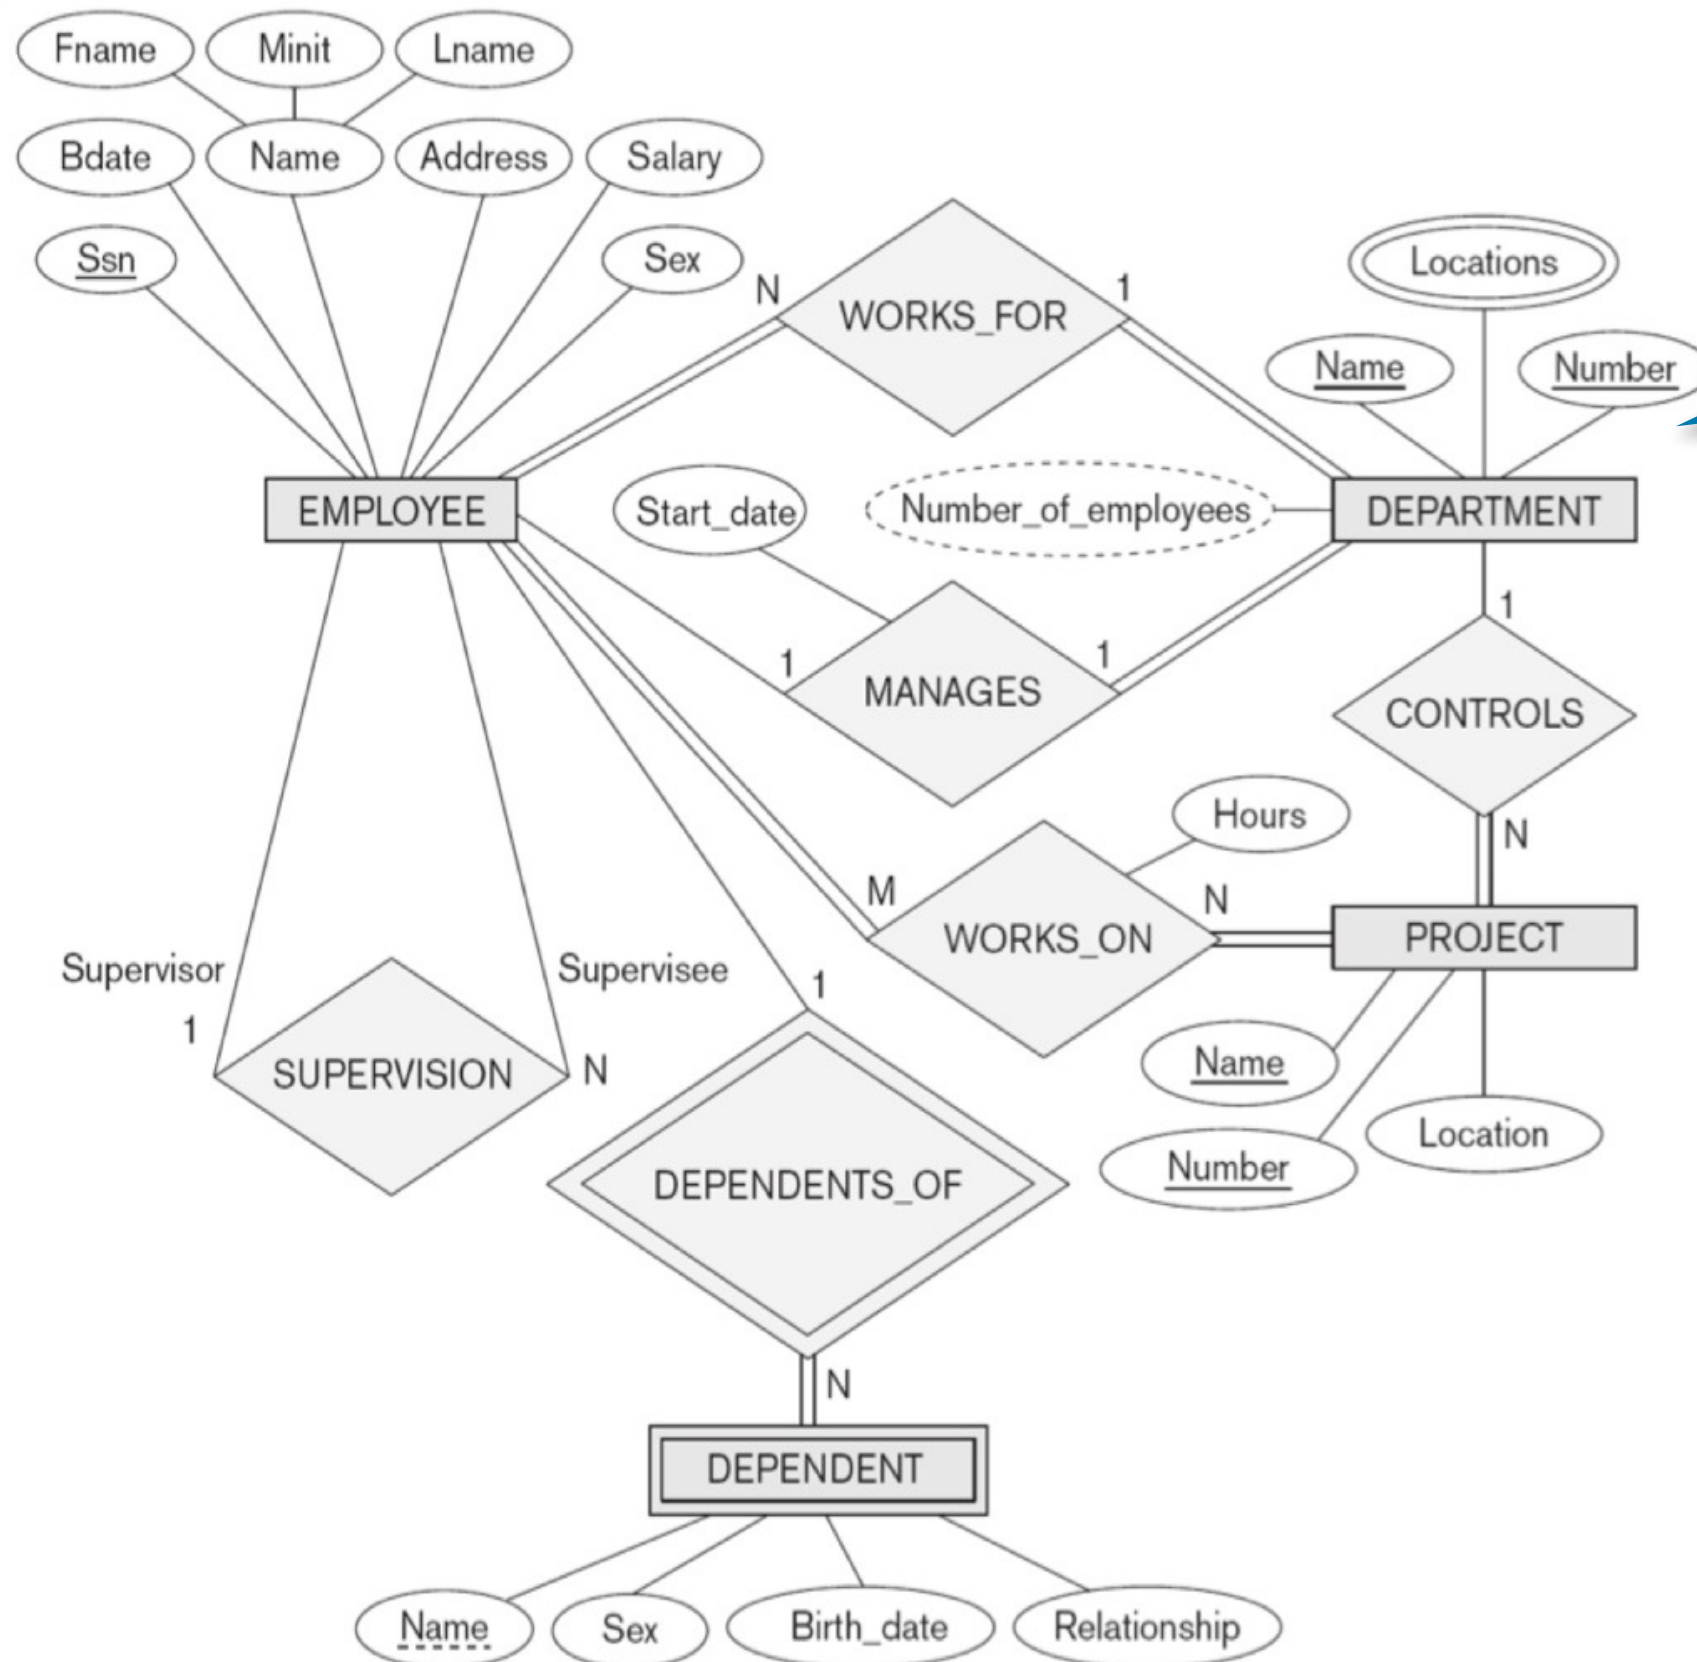
Relational Database Design through (E)ER-to-Relational Mapping

5

- The goal is to derive a relational schema (i.e., logical, implementable) from an (E)ER conceptual schema
- As we have seen, a conceptual schema (because it is more abstract and expressive) is better suited to serve as a bridge between the DB designer and the various kinds of DB users when the latter are expressing interactively and iteratively their data requirements to the former.
- There is a seven-step semi-formal procedure to convert an ER model into a relational schema.
- Then, two further steps suffice for converting an EER model into a relational schema.

An Example ER Diagram for a COMPANY DB

6



A
model like this one is the
starting point.

An Example Relational Schema for a COMPANY DB

7

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|------------------|
| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|

A schema like this one is the end point.

Step 1: Mapping of Regular Entity Types

8

- For each regular (i.e., not weak) entity type E, create a relation R that includes all the simple attributes of E.
- Then include the leaf-level component attributes of a composite attribute.
- Then choose one of the keys in E as the primary key (PK) of R.
- If the chosen key is composite in E, then make the set of component attributes in it the composite PK of R.
- Then make the other keys in E secondary keys in R (i.e., assert them as unique in the DDL).
- Such relations are called **entity relations** and *each tuple represents an entity instance*.

After Step 1: Entity Relations

9

EMPLOYEE

| | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary |
|-------|-------|-------|------------|-------|---------|-----|--------|

DEPARTMENT

| | |
|-------|----------------|
| Dname | <u>Dnumber</u> |
|-------|----------------|

PROJECT

| | | |
|-------|----------------|-----------|
| Pname | <u>Pnumber</u> | Plocation |
|-------|----------------|-----------|

Step 2: Mapping of Weak Entity Types

10

- For each weak entity type W with owner entity type E , create a relation R that includes all the simple attributes (or components of a composite attribute) of W as attributes of R .
- Then include the PK attributes of E as foreign key (FK) attributes of R .
- Then make the PK of R the composition of the PK of E and the weak/partial key of W .
- If the owner of W is another weak entity type W' with owner E , then map W' first and map W subsequently.

After Steps 1-2

11

EMPLOYEE

| | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary |
|-------|-------|-------|------------|-------|---------|-----|--------|

DEPARTMENT

| | |
|-------|----------------|
| Dname | <u>Dnumber</u> |
|-------|----------------|

PROJECT

| | | |
|-------|----------------|-----------|
| Pname | <u>Pnumber</u> | Plocation |
|-------|----------------|-----------|

DEPENDENT

| | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|



Step 3: Mapping of Binary 1:1 Relationship Types

12

- For each binary 1:1 relationship type **R**, identify the relations S and T that correspond to the entity types participating in it.
- There are three possible approaches:
 - ▶ Foreign key approach (FKA)
 - ▶ Merged relationship approach (MRA)
 - ▶ Cross-reference or relationship relation approach (RRA)

Step 3: Mapping of Binary 1:1 Relationship Types

13

- The foreign key approach is very useful, in general.
- In **FKA**:
 - ▶ Choose one of the relations, say, S, and include the PK of T (the parent, or exporter, as it were) as FK in S (the child, or importer, as it were).
 - ▶ Then include all the simple (or component of composite) attributes of the relationship type as attributes in S.
 - ▶ It is better to choose an entity type with total participation in **R** for the role of S (i.e., as the PK importer).

Step 3: Mapping of Binary 1:1 Relationship Types

14

- In **MRA**:
 - ▶ Merge the two entity types and the relationship into a single relation.
 - ▶ This is possible when both participations are total, since the two relations will have the same number of tuples at all times.

Step 3: Mapping of Binary 1:1 Relationship Types

15

- In **RRA**:
 - ▶ Create a new, distinct relation R' for capturing the cross-referencing between S and T .
 - ▶ Then include as FKs of R' the PK of S and the PK of T .
 - ▶ Then choose one of them to be the PK of R' and assert the other to be unique (i.e., a secondary key).
 - ▶ This relation is then called a **relationship relation** (or a lookup table).

Step 4: Mapping of Binary 1:N Relationship Types

16

- For each regular binary 1:N relationship type **R** between entity types S and T, use FKA:
 - ▶ Identify the relation, say S, that stands at the N-side of R (i.e., is the child, or importer)
 - ▶ Then include the PK of T (i.e., the parent, or exporter) as FK in S
 - ▶ Then include the simple (or component of composite) attributes of **R** as attributes of S.
- Alternatively, use RRA.

Step 5: Mapping of Binary M:N Relationship Types

17

- For each binary M:N relationship type **R** between entity types S and T, use RRA:
 - ▶ Create a new relation R'.
 - ▶ Then include as FKs of R', the PKs of S and T.
 - ▶ Then make the PK of R' the composition of the PKs of S and T.
 - ▶ Then include any simple (or components of composite) attributes of **R** in R'.

After Steps 1-5

18

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

PROJECT

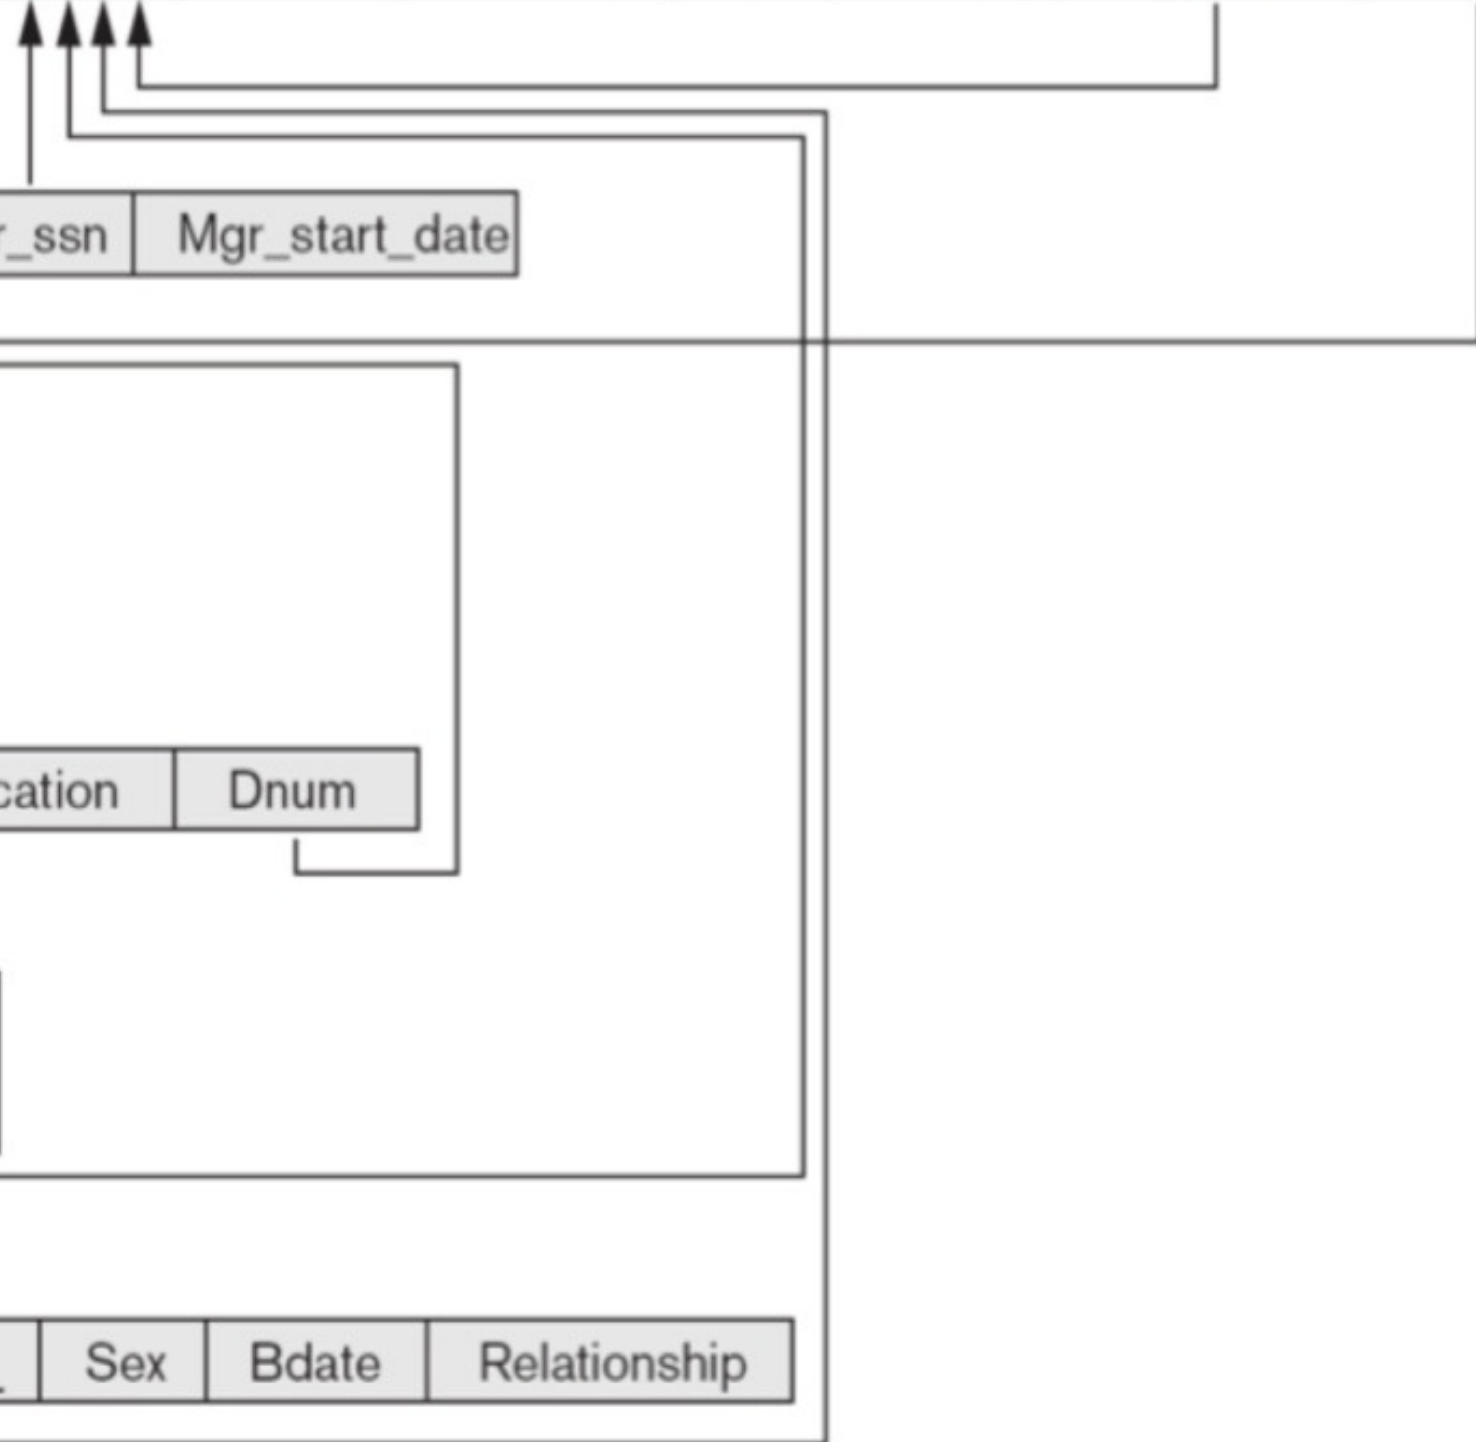
| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|



Step 6: Mapping of Multivalued Attributes

19

- For each multivalued attribute A in an entity type S
 - ▶ Create a new relation A' with the attribute A .
 - ▶ Then include the PK of S as an FK of A' .
 - ▶ Then make the PK of A' the composition of A and the PK of S .
 - ▶ Then, if the multivalued attribute is composite, include its components.

After Steps 1-6

20

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|------------------|
| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|

PROJECT

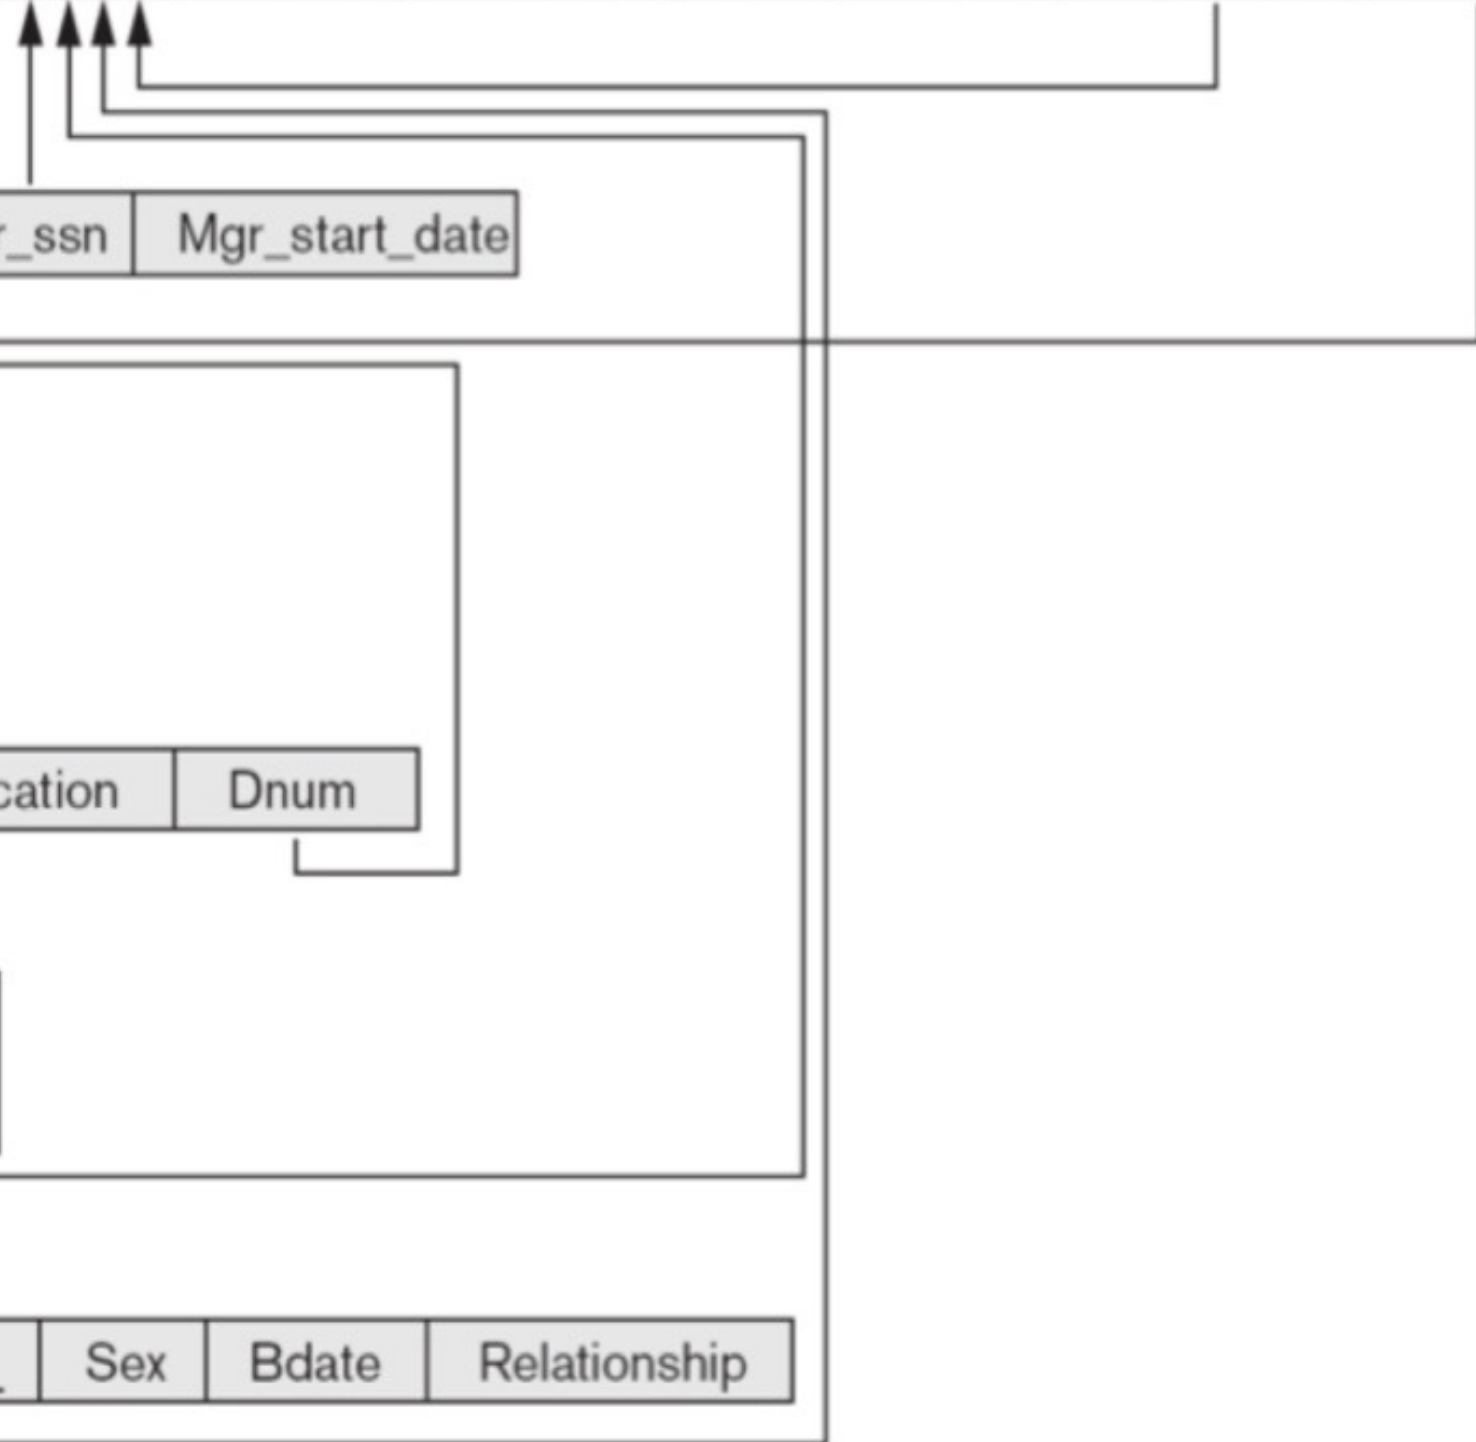
| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|



Step 7: Mapping of N-ary Relationship Types

21

- For each n-ary, $n > 2$, relationship type **R**:
 - ▶ Create a new relation R' to represent **R**.
 - ▶ Then include the PKs of participating entity types as FKs in R' .
 - ▶ Then make the PK of R' the concatenation of the PKs of participating entity types.
 - ▶ But, if any participating entity type X does so with a cardinality constraint of 1, then do not include the PK of X in the PK of R' .
 - ▶ Then include any simple (or components of composite) attributes of **R** in R' .

Correspondence between ER and Relational

22

ER MODEL

Entity type

1:1 or 1:N relationship type

M:N relationship type

n -ary relationship type

Simple attribute

Composite attribute

Multivalued attribute

Value set

Key attribute

RELATIONAL MODEL

Entity relation

Foreign key (or *relationship* relation)

Relationship relation and *two* foreign keys

Relationship relation and n foreign keys

Attribute

Set of simple component attributes

Relation and foreign key

Domain

Primary (or secondary) key

Major Contrasts

23

- In a relational schema, relationship types are not represented explicitly.
- Instead, they are represented by having two attributes A and B, where one is a primary key and the other a foreign key.
- Composite and multivalued attributes are decomposed in the relational schema .
- Derived attributes need triggers and procedures (studied later in the course) to be supported.
- Specialization and generalization are only indirectly supported (as we shall now explore).

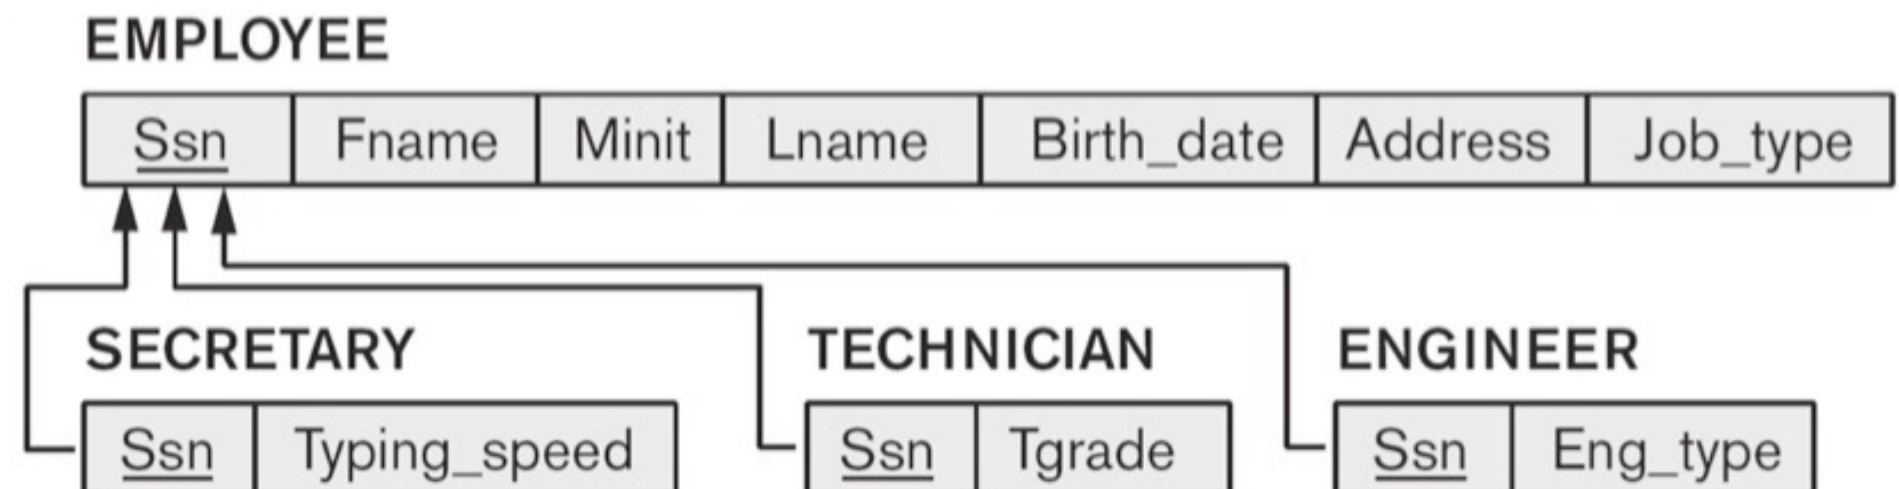
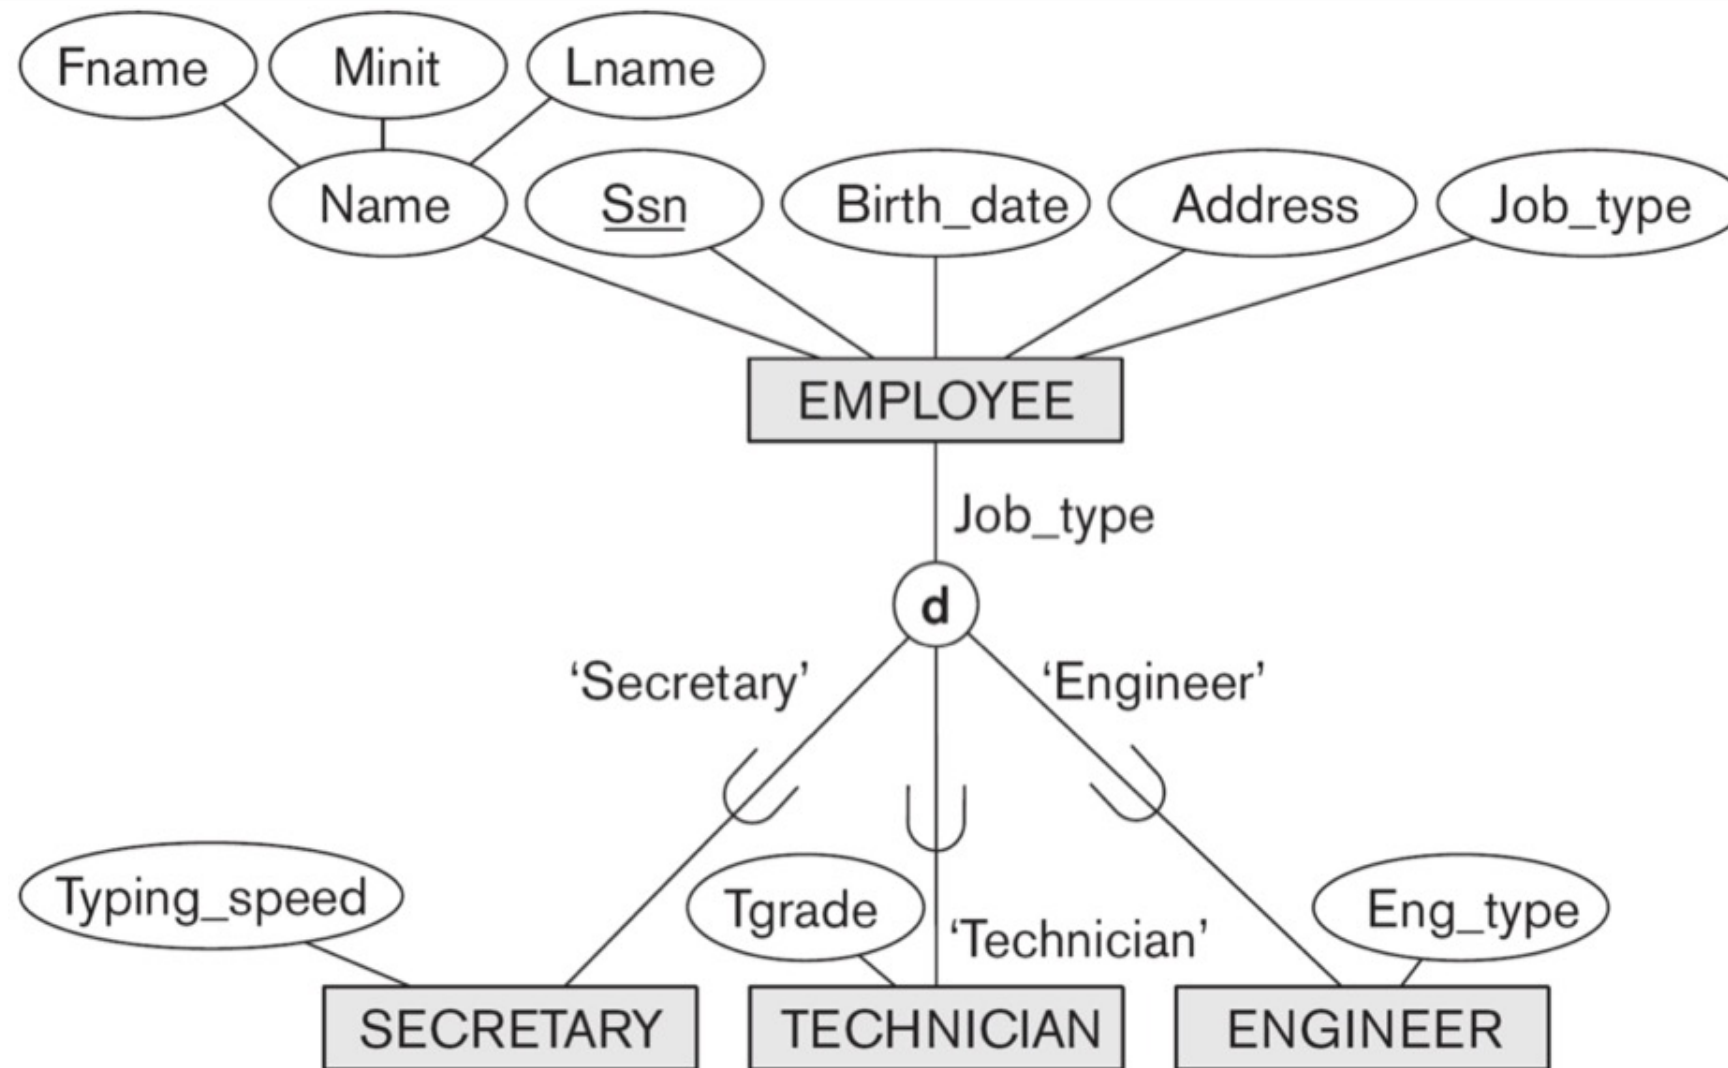
Step 8: Options for Mapping Specialization or Generalization

24

- Option 8A: Using *multiple relations* (superclass and subclasses) works for any specialization (total or partial, disjoint or overlapping).

Step 8 Example: Option 8A

25



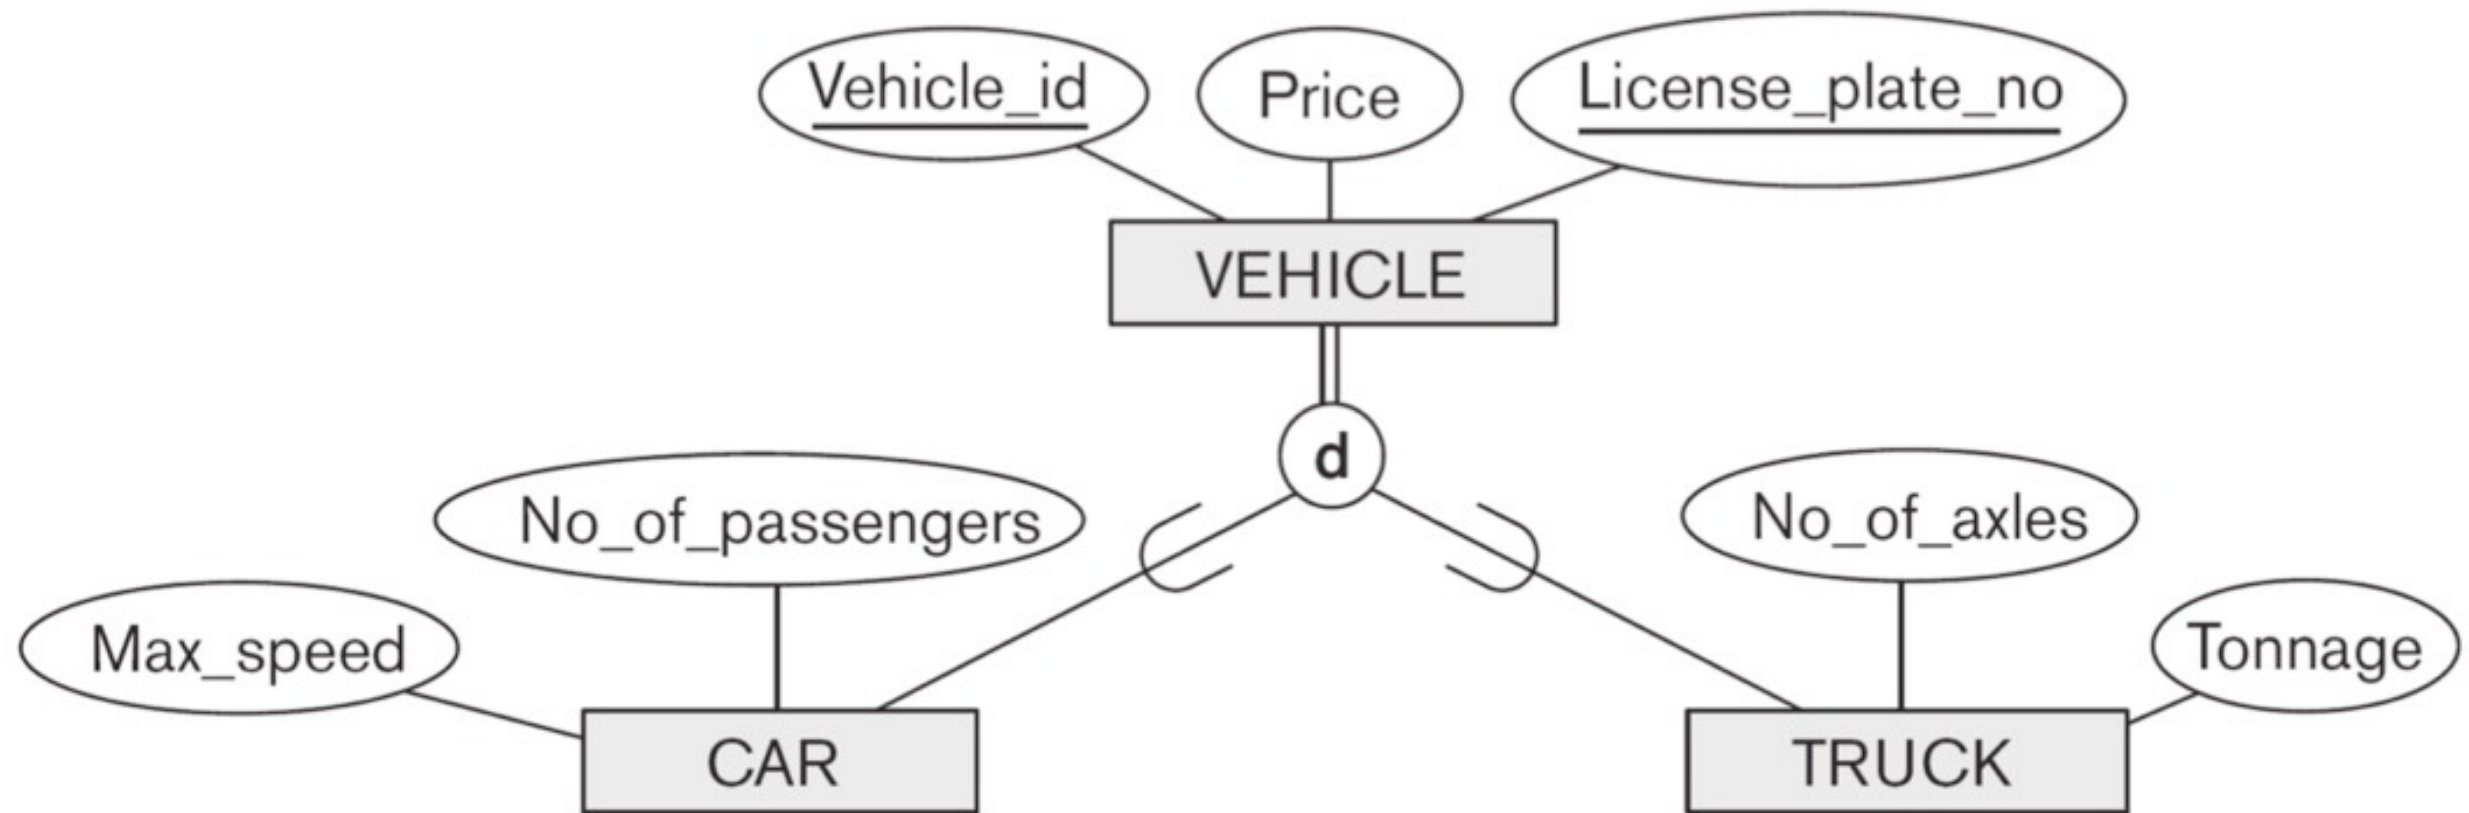
Step 8: Options for Mapping Specialization or Generalization

26

- Option 8B: Using *multiple relations on subclass relations* only works if subclasses are total.
- It is only recommended if the specialization has disjointedness constraint.
- Otherwise, the same entity may be duplicated in several relations.

Step 8 Example: Option 8B

27



CAR

| | | | | |
|-------------------|------------------|-------|-----------|------------------|
| <u>Vehicle_id</u> | License_plate_no | Price | Max_speed | No_of_passengers |
|-------------------|------------------|-------|-----------|------------------|

TRUCK

| | | | | |
|-------------------|------------------|-------|-------------|---------|
| <u>Vehicle_id</u> | License_plate_no | Price | No_of_axles | Tonnage |
|-------------------|------------------|-------|-------------|---------|

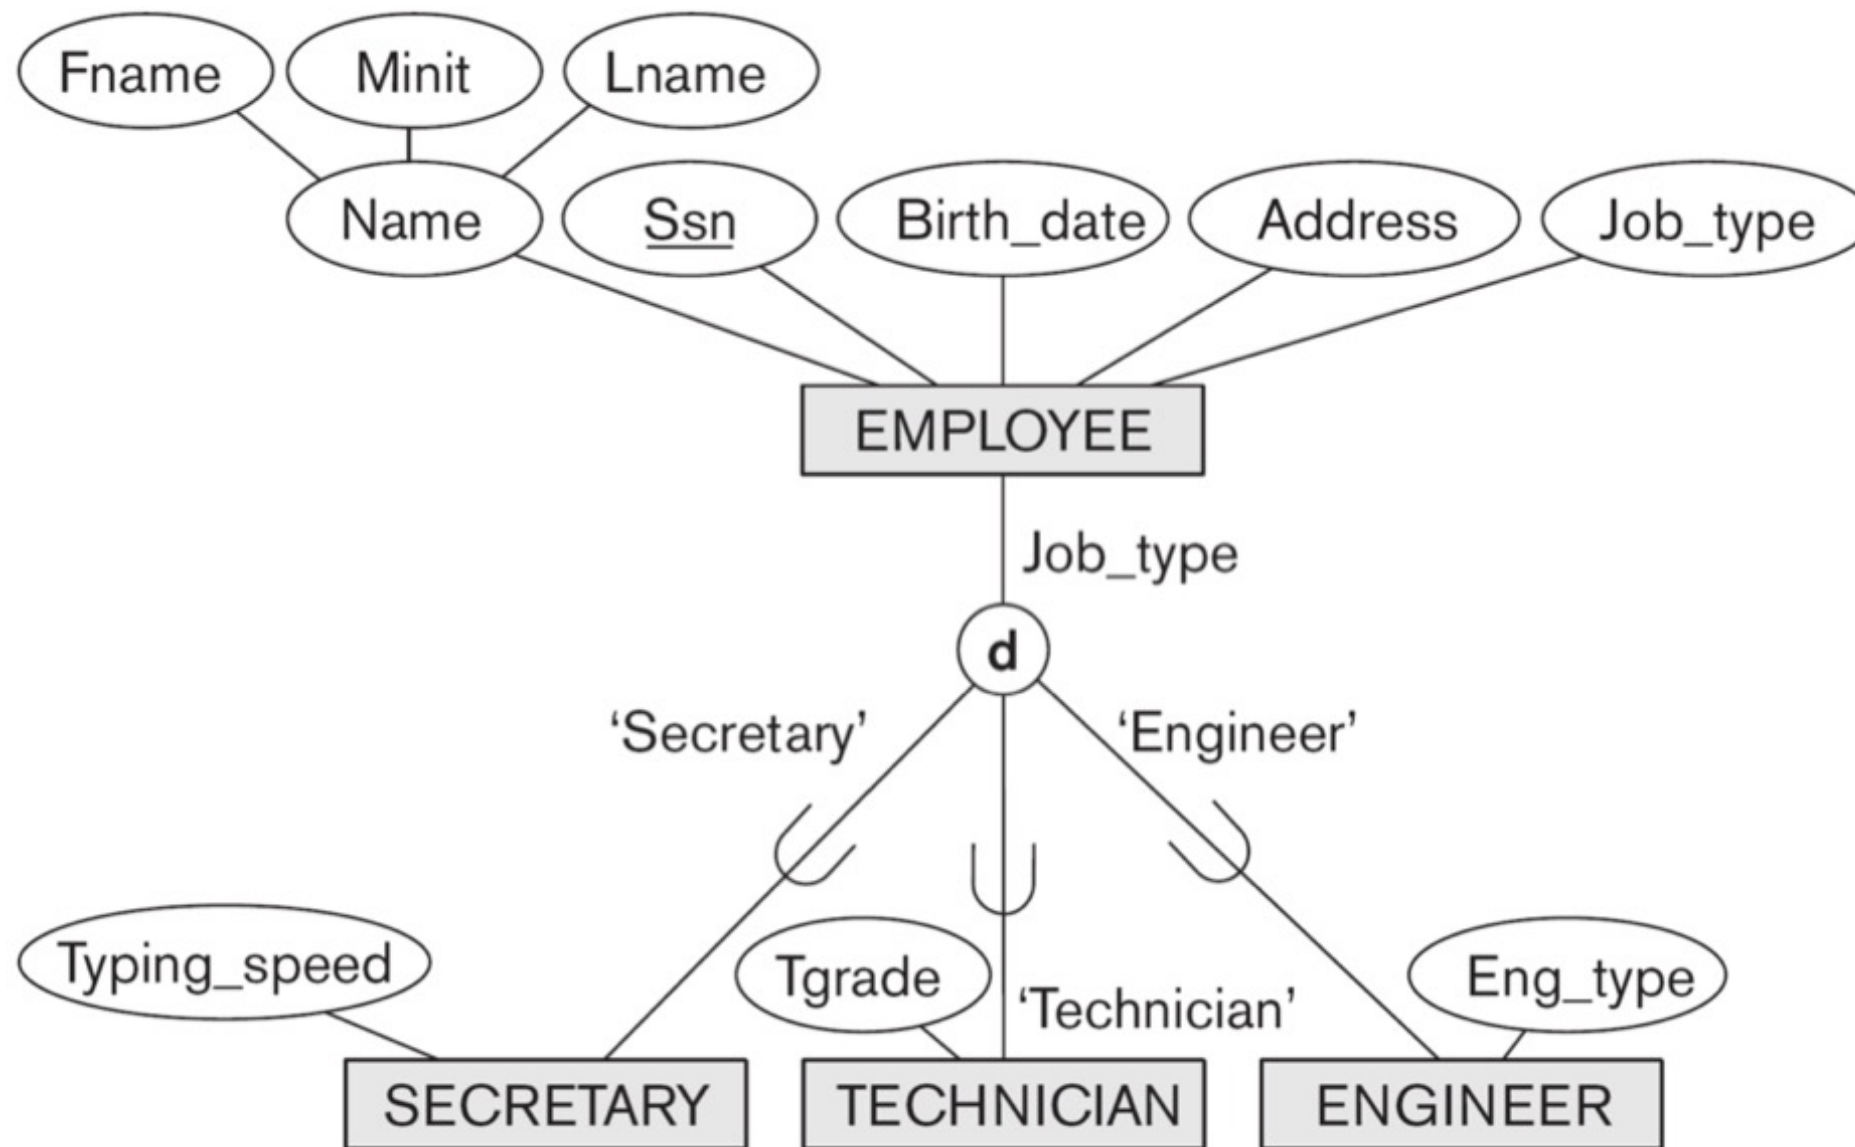
Step 8: Options for Mapping Specialization or Generalization

28

- Option 8C: Using a *single relation with one so-called type (or discriminating) attribute* whose value in a tuple indicates which subclass the tuple belongs to.
- This works only if subclasses are disjoint and still has the potential for generating many NULL values if many specific attributes exist in the subclasses.

Step 8 Example: Option 8C

29



EMPLOYEE

| <u>Ssn</u> | Fname | Minit | Lname | Birth_date | Address | Job_type | Typing_speed | Tgrade | Eng_type |
|------------|-------|-------|-------|------------|---------|----------|--------------|--------|----------|
|------------|-------|-------|-------|------------|---------|----------|--------------|--------|----------|

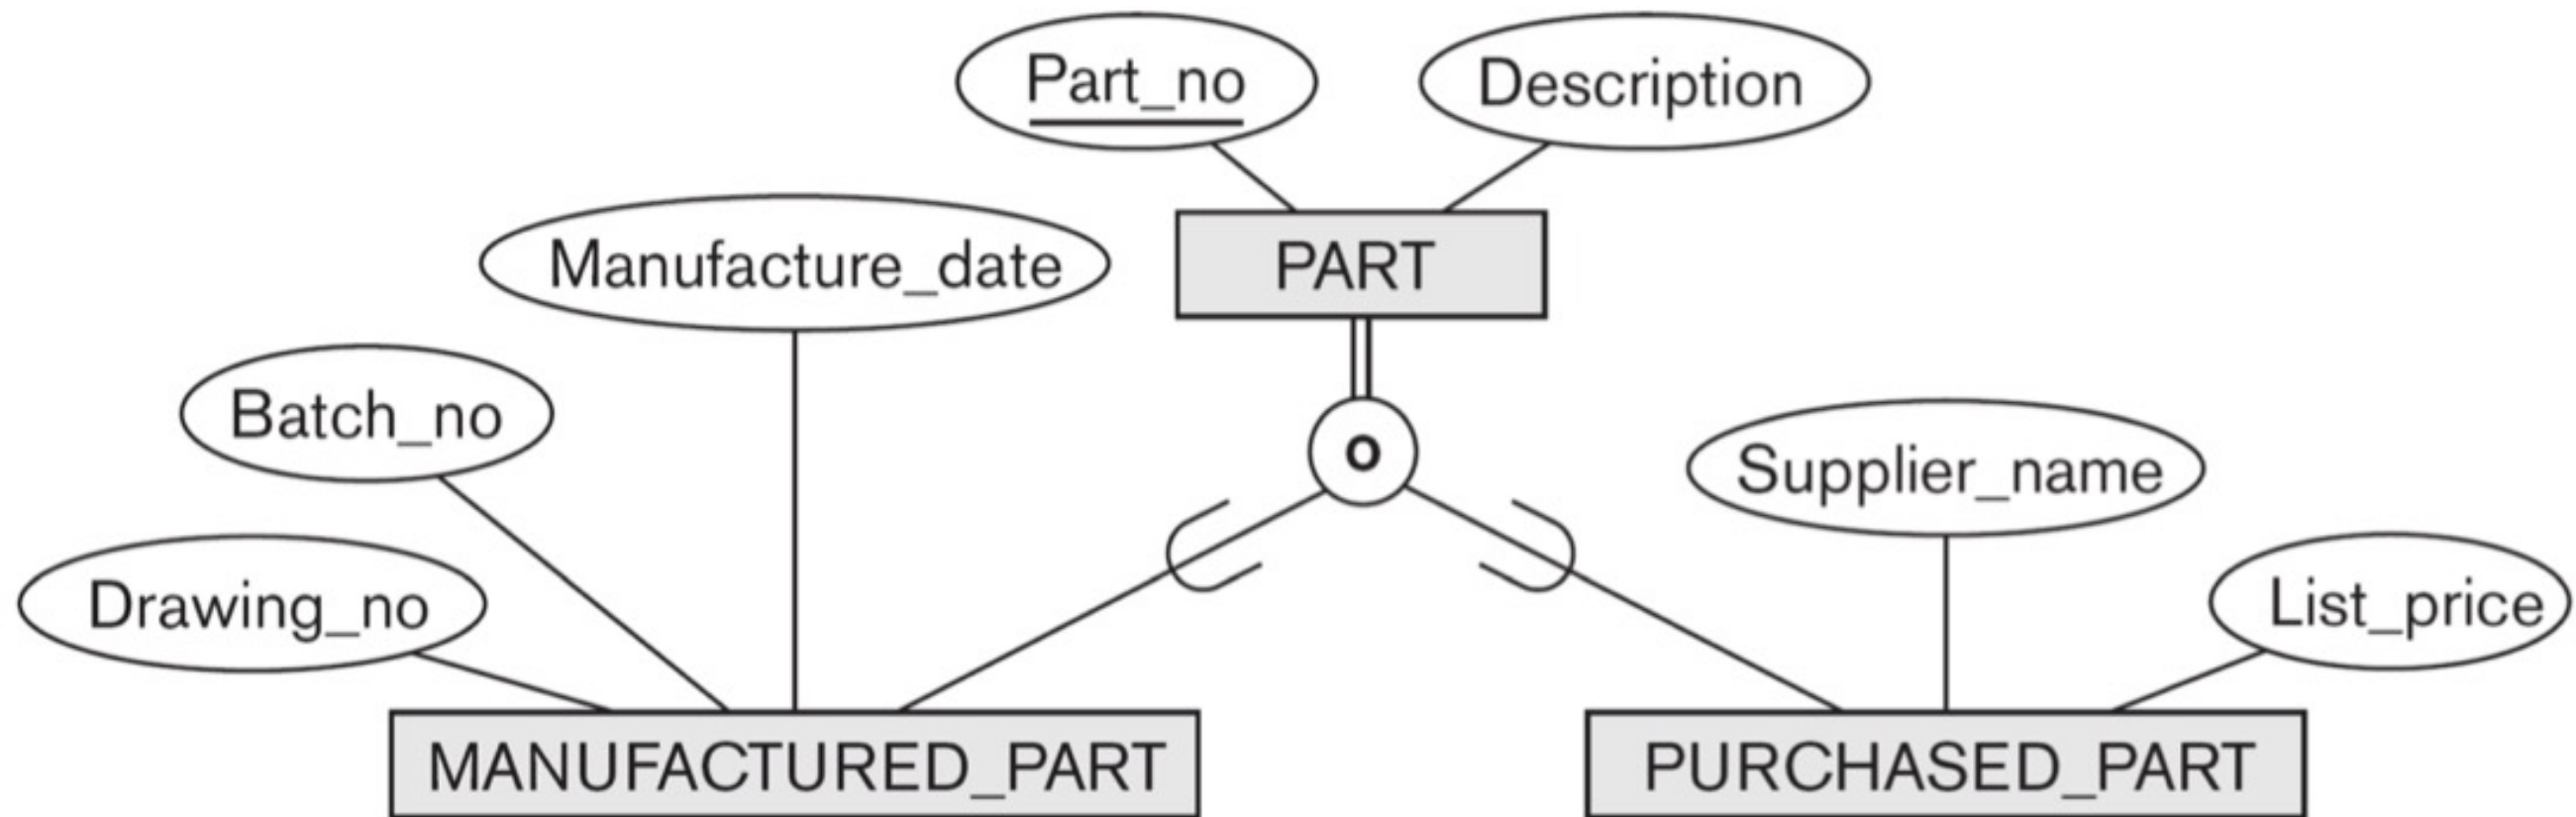
Step 8: Options for Mapping Specialization or Generalization

30

- Option 8D: Using a *single relation with multiple Boolean type attributes*, each of which indicates whether or not the tuple belongs to the corresponding subclass.
- This works when subclasses are overlapping and will also work for a disjoint specialization.

Step 8 Example: Option 8D

31

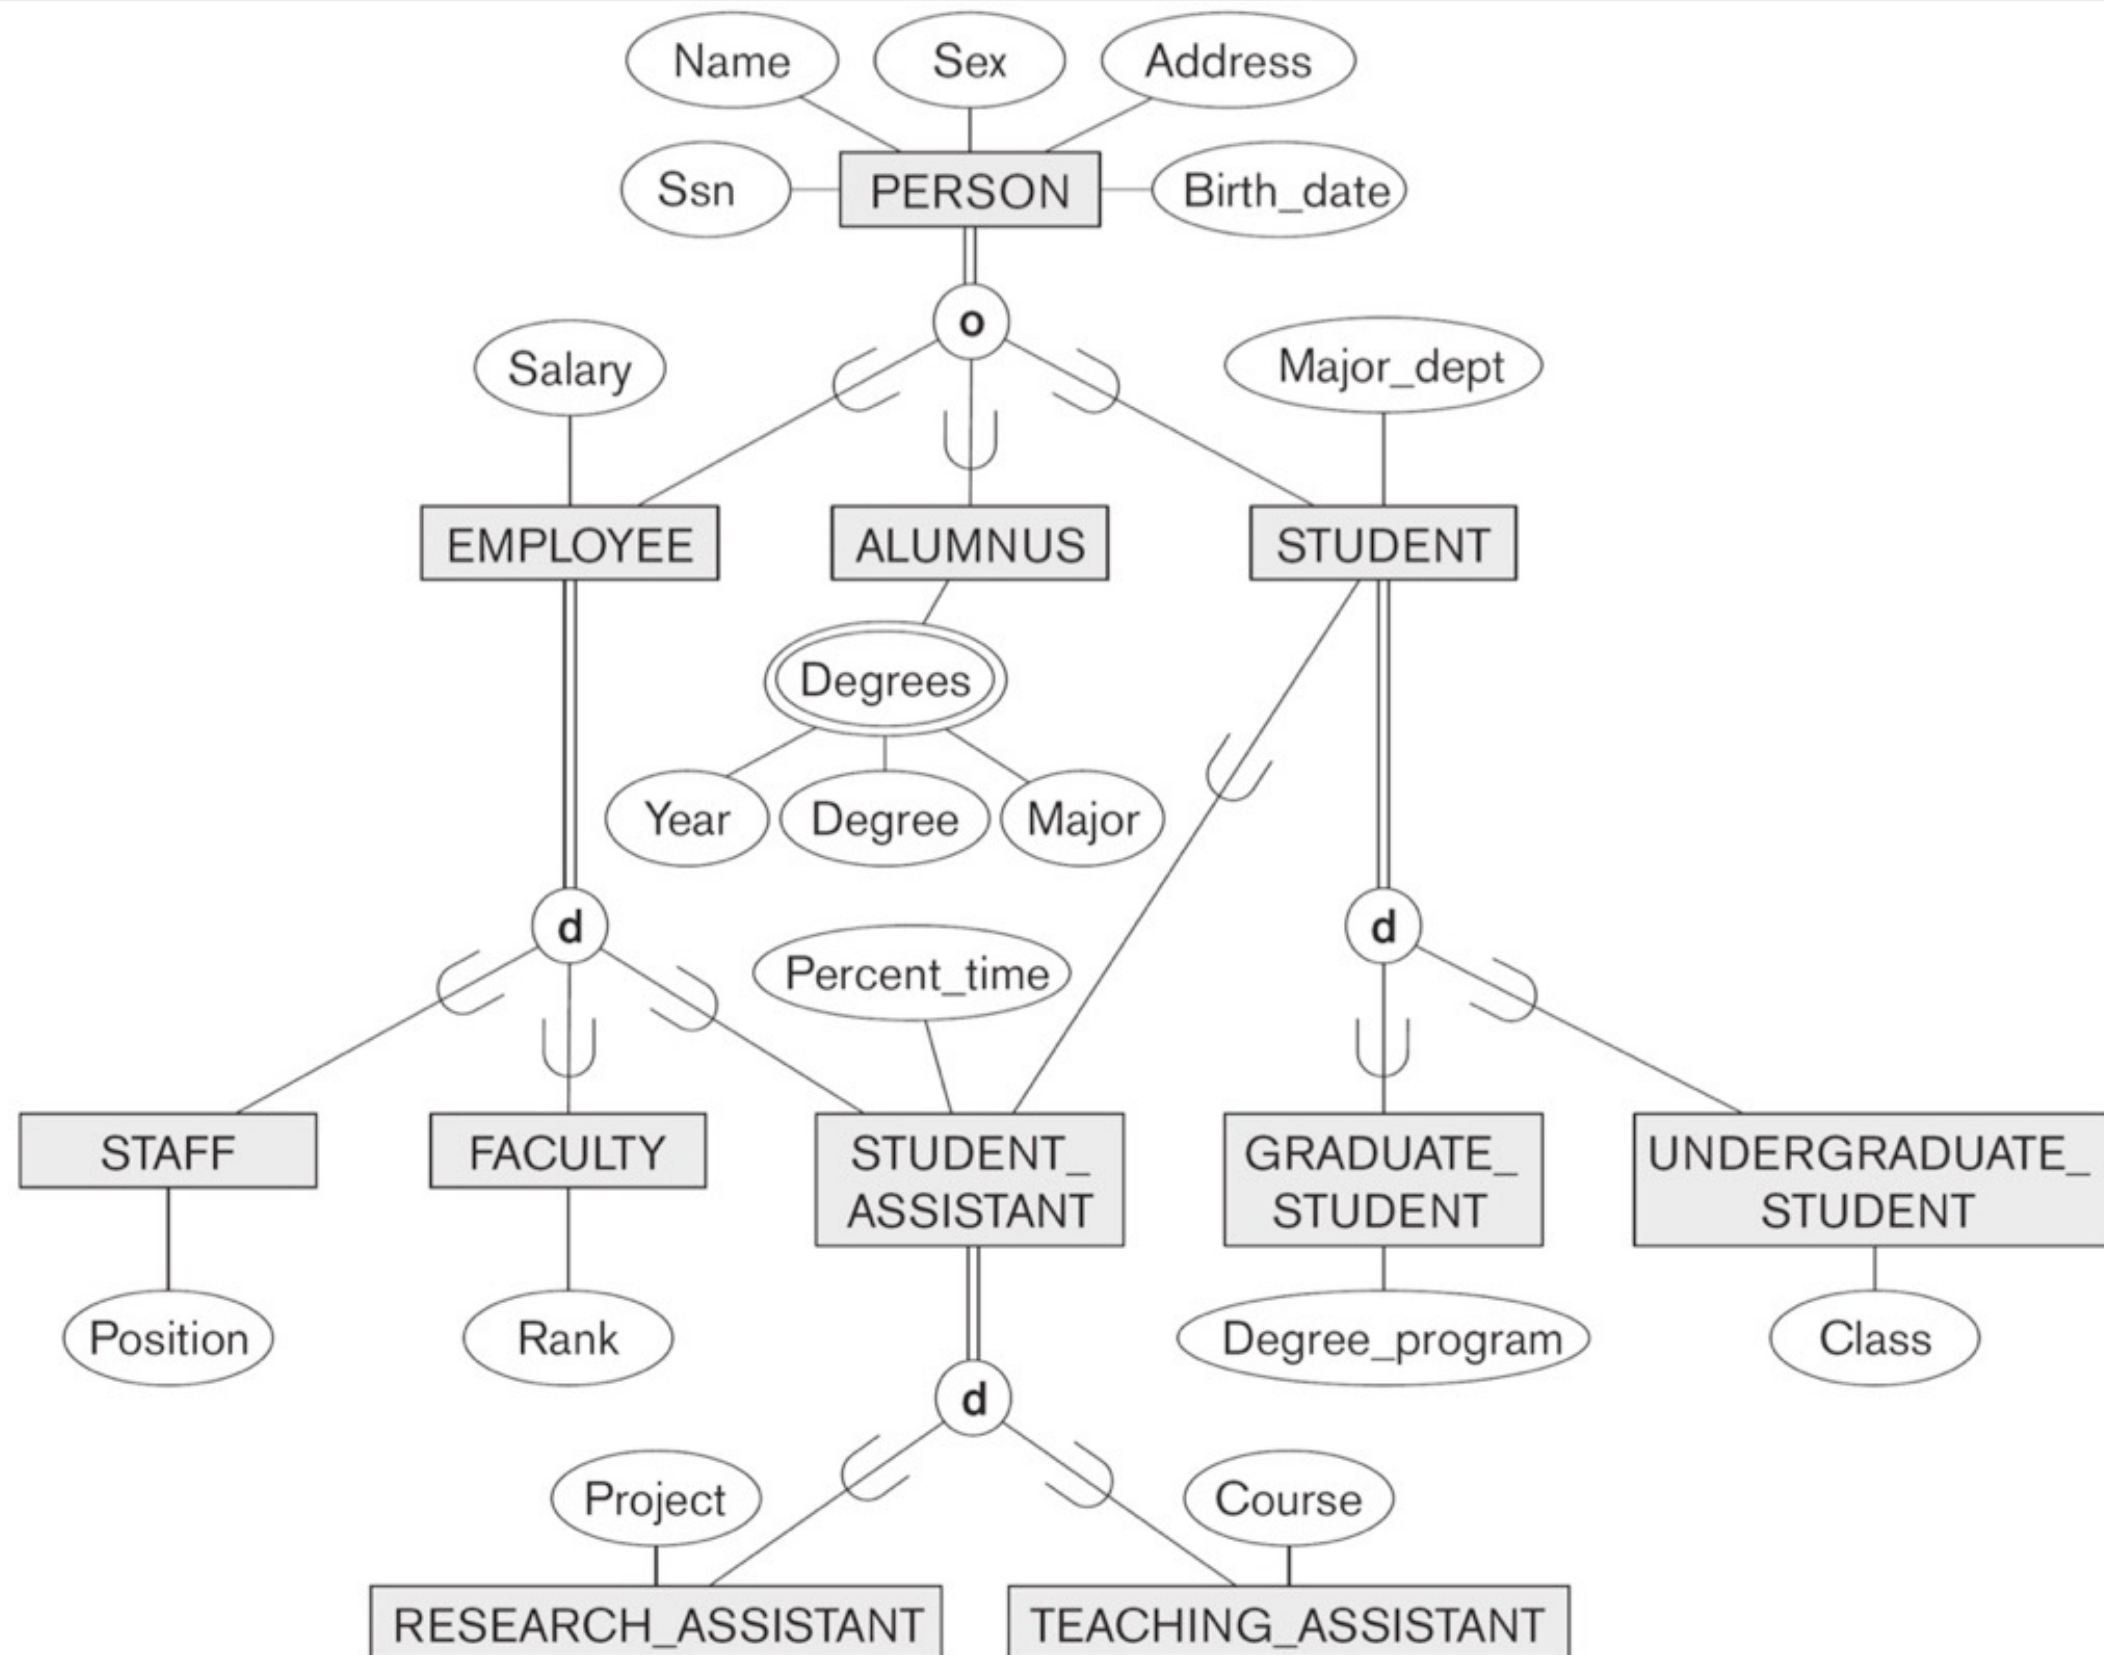


PART

| | | | | | | | | |
|----------------|-------------|-------|------------|------------------|----------|-------|---------------|------------|
| <u>Part_no</u> | Description | Mflag | Drawing_no | Manufacture_date | Batch_no | Pflag | Supplier_name | List_price |
|----------------|-------------|-------|------------|------------------|----------|-------|---------------|------------|

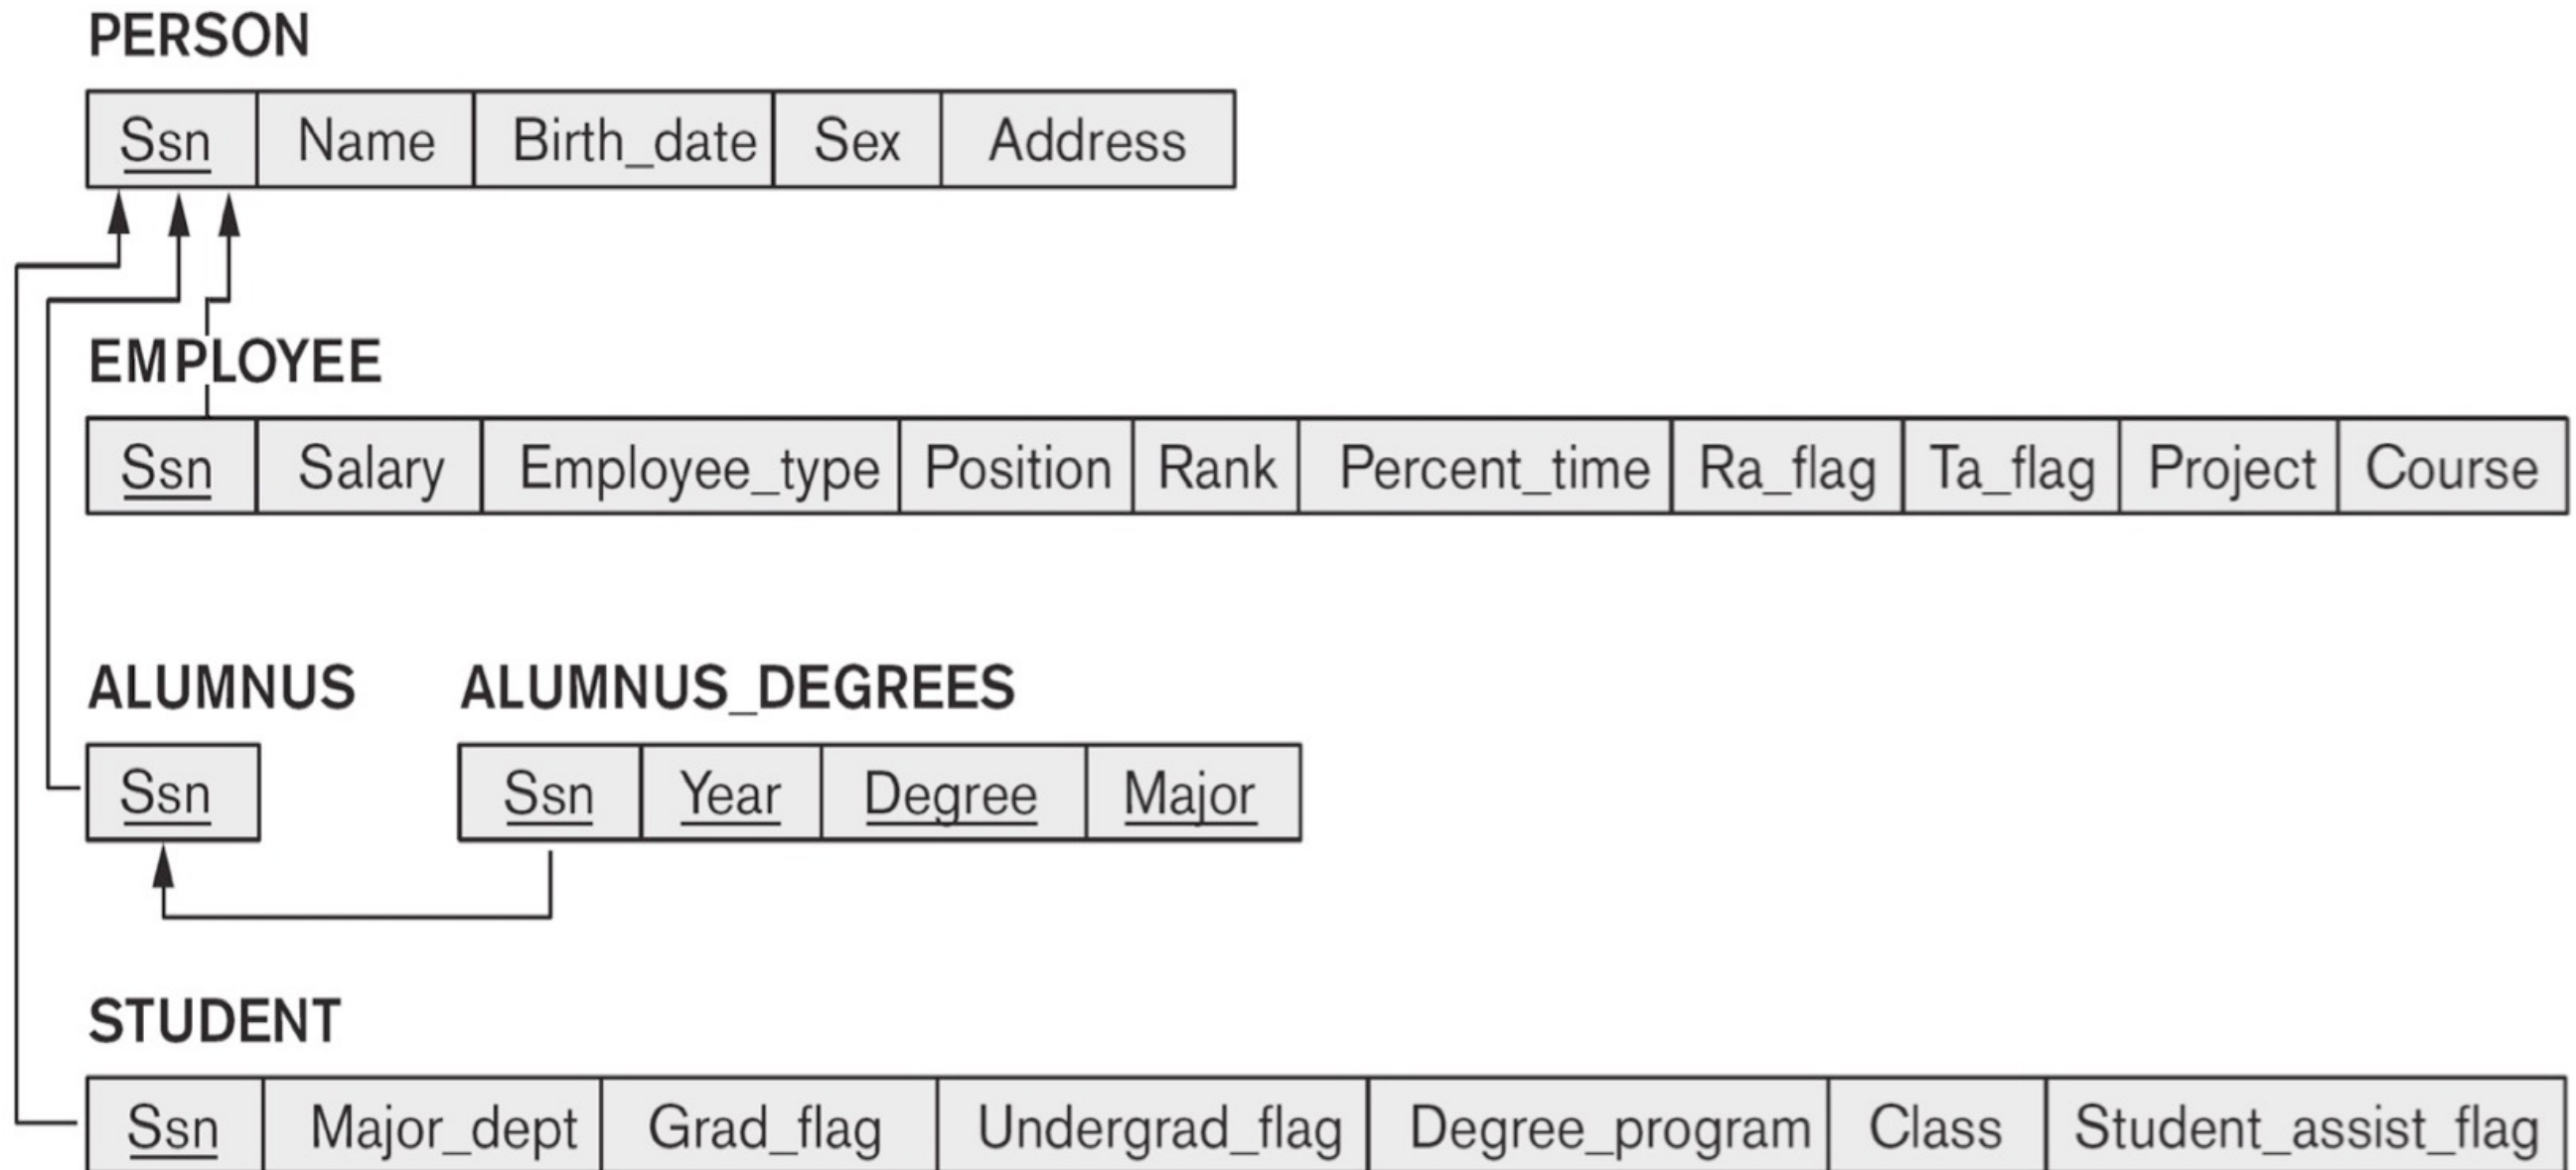
A Specialization Lattice with Multiple Inheritance

32



Mapping the Lattice Using Various Options

33



Step 9: Mapping of Union Types (Categories)

34

- A category (or union type) is a subclass of the union of two or more superclasses that have different keys because they can be of different entity types.
- If the defining superclasses have different keys, the custom is to specify a new key attribute, called a **surrogate key**.

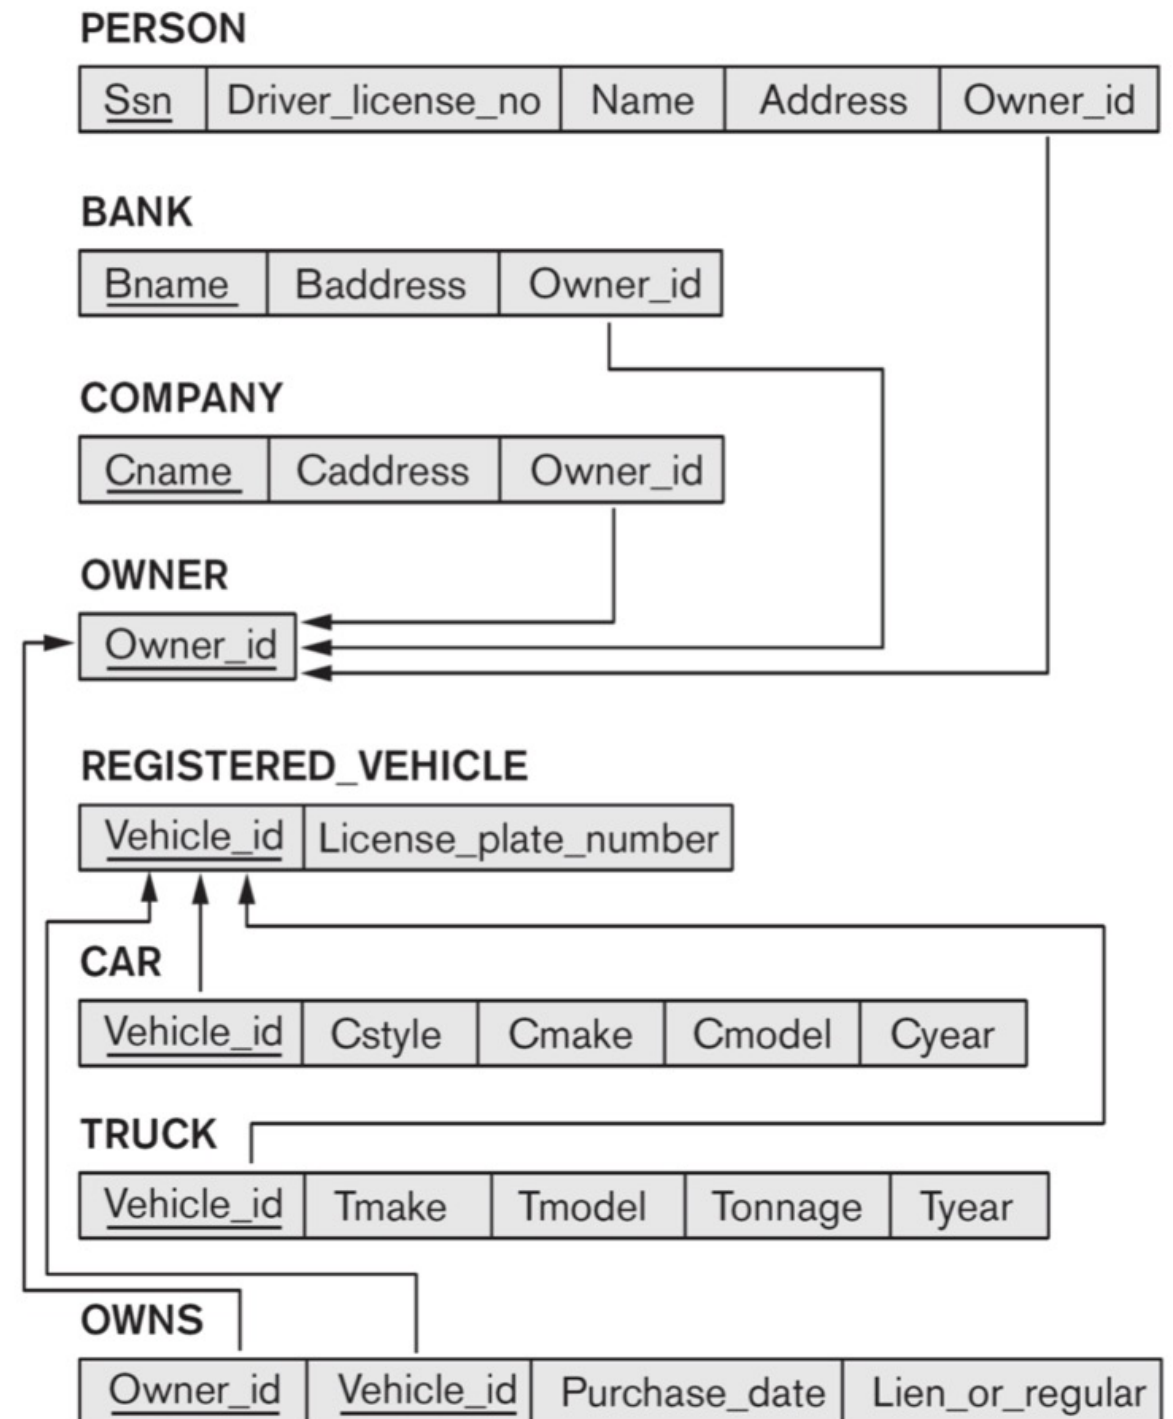
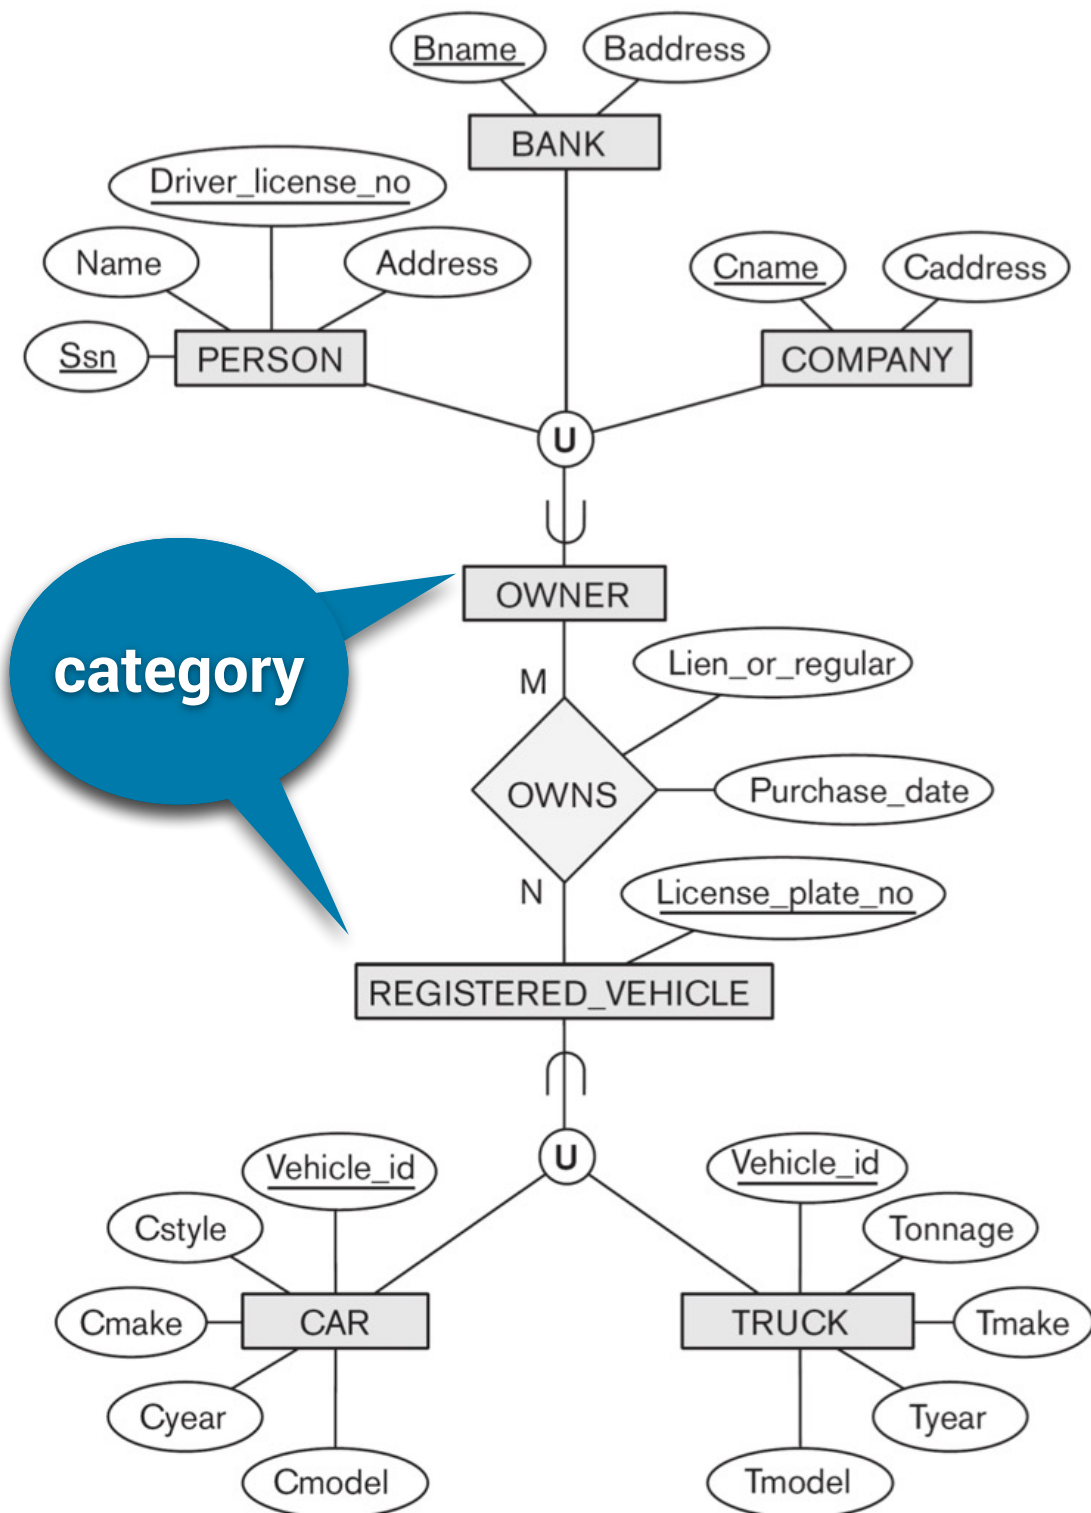
Step 9: Mapping of Union Types (Categories)

35

- For a category C , create a relation C' that includes all the attributes of C .
- Then make the PK of C' the surrogate key of C .
- Then include the PK of C' as FK in the relations corresponding to each of the superclasses of C .
- Then, optionally, add a type (or discriminating) attribute to C' to indicate the entity type (i.e., what is the superclass) of each tuple.
- Note that if a particular entity in a superclass is not a member of C , it would have a NULL surrogate key, i.e., it does not appear as a tuple in C' .

Step 9 Example: Mapping Categories

36



In The Next Handout

37

- We'll learn how to analyze, evaluate and improve a logical model.
- For that we will study the theory of normalization, based on functional dependencies.