



Core SQL

Karen Morton

Whether you're relatively new to writing SQL or you've been writing it for years, learning to write "good" SQL is a process that requires a strong foundation knowledge of core syntax and concepts. This chapter provides a review of the core concepts of the SQL language and its capabilities along with descriptions of the common SQL commands with which you should already be familiar. For those of you who have worked with SQL previously and have a good grasp on the basics, this will be a brief refresher, and it will prepare you for the more detailed treatment of SQL we'll cover in the chapters ahead. If you're a new SQL user, you may want to read *Beginning Oracle SQL* first to make sure you're comfortable with the basics. Either way, Chapter 1 is intended to "level set" you with a whirlwind tour of the five core SQL statements and provide a quick review of the tool we'll be using to execute SQL, SQL*Plus.

The SQL Language

The SQL language was originally developed in the 1970s by IBM and called Structured English QUery Language, or SEQUEL. The language was based on the model for relational database management systems (RDBMS) developed by E.F. Codd in 1969. The acronym was later shortened to SQL due to a trademark dispute. In 1986, ANSI adopted SQL as a standard, and in 1987, ISO did so as well. A piece of not-so-common knowledge is that the official pronunciation of the language was declared to be "ess queue ell" by ANSI. Most people, including me, still use the *sequel* pronunciation just because it flows a bit easier linguistically.

The purpose of SQL is to simply provide an interface to the database, in our case, Oracle. Every SQL statement is a command, or instruction, to the database. It differs from other programming languages like C and Java in that it is intended to process data in sets, not individual rows. The language also doesn't require that you provide instructions on how to navigate to the data—that happens transparently under the covers. But, as you'll see in the chapters ahead, knowing about your data and how and where it is stored can be very important if you want to write efficient SQL in Oracle.

While there are minor differences in how vendors (like Oracle, IBM and Microsoft) implement the core functionality of SQL, the skills you learn in one database will transfer to another. You will be able to use basically the same SQL statements to query, insert, update, and delete data and create, alter, and drop objects regardless of the database vendor.

Although SQL is the standard for use with various RDBMS, it is not particularly relational in practice. I'll expand on this a bit later in the book; I would also recommend that you read C.J. Date's book entitled *SQL and Relational Theory* for a more detailed review. Keep in mind that the SQL language doesn't always follow the relational model precisely—it doesn't implement some elements of the relational model at all while implementing other elements improperly. The fact remains that

since SQL is based on this model you must not only understand SQL but you must understand the relational model as well in order to write SQL as correctly and efficiently as possible.

Interfacing to the Database

Numerous ways have been developed over the years for transmitting SQL to a database and getting results back. The native interface to the Oracle database is the Oracle Call Interface (OCI). The OCI powers the queries that are sent by the Oracle kernel internally to the database. You use the OCI anytime you use one of Oracle's tools like SQL*Plus or SQL Developer. Various other Oracle tools like SQL*Loader, Data Pump, and Real Application Testing (RAT) use OCI as well as language specific interfaces such as Oracle JDBC-OCI, ODP.Net, Oracle Precompilers, Oracle ODBC, and the Oracle C++ Call Interface (OCCI) drivers.

When you use programming languages like COBOL or C, the statements you write are known as Embedded SQL statements and are preprocessed by a SQL preprocessor before the application program is compiled. Listing 1-1 shows an example of a SQL statement that could be used within a C/C++ block.

Listing 1-1. *Embedded SQL Statement Used Within C/C++ Block*

```
{
    int a;
    /* ... */
    EXEC SQL SELECT salary INTO :a
           FROM hr.employees
           WHERE employee_id = 108;
    /* ... */
    printf("The salary is %d\n", a);
    /* ... */
}
```

Other tools, like SQL*Plus and SQL Developer, are interactive tools. You enter and execute commands, and the output is displayed back to you. Interactive tools don't require you to explicitly compile your code before running it. You simply enter the command you wish to execute. Listing 1-2 shows an example of using SQL*Plus to execute a statement.

Listing 1-2. *Using SQL*Plus to Execute a SQL Statement*

```
SQL> select salary
2  from hr.employees
3  where employee_id = 108;

      SALARY
-----
      12000
```

In this book, we'll use SQL*Plus for our example listings just for consistency, but keep in mind that whatever method or tool you use to enter and execute SQL statements, everything ultimately goes

through the OCI. The bottom line is that the tool you use doesn't matter, the native interface is the same for all.

Review of SQL*Plus

SQL*Plus is a command line tool provided with every Oracle installation regardless of platform (Windows, Unix). It is used to enter and execute SQL commands and to display the resulting output in a text-only environment. The tool allows you to enter and edit commands, save and execute commands individually or via script files, and display the output in nicely formatted report form. To start SQL*Plus you simply start `sqlplus` from your host's command prompt.

Connect to a Database

There are multiple ways to connect to a database from SQL*Plus. Before you can connect however, you will likely need to have entries for the databases you will need to connect to entered in the `$ORACLE_HOME/network/admin/tnsnames.ora` file. Two common ways are to supply your connection information when you start SQL*Plus, as shown in Listing 1-3; the other is to use the SQL*Plus connect command after SQL*Plus starts, as shown in Listing 1-4.

Listing 1-3. *Connecting to SQL*Plus from the Windows Command Prompt*

```
E:\pro_oracle_sql>sqlplus hr@ora11r2

SQL*Plus: Release 11.2.0.1.0 - Production on Sun Jun 6 11:22:24 2010

Copyright (c) 1982, 2010, Oracle. All rights reserved.
Enter password:

Connected to:
Oracle Database 11g Enterprise Edition Release 11.2.0.1.0 - Production
With the Partitioning, OLAP and Data Mining and Real Appliation Testing
options

SQL>
```

To start SQL*Plus without being prompted to login to a database, you can start SQL*Plus with the `/nolog` option.

Listing 1-4. *Connecting to SQL*Plus and Logging into the Database from The SQL> Prompt*

```
E:\pro_oracle_sql>sqlplus /nolog

SQL*Plus: Release 11.2.0.1.0 - Production on Sun Jun 6 11:22:24 2010

Copyright (c) 1982, 2010, Oracle. All rights reserved.

SQL> connect hr@ora11r2
Enter password:
Connected.
SQL>
```

Configuring the SQL*Plus environment

SQL*Plus has numerous commands that allow you to customize the working environment and display options. Listing 1-5 shows the SQL*Plus commands available after entering the SQL*Plus help index command at the SQL> prompt.

Listing 1-5. *SQL*Plus Command List*

```
SQL> help index
```

Enter Help [topic] for help.

@	COPY	PAUSE	SHUTDOWN
@@	DEFINE	PRINT	SPOOL
/	DEL	PROMPT	SQLPLUS
ACCEPT	DESCRIBE	QUIT	START
APPEND	DISCONNECT	RECOVER	STARTUP
ARCHIVE LOG	EDIT	REMARK	STORE
ATTRIBUTE	EXECUTE	REPFOOTER	TIMING
BREAK	EXIT	REPHEADER	TTITLE
BTITLE	GET	RESERVED WORDS (SQL)	UNDEFINE
CHANGE	HELP	RESERVED WORDS (PL/SQL)	VARIABLE
CLEAR	HOST	RUN	WHENEVER OSERROR
COLUMN	INPUT	SAVE	WHENEVER SQLERROR
COMPUTE	LIST	SET	XQUERY
CONNECT	PASSWORD	SHOW	

The set command is the primary command used for customizing your environment settings. Listing 1-6 shows the help text for the set command.

Listing 1-6. *SQL*Plus SET Command*

```
SQL> help set
```

```
SET
---
```

Sets a system variable to alter the SQL*Plus environment settings for your current session. For example, to:

- set the display width for data
- customize HTML formatting
- enable or disable printing of column headings
- set the number of lines per page

```
SET system_variable value
```

where `system_variable` and `value` represent one of the following clauses:

APPI[NFO] {OFF ON text}	NEWP[AGE] {1 n NONE}
ARRAY[SIZE] {15 n}	NULL text
AUTO[COMMIT] {OFF ON IMM[EDIATE] n}	NUMF[ORMAT] format
AUTOP[RINT] {OFF ON}	NUM[WIDTH] {10 n}
AUTORECOVERY {OFF ON}	PAGES[IZE] {14 n}
AUTOT[RACE] {OFF ON TRACE[ONLY]}	PAU[SE] {OFF ON text}
[EXP[LAIN]] [STAT[ISTICS]]	RECSEP {WR[APPED] EA[CH] OFF}
BLO[CKTERMINATOR] {. c ON OFF}	RECSEPCHAR {_ c}
CMDS[EP] {; c OFF ON}	SERVEROUT[PUT] {ON OFF}
COLSEP {_ text}	[SIZE {n UNLIMITED}]
CON[CAT] {. c ON OFF}	[FOR[MAT] {WRA[PPED]
COPYC[OMMIT] {o n}	WOR[D_WAPPED]
COPYTYPECHECK {ON OFF}	TRU[NCATED]}}
DEF[INE] {& c ON OFF}	SHIFT[INOUT] {VIS[IBLE]
DESCRIBE [DEPTH {1 n ALL}]	INV[ISIBLE]]}
[LINENUM {OFF ON}] [INDENT {OFF ON}]	SHOW[MODE] {OFF ON}
ECHO {OFF ON}	SQLBL[ANKLINES] {OFF ON}
EDITF[ILE] file_name[.ext]	SQLC[ASE] {MIX[ED]
EMB[EDDED] {OFF ON}	LO[WER] UP[PER]]}
ERRORL[OGGING] {ON OFF}	SQLCO[NTINUE] {> text}
[TABLE [schema.]tablename]	SQLN[UMBER] {ON OFF}
[TRUNCATE] [IDENTIFIER identifier]	SQLPLUSCOMPAT[IBILITY]
	{x.y[z.]}
ESC[APE] {\\ c OFF ON}	SQLPRE[FIX] {# c}
ESCCHAR {@ ? % \$ OFF}	SQLP[ROMPT] {SQL> text}
EXITC[OMMIT] {ON OFF}	SQLT[ERMINATOR] {; c ON OFF}
FEED[BACK] {6 n ON OFF}	SUF[FIX] {SQL text}
FLAGGER {OFF ENTRY INTERMED[IATE] FULL} TAB {ON OFF}	
FLU[SH] {ON OFF}	TERM[OUT] {ON OFF}
HEA[DING] {ON OFF}	TI[ME] {OFF ON}
HEADS[EP] { c ON OFF}	TIMI[NG] {OFF ON}
INSTANCE [instance_path LOCAL]	TRIM[OUT] {ON OFF}
LIN[ESIZE] {80 n}	TRIMS[POOL] {OFF ON}
LOBOF[FSET] {1 n}	UND[ERLINE] {- c ON OFF}
LOGSOURCE [pathname]	VER[IFY] {ON OFF}
LONG {80 n}	WRA[P] {ON OFF}
LONGC[HUNKSIZE] {80 n}	XQUERY {BASEURI text
MARK[UP] HTML [OFF ON]	ORDERING{UNORDERED
[HEAD text] [BODY text] [TABLE text]	ORDERED DEFAULT]
[ENTMAP {ON OFF}]	NODE{BYVALUE BYREFERENCE
[SPOOL {OFF ON}]	DEFAULT}
[PRE[FORMAT] {OFF ON}]	CONTEXT text}
SQL>	

Given the number of commands available, you can easily customize your environment to best suit you. One thing to keep in mind is that the set commands aren't retained by SQL*Plus when you exit/close the tool. Instead of typing in each of the set commands you want to apply each time you use SQL*Plus, you can create a file named `login.sql`. There are actually two files which SQL*Plus reads by default each time you start it. The first is `glogin.sql` and it can be found in the directory `$ORACLE_HOME/sqlplus/admin`. If this file is found, it is read and the statements it contains are executed. This will allow you to store the SQL*Plus commands and SQL statements that customize your experience across SQL*Plus sessions.

After reading `glogin.sql`, SQL*Plus looks for the `login.sql` file. This file must exist in either the directory from which SQL*Plus was started or in a directory included in the path the environment variable `SQLPATH` points to. Any commands in `login.sql` will take precedence over those in `glogin.sql`. Since 10g, Oracle reads both `glogin.sql` and `login.sql` each time you either start SQL*Plus or execute the `connect` command from within SQL*Plus. Prior to 10g, the `login.sql` script was only executed when SQL*Plus started. The contents of a common `login.sql` file are shown in Listing 1-7.

Listing 1-7. *A Common login.sql File*

```
SET LINES 3000
Sets width of display line (default 80 characters)
SET PAGES 1000
Sets number of lines per page (default 14 lines)
SET TIMING ON
Sets display of elapsed time (default OFF)
SET NULL <null>
Sets display of nulls to show <null> (default empty)
SET SQLPROMPT '&_user@&_connect_identifier> '
Sets the prompt to show connected user and instance
```

Note the use of the variables `_user` and `_connect_identifier` in the `SET SQLPROMPT` command. These are two examples of predefined variables. You may use any of the following predefined variables in your `login.sql` file or in any other script file you may create:

- `_connect_identifier`
- `_date`
- `_editor` (This variable specifies the editor which is started when you use the `edit` command.)
- `_o_version`
- `_o_release`
- `_privilege`
- `_sqlplus_release`
- `_user`

Executing Commands

There are two types of commands that can be executed within SQL*Plus: SQL statements and SQL*Plus commands. The SQL*Plus commands shown in Listing 1-5 and 1-6 are specific to SQL*Plus and can be used for customizing the environment and executing commands that are specific to SQL*Plus, like

DESCRIBE and CONNECT. Executing a SQL*Plus command requires only that you type the command at the prompt and hit Enter. The command is automatically executed. On the other hand, in order to execute SQL statements, you must use a special character to indicate you wish to execute the entered command. You may use either a semi-colon (;) or a forward slash (/). A semi-colon may be placed directly at the end of the typed command or on a following blank line. The forward slash must be placed on a blank line in order to be recognized. Listing 1-8 shows how these two execution characters are used.

Listing 1-8. *Execution Character Usage*

```
SQL>select empno, deptno from scott.emp where ename = 'SMITH' ;
      EMPNO      DEPTNO
-----
       7369         20
SQL>select empno, deptno from scott.emp where ename = 'SMITH'
 2 ;
      EMPNO      DEPTNO
-----
       7369         20
SQL>select empno, deptno from scott.emp where ename = 'SMITH'
 2 /
      EMPNO      DEPTNO
-----
       7369         20
SQL>select empno, deptno from scott.emp where ename = 'SMITH'
 2
SQL>/
      EMPNO      DEPTNO
-----
       7369         20
SQL>select empno, deptno from scott.emp where ename = 'SMITH' /
 2
SQL>l
 1* select empno, deptno from scott.emp where ename = 'SMITH' /
SQL>/
select empno, deptno from scott.emp where ename = 'SMITH' /
      *
```

ERROR at line 1:
ORA-00936: missing expression

Notice the fifth example that puts the / at the end of the statement. The cursor moves to a new line instead of executing the command immediately. Then, if you press Enter again, the statement is entered into the SQL*Plus buffer but not executed. In order to view the contents of the SQL*Plus buffer, the list command is used (abbreviated to only l). If you then attempt to execute the statement in the buffer using /, which is how the / command is intended to be used, you get an error. That's because you had typed in the / on the end of the SQL statement line originally. The / is not a valid SQL command and thus causes an error when the statement attempts to execute.

Another way to execute commands is to place them in a file. You can produce these files with the text editor of your choice outside of SQL*Plus or you may invoke an editor directly from SQL*Plus using the EDIT command. The EDIT command will either open a named file or create a file if it doesn't exist. The file must be in the default directory or you must specify the full path. To set the editor to one of your choice, you simply set the predefined `_editor` variable using the following command: `define _editor='<full path>/myeditor.exe'`. Files with the extension of `.sql` will execute without having to include the extension and can be ran using either the `@` or `START` command. Listing 1-9 shows the use of both commands.

Listing 1-9. *Executing .sql Script Files*

```
SQL> @list_depts
      DEPTNO DNAME          LOC
-----
      10 ACCOUNTING      NEW YORK
      20 RESEARCH        DALLAS
      30 SALES            CHICAGO
      40 OPERATIONS      BOSTON

SQL>
SQL> start list_depts
      DEPTNO DNAME          LOC
-----
      10 ACCOUNTING      NEW YORK
      20 RESEARCH        DALLAS
      30 SALES            CHICAGO
      40 OPERATIONS      BOSTON

SQL>
SQL> l
      1* select * from scott.dept
SQL>
```

SQL*Plus has many features and options—way too many to cover here. For what we'll need in this book, this brief overview should suffice. However, the Oracle documentation provides guides for SQL*Plus usage and there are numerous books, including *Beginning Oracle SQL*, that go into more depth if you're interested.

The Five Core SQL Statements

The SQL language contains many different statements. In your professional career you may end up using just a small percentage of what is available to you. But isn't that the case with almost any product you use? I once heard a statistic quoted stating that most people use 20 percent or less of the functionality available in the software products or programming languages they regularly use. I don't know if that's actually true or not, but in my experience, it seems fairly accurate. I have found the same basic SQL statement formats in use within most applications for almost 20 years. Very few people ever use everything SQL has to offer—and often improperly implement those they do use frequently. Obviously, we will not be able to cover all the statements and their options found in the SQL language. This book is intended to provide you deeper insight into the most commonly used statements and help you learn to apply them more effectively.

In this book, we will examine five of the most frequently used SQL statements. These statements are `SELECT`, `INSERT`, `UPDATE`, `DELETE`, and `MERGE`. Although we'll address each of these core statements in some fashion, the focus will be primarily on the `SELECT` statement. Developing a good command of these five statements will provide a strong foundation for your day-to-day work with the SQL language.

The `SELECT` Statement

The `SELECT` statement is used to retrieve data from one or more tables or other database objects. You should already be familiar with the basics of the `SELECT` statement so instead of reviewing the statement from that beginner point of view, I wanted to review how a `SELECT` statement processes logically. You should have already learned the basic clauses that form a common `SELECT` statement, but in order to build the foundation mindset you'll need to write well-formed and efficient SQL consistently, you need to understand how SQL processes.

How a query statement is processed logically may be quite different from its actual physical processing. The Oracle cost-based optimizer (CBO) is responsible for generating the actual execution plan for a query and we will cover what the optimizer does, how it does it, and why in the chapters ahead. For now, note that the optimizer will determine how to access tables and in which order to process them, and how to join multiple tables and apply filters. The logical order of query processing occurs in a very specific order. However, the steps the optimizer chooses for the physical execution plan can end up actually processing the query in a very different order. Listing 1-10 shows a query stub containing the main clauses of a `SELECT` statement with step numbers assigned to each clause in the order it is logically processed.

Listing 1-10. *Logical Query Processing Order*

```

5      SELECT  <column list>
1      FROM    <source object list>
1.1    FROM    <left source object> <join type>
          JOIN <right source object> ON <on predicates>
2      WHERE   <where predicates>
3      GROUP BY <group by expression(s)>
4      HAVING  <having predicates>
6      ORDER BY <order by list>
```

You should notice right away that SQL differs from other programming languages in that the first written statement (the `SELECT`) is not the first line of code that is processed; the `FROM` clause is processed first. Note that I have shown two different `FROM` clauses in this listing. The one marked as 1.1 is provided to show the difference when ANSI syntax is used. It may be helpful to imagine that each step in the processing order creates a temporary dataset. As each step is processed, the dataset is manipulated until a final result is formulated. It is this final result set of data that the query returns to the caller.

In order to walk through each part of the `SELECT` statement in more detail, you'll use the query in Listing 1-11 that returns a result set containing a list of female customers that have placed more than four orders.

Listing 1-11. *Female Customers Who Have Placed More Than Four Orders*

```
SQL> select c.customer_id, count(o.order_id) as orders_ct
  2   from oe.customers c
  3   join oe.orders o
  4   on c.customer_id = o.customer_id
  5   where c.gender = 'F'
  6   group by c.customer_id
  7   having count(o.order_id) > 4
  8   order by orders_ct, c.customer_id
  9   /
```

CUSTOMER_ID	ORDERS_CT
146	5
147	5

The FROM Clause

The FROM clause lists the source objects from which data is selected. This clause can contain tables, views, materialized views, partitions or subpartitions, or may specify a subquery that identifies objects. If multiple source objects are used, this logical processing phase also applies each join type and ON predicates (shown as step 1.1). You'll cover join types in more detail later but note that as joins are processed, they occur in the following order:

1. Cross join, also called a Cartesian product
2. Inner join
3. Outer join

In the example query in Listing 1-11, the FROM clause lists two tables: customers and orders. They are joined on the customer_id column. So, when this information is processed, the initial dataset that will be produced by the FROM clause will include rows where the customer_id matches in both tables. The result set will contain 105 rows at this point. To verify this is true, simply execute only the first four lines of the example query as shown in Listing 1-12.

Listing 1-12. *Partial Query Execution Through the FROM Clause Only*

```
SQL> select c.customer_id cust_id, o.order_id ord_id, c.gender
  2   from oe.customers c
  3   join oe.orders o
  4   on c.customer_id = o.customer_id;
```

CUST_ID	ORD_ID	G	CUST_ID	ORD_ID	G	CUST_ID	ORD_ID	G
147	2450	F	101	2430	M	109	2394	M
147	2425	F	101	2413	M	116	2453	M
147	2385	F	101	2447	M	116	2428	M
147	2366	F	101	2458	M	116	2369	M
147	2396	F	102	2431	M	116	2436	M

148	2451	M	102	2414	M	117	2456	M
148	2426	M	102	2432	M	117	2429	M
148	2386	M	102	2397	M	117	2370	M
148	2367	M	103	2437	F	117	2446	M
148	2406	M	103	2415	F	118	2457	M
149	2452	M	103	2433	F	118	2371	M
149	2427	M	103	2454	F	120	2373	M
149	2387	M	104	2438	F	121	2374	M
149	2368	M	104	2416	F	122	2375	M
149	2434	M	104	2355	F	123	2376	F
150	2388	M	104	2354	F	141	2377	M
151	2389	M	105	2439	F	143	2380	M
152	2390	M	105	2417	F	144	2445	M
153	2391	M	105	2356	F	144	2422	M
154	2392	F	105	2358	F	144	2382	M
155	2393	M	106	2441	M	144	2363	M
156	2395	F	106	2418	M	144	2435	M
157	2398	M	106	2359	M	145	2448	M
158	2399	M	106	2381	M	145	2423	M
159	2400	M	107	2442	F	145	2383	M
160	2401	M	107	2419	F	145	2364	M
161	2402	M	107	2360	F	145	2455	M
162	2403	M	107	2440	F	119	2372	M
163	2404	M	108	2443	M	142	2378	M
164	2405	M	108	2420	M	146	2449	F
165	2407	M	108	2361	M	146	2424	F
166	2408	F	108	2357	M	146	2384	F
167	2409	M	109	2444	M	146	2365	F
169	2411	F	109	2421	M	146	2379	F
170	2412	M	109	2362	M	168	2410	M

105 rows selected.

■ **NOTE** I formatted the result of this output manually to make it fit nicely on the page. The actual output was displayed over 105 separate lines.

The WHERE Clause

The WHERE clause provides a way to conditionally limit the rows emitted to the query's final result set. Each condition, or predicate, is entered as a comparison of two values or expressions. The comparison will match (evaluate to TRUE) or it will not match (evaluate to FALSE). If the comparison is FALSE, then the row will not be included in the final result set.

I need to digress just a bit to cover an important aspect of SQL related to this step. Actually, the possible values of a logical comparison in SQL are TRUE, FALSE, and UNKNOWN. The UNKNOWN value occurs

when a null is involved. Nulls compared to anything or nulls used in expressions evaluate to null, or UNKNOWN. A null represents a missing value and can be confusing due to inconsistencies in how nulls are treated within different elements of the SQL language. We'll address how nulls effect the execution of SQL statements throughout the book, but I didn't want to ignore mentioning the topic at this point. What I stated previously is still basically true, that comparisons will either return TRUE or FALSE. What you'll find is that when a null is involved in a filter comparison, it is treated as if it were FALSE.

In our example, there is a single predicate used to limit the result to only females who have placed orders. If you review the intermediate result after the FROM clause was processed (see Listing 1-12), you'll note that only 31 of the 105 rows were placed by female customers (gender = 'F'). Therefore, after the WHERE clause is applied, the intermediate result set would be reduced from 105 down to 31 rows.

After the WHERE clause is applied, the detailed result set is ready. Note that I use the phrase "detailed result set." What I mean is the rows that satisfy your query requirements are now available. Other clauses may be applied (GROUP BY, HAVING) that will aggregate and further limit the final result set that the caller will receive, but it is important to note that at this point, all the data your query needs to compute the final answer is available.

The WHERE clause is intended to restrict, or reduce, the result set. The less restrictions you include, the more data your final result set will contain. The more data you need to return, the longer the query will take to execute.

The GROUP BY Clause

The GROUP BY clause aggregates the filtered result set available after processing the FROM and WHERE clauses. The selected rows are grouped by the expression(s) listed in this clause to produce a single row of summary information for each group. You may group by any column of any object listed in the FROM clause even if you don't intend to display that column in the list of output columns. Conversely, any non-aggregate column in the select list must be included in the GROUP BY expression.

There are two additional operations that can be included in a GROUP BY clause: ROLLUP and CUBE. The ROLLUP operation is used to produce subtotal values. The CUBE operation is used to produce cross-tabulation values. If you use either of these operations, you'll get more than one row of summary information. Both of these operations will be discussed in detail in Chapter 7: Advanced Grouping.

In the example query, the requested grouping is by customer_id. This means that there will only be one row for each distinct customer_id. Of the 31 rows that represent the females who have placed orders that have made it through the WHERE clause processing, there are 11 distinct customer_id values, as shown in Listing 1-13.

Listing 1-13. *Partial Query Execution Through the GROUP BY Clause*

```
SQL> select c.customer_id, count(o.order_id) as orders_ct
  2  from oe.customers c
  3  join oe.orders o
  4  on c.customer_id = o.customer_id
  5  where gender = 'F'
  6  group by c.customer_id;
```

CUSTOMER_ID	ORDERS_CT
156	1
123	1
166	1

154	1
169	1
105	4
103	4
107	4
104	4
147	5
146	5

11 rows selected.

You’ll notice that the output from the query, while grouped, is not ordered. The display makes it appear as though the rows are ordered by `order_ct`, but this is more coincidence and not guaranteed behavior. This is an important item to remember: the `GROUP BY` clause does not insure ordering of data. If you want the list to display in a specific order, you have to specify an `ORDER BY` clause.

The HAVING Clause

The `HAVING` clause restricts the grouped summary rows to those where the condition(s) in the clause are `TRUE`. Unless you include a `HAVING` clause, all summary rows are returned. The `GROUP BY` and `HAVING` clauses are actually interchangeable positionally; it doesn’t matter which one comes first. However, it seems to make more sense to code them with the `GROUP BY` first since `GROUP BY` is logically processed first. Essentially, the `HAVING` clause is a second `WHERE` clause that is evaluated after the `GROUP BY` occurs and is used to filter on grouped values.

In our example query, the `HAVING` clause, `HAVING COUNT(o.order_id) > 4`, limits the grouped result data of 11 rows down to 2. You can confirm this by reviewing the list of rows returned after the `GROUP BY` is applied, as shown in Listing 1-13. Note that only customers 146 and 147 have placed more than four orders. The two rows that make up the final result set are now ready.

The SELECT List

The `SELECT` list is where the columns included in the final result set from your query are provided. A column can be an actual column from a table, an expression, or even the result of a `SELECT` statement, as shown in Listing 1-14.

Listing 1-14. *Example Query Showing Select List Alternatives*

```
SQL> select c.customer_id, c.cust_first_name||' '||c.cust_last_name,
2  (select e.last_name from hr.employees e where e.employee_id = c.account_mgr_id)
acct_mgr)
3  from oe.customers c;
```

CUSTOMER_ID	CUST_NAME	ACCT_MGR

147	Ishwarya Roberts	Russell
148	Gustav Steenburgen	Russell
...		

931 Buster Edwards	Cambrault
981 Daniel Gueney	Cambrault

319 rows selected.

When another SELECT statement is used to produce the value of a column, the query must return only one row and one column value. These types of subqueries are referred to as scalar subqueries. While this can be very useful syntax, keep in mind that the scalar subquery will be executed once for each row in the result set. There are optimizations available that may eliminate some duplicate executions of the subquery, but the worse case scenario is that each row will require this scalar subquery to be executed. Imagine the possible overhead involved if your result set had thousands, or millions, of rows! We'll review scalar subqueries later in the book and discuss how to use them optimally.

Another option you may need to use in the SELECT list is the DISTINCT clause. The example doesn't use it, but I wanted to mention it briefly. The DISTINCT clause causes duplicate rows to be removed from the data set produced after the other clauses have been processed.

After the select list is processed, you now have the final result set for your query. The only thing that remains to be done, if it is included, is to sort the result set into a desired order.

The ORDER BY Clause

The ORDER BY clause is used to order the final set of rows returned by the statement. In this case, the requested sort order was to be by orders_ct and customer_id. The orders_ct column is the value computed using the COUNT aggregate function in the GROUP BY clause. As shown in Listing 1-13, there were two customers that each placed more than four orders. Since each customer placed five orders, the order_ct is the same, so the second ordering column determines the final display order. As shown in Listing 1-15, the final sorted output of the query is a two row data set ordered by customer_id.

Listing 1-15. Example Query Final Output

```
SQL> select c.customer_id, count(o.order_id) as orders_ct
  2  from oe.customers c
  3  join oe.orders o
  4  on c.customer_id = o.customer_id
  5  where c.gender = 'F'
  6  group by c.customer_id
  7  having count(o.order_id) > 4
  8  order by orders_ct, c.customer_id
  9  /
```

CUSTOMER_ID	ORDERS_CT
146	5
147	5

When ordered output is requested, Oracle must take the final set of data after all other clauses have been processed and sort them as specified. The size of the data that needs to be sorted is important. When I say size, I mean total bytes of data that is in the result set. To estimate the size of the data set, you multiply the number of rows by the number of bytes per row. The bytes per row are determined by summing the average column lengths of each of the columns in the select list.

The example query requests only the `customer_id` and `orders_ct` column values in the select list. Let's use 10 as our estimated bytes per row value. I'll show you in Chapter 6 where to find the optimizer's estimate for this value. So, given that we only have two rows in the result set, the sort size is actually quite small, approximately 20 bytes. Remember that this is only an estimate, but the estimate is an important one.

Small sorts should be accomplished entirely in memory while large sorts may have to use temporary disk space to complete the sort. As you can likely deduce, a sort that occurs in memory will be faster than a sort that must use disk. Therefore, when the optimizer estimates the effect of sorting data, it has to consider how big the sort is in order to adjust how to accomplish getting the query result in the most efficient way. In general, consider sorts as a fairly expensive overhead to your query processing time, particularly if the size of your result set is large.

The INSERT Statement

The INSERT statement is used to add rows to a table, partition, or view. Rows can be inserted in either a single-table or multi-table method. A single-table insert will insert values into one row of one table by either explicitly specifying the values or by retrieving the values using a subquery. The multi-table insert will insert rows into one or more tables and will compute the row values it inserts by retrieving the values using a subquery.

Single-table Inserts

The first example in Listing 1-16 illustrates a single-table insert using the values clause. Each column value is explicitly entered. The column list is optional if you include values for each column defined in the table. However, if you only want to provide values for a subset of the columns, you must specify the column names in the column list. A good practice is to include the column list regardless of whether or not you specify values for all the columns. Doing so acts to self-document the statement and also can help reduce possible errors that might happen in the future should someone add a new column to the table.

Listing 1-16. *Single-Table Insert*

```
SQL> insert into hr.jobs (job_id, job_title, min_salary, max_salary)
  2  values ('IT_PM', 'Project Manager', 5000, 11000) ;
```

1 row created.

```
SQL> insert into scott.bonus (ename, job, sal)
  2  select ename, job, sal * .10
  3  from scott.emp;
```

14 rows created.

The second example illustrates an insert using a subquery. This is a very flexible option for inserting rows. The subquery can be written to return one or more rows. Each row returned will be used to supply column values for the new rows to be inserted. The subquery can be as simple or complex as needed to satisfy your needs. In this example, we use the subquery to compute a 10% bonus for each employee based on their current salary. The bonus table actually has four columns, but we only populate three of them with this insert. The `comm` column isn't populated with a value from the

subquery and we do not include it in the column list. Since we don't include this column, the value for that column will be null. Note that if the `comm` column had a not null constraint, you would have gotten a constraint error and the statement would have failed.

Multi-table Inserts

The multi-table insert example in Listing 1-17 illustrates how rows returned from a single subquery can be used to insert rows into more than one table. We start with three tables: `small_customers`, `medium_customers`, and `large_customers`. We'd like to populate these tables with customer data based on the total amount of orders a customer has placed. The subquery sums the `order_total` column for each customer and then the insert conditionally places a row in the proper table based on whether the customer is considered to be small (less than \$10,000 of total orders), medium (between \$10,000 and \$99,999.99), and large (greater than or equal to \$100,000).

Listing 1-17. Multi-Table Insert

```
SQL> select * from small_customers ;

no rows selected

SQL> select * from medium_customers ;

no rows selected

SQL> select * from large_customers ;

no rows selected

SQL> insert all
  2  when sum_orders < 10000 then
  3  into small_customers
  4  when sum_orders >= 10000 and sum_orders < 100000 then
  5  into medium_customers
  6  else
  7  into large_customers
  8  select customer_id, sum(order_total) sum_orders
  9  from oe.orders
 10  group by customer_id ;

47 rows created.

SQL> select * from small_customers ;
```


CUSTOMER_ID	SUM_ORDERS
120	416
121	4797
152	7616.8
157	7110.3
160	969.2
161	600
162	220
163	510
164	1233
165	2519
166	309
167	48

12 rows selected.

```
SQL> select * from medium_customers ;
```

CUSTOMER_ID	SUM_ORDERS
102	69211.4
103	20591.4
105	61376.5
106	36199.5
116	32307
119	16447.2
123	11006.2
141	38017.8
142	25691.3
143	27132.6
145	71717.9
146	88462.6
151	17620
153	48070.6
154	26632
155	23431.9
156	68501
158	25270.3
159	69286.4
168	45175
169	15760.5
170	66816

22 rows selected.

```
SQL> select * from large_customers ;
CUSTOMER_ID SUM_ORDERS
```

```
-----
101 190395.1
104 146605.5
107 155613.2
108 213399.7
109 265255.6
117 157808.7
118 100991.8
122 103834.4
144 160284.6
147 371278.2
148 185700.5
149 403119.7
150 282694.3
```

13 rows selected.

Note the use of the ALL clause following the INSERT keyword. When ALL is specified, the statement will perform unconditional multi-table inserts. This means that each WHEN clause is evaluated for each row returned by the subquery regardless of the outcome of a previous condition. Therefore, you need to be careful about how you specify each condition. For example, if I had used WHEN sum_orders < 100000 instead of the range I specified, the medium_customers table would have included the rows that were also inserted into small_customers.

You should specify the FIRST option to cause each WHEN to be evaluated in the order it appears in the statement and to skip subsequent WHEN clause evaluations for a given subquery row. The key is to remember which option, ALL or FIRST, best meets your needs and then use the one most suitable.

The UPDATE Statement

The UPDATE statement is used to change the column values of existing rows in a table. The syntax for this statement is composed of three parts: UPDATE, SET, and WHERE. The UPDATE clause specifies the table to update. The SET clause specifies which columns are changed and the modified values. The WHERE clause is used to conditionally filter which rows will be updated. It is optional and if it is omitted, the update operation will be applied to all rows of the specified table.

Listing 1-18 demonstrates several different ways an UPDATE statement can be written. First, I create a duplicate of the employees table called employees2, then I execute several different updates that accomplish basically the same task: the employees in department 90 are updated to have a 10% salary increase and, in the case of Example 5, the commission_pct column is also updated. Following are the different approaches taken:

Example 1: Update a single column value using an expression.

Example 2: Update a single column value using a subquery.

Example 3: Update single column using subquery in WHERE clause to determine which rows to update.

Example 4: Update a table using a SELECT statement to define the table and column values.

Example 5: Update multiple columns using a subquery.

Listing 1-18. *UPDATE Statement Examples*

```
SQL> -- create a duplicate employees table
SQL> create table employees2 as select * from employees ;
Table created.

SQL> -- add a primary key
SQL> alter table employees2
  1 add constraint emp2_emp_id_pk primary key (employee_id) ;
```

Table altered.

```
SQL> -- retrieve list of employees in department 90
SQL> select employee_id, last_name, salary
  2 from employees where department_id = 90 ;
```

EMPLOYEE_ID	LAST_NAME	SALARY
100	King	24000
101	Kochhar	17000
102	De Haan	17000

3 rows selected.

```
SQL> -- Example 1: Update a single column value using an expression
```

```
SQL> update employees2
  2 set salary = salary * 1.10 -- increase salary by 10%
  3 where department_id = 90 ;
```

3 rows updated.

```
SQL> commit ;
```

Commit complete.

```
SQL> select employee_id, last_name, salary
2   from employees2 where department_id = 90 ;
```

EMPLOYEE_ID	LAST_NAME	SALARY	
100	King	26400	-- previous value 24000
101	Kochhar	18700	-- previous value 17000
102	De Haan	18700	-- previous value 17000

3 rows selected.

```
SQL> -- Example 2: Update a single column value using a subquery
```

```
SQL> update employees
2   set salary = (select employees2.salary
3                  from employees2
4                  where employees2.employee_id = employees.employee_id
5                      and employees.salary != employees2.salary)
6   where department_id = 90 ;
```

3 rows updated.

```
SQL> select employee_id, last_name, salary
2   from employees where department_id = 90 ;
```

EMPLOYEE_ID	LAST_NAME	SALARY
100	King	26400
101	Kochhar	18700
102	De Haan	18700

3 rows selected.

```
SQL> rollback ;
```

Rollback complete.

```
SQL> -- Example 3: Update single column using subquery in
SQL> -- WHERE clause to determine which rows to update
```

```
SQL> update employees
2   set salary = salary * 1.10
3   where department_id in (select department_id
4                             from departments
5                             where department_name = 'Executive') ;
```

3 rows updated.

```
SQL> select employee_id, last_name, salary
  2   from employees
  3   where department_id in (select department_id
  4                           from departments
  5                           where department_name = 'Executive') ;
```

EMPLOYEE_ID	LAST_NAME	SALARY
100	King	26400
101	Kochhar	18700
102	De Haan	18700

3 rows selected.

```
SQL> rollback ;
```

Rollback complete.

```
SQL> -- Example 4: Update a table using a SELECT statement
```

```
SQL> -- to define the table and column values
```

```
SQL> update (select e1.salary, e2.salary new_sal
  2           from employees e1, employees2 e2
  3           where e1.employee_id = e2.employee_id
  4           and e1.department_id = 90)
  5 set salary = new_sal;
```

3 rows updated.

```
SQL> select employee_id, last_name, salary, commission_pct
  2   from employees where department_id = 90 ;
```

EMPLOYEE_ID	LAST_NAME	SALARY	COMMISSION_PCT
100	King	26400	
101	Kochhar	18700	
102	De Haan	18700	

3 rows selected.

```
SQL> rollback ;
```

Rollback complete.

```
SQL> -- Example 5: Update multiple columns using a subquery
```

```
SQL> update employees
  2  set (salary, commission_pct) = (select employees2.salary, .10 comm_pct
  3      from employees2
  4      where employees2.employee_id = employees.employee_id
  5      and employees.salary != employees2.salary)
  6  where department_id = 90 ;
```

3 rows updated.

```
SQL> select employee_id, last_name, salary, commission_pct
  2  from employees where department_id = 90 ;
```

EMPLOYEE_ID	LAST_NAME	SALARY	COMMISSION_PCT
100	King	26400	.1
101	Kochhar	18700	.1
102	De Haan	18700	.1

3 rows selected.

```
SQL> rollback ;
```

Rollback complete.

```
SQL>
```

The DELETE Statement

The DELETE statement is used to remove rows from a table. The syntax for this statement is composed of three parts: DELETE, FROM, and WHERE. The DELETE keyword stands alone. Unless you decide to use a hint, which we'll discuss later, there are no other options associated with the DELETE keyword. The FROM clause identifies the table from which rows will be deleted. As the examples in Listing 1-19 demonstrate, the table can be specified directly or via a subquery. The WHERE clause provides any filter conditions to help determine which rows are deleted. If the WHERE clause is omitted, the delete operation will delete all rows in the specified table.

Listing 1-19 demonstrates several different ways a DELETE statement can be written. Note that I am using the employees2 table created in Listing 1-18 for these examples. Here are the different delete methods that you'll see:

Example 1: Delete rows from specified table using a filter condition in the WHERE clause.

Example 2: Delete rows using a subquery in the FROM clause.

Example 3: Delete rows from specified table using a subquery in the WHERE clause.

Listing 1-19. DELETE Statement Examples

```
SQL> select employee_id, department_id, last_name, salary
  2  from employees2
  3  where department_id = 90;
```

EMPLOYEE_ID	DEPARTMENT_ID	LAST_NAME	SALARY
100	90	King	26400
101	90	Kochhar	18700
102	90	De Haan	18700

3 rows selected.

SQL> -- Example 1: Delete rows from specified table using

SQL> -- a filter condition in the WHERE clause

SQL> delete from employees2

2 where department_id = 90;

3 rows deleted.

SQL> select employee_id, department_id, last_name, salary

2 from employees2

3 where department_id = 90;

no rows selected

SQL> rollback;

Rollback complete.

SQL> select employee_id, department_id, last_name, salary

2 from employees2

3 where department_id = 90;

EMPLOYEE_ID	DEPARTMENT_ID	LAST_NAME	SALARY
100	90	King	26400
101	90	Kochhar	18700
102	90	De Haan	18700

3 rows selected.

SQL> -- Example 2: Delete rows using a subquery in the FROM clause

SQL> delete from (select * from employees2 where department_id = 90);

3 rows deleted.

SQL> select employee_id, department_id, last_name, salary

2 from employees2

3 where department_id = 90;

no rows selected

```
SQL> rollback;
```

Rollback complete.

```
SQL> select employee_id, department_id, last_name, salary
2   from employees2
3   where department_id = 90;
```

EMPLOYEE_ID	DEPARTMENT_ID	LAST_NAME	SALARY
100	90	King	26400
101	90	Kochhar	18700
102	90	De Haan	18700

3 rows selected.

```
SQL> -- Example 3: Delete rows from specified table using
```

```
SQL> -- a subquery in the WHERE clause
```

```
SQL> delete from employees2
2   where department_id in (select department_id
3                           from departments
4                           where department_name = 'Executive');
```

3 rows deleted.

```
SQL> select employee_id, department_id, last_name, salary
2   from employees2
3   where department_id = 90;
```

no rows selected

```
SQL> rollback;
```

Rollback complete.

```
SQL>
```

The MERGE Statement

The MERGE statement is a single command that combines the ability to update or insert rows into a table by conditionally deriving the rows to be updated or inserted from one or more sources. It is most frequently used in data warehouses to move large amounts of data but its use is not limited to only data warehouse environments. The big value-add this statement provides is that you have a convenient way to combine multiple operations into one. This allows you to avoid issuing multiple INSERT, UPDATE, and DELETE statements. And, as you'll see later in the book, if you can avoid doing work you really don't have to do, your response times will likely improve.

The syntax for the MERGE statement is:

```

MERGE <hint>
  INTO <table_name>
  USING <table_view_or_query>
  ON (<condition>)
  WHEN MATCHED THEN <update_clause>
  DELETE <where_clause>
  WHEN NOT MATCHED THEN <insert_clause>
  [LOG ERRORS <log_errors_clause> <reject limit <integer | unlimited>];

```

In order to demonstrate the use of the MERGE statement, Listing 1-20 shows how to create a test table and then appropriately insert or update rows into that table based on the MERGE conditions.

Listing 1-20. *MERGE Statement Example*

```

SQL> create table dept60_bonuses
  2  (employee_id      number
  3  ,bonus_amt        number);

```

Table created.

```
SQL> insert into dept60_bonuses values (103, 0);
```

1 row created.

```
SQL> insert into dept60_bonuses values (104, 100);
```

1 row created.

```
SQL> insert into dept60_bonuses values (105, 0);
```

1 row created.

```
SQL> commit;
```

Commit complete.

```

SQL> select employee_id, last_name, salary
  2  from employees
  3  where department_id = 60 ;

```

EMPLOYEE_ID	LAST_NAME	SALARY
103	Hunold	9000
104	Ernst	6000
105	Austin	4800

106 Pataballa	4800
107 Lorentz	4200

5 rows selected.

SQL> select * from dept60_bonuses;

EMPLOYEE_ID	BONUS_AMT
-----	-----
103	0
104	100
105	0

3 rows selected.

```
SQL> merge into dept60_bonuses b
  2  using (
  3    select employee_id, salary, department_id
  4    from employees
  5    where department_id = 60) e
  6  on (b.employee_id = e.employee_id)
  7  when matched then
  8    update set b.bonus_amt = e.salary * 0.2
  9    where b.bonus_amt = 0
 10  delete where (e.salary > 7500)
 11  when not matched then
 12    insert (b.employee_id, b.bonus_amt)
 13    values (e.employee_id, e.salary * 0.1)
 14    where (e.salary < 7500);
```

4 rows merged.

SQL> select * from dept60_bonuses;

EMPLOYEE_ID	BONUS_AMT
-----	-----
104	100
105	960
106	480
107	420

4 rows selected.

SQL> rollback;

Rollback complete.

SQL>

The MERGE accomplished the following:

- Two rows were inserted (employee_ids 106 and 107).
- One row was updated (employee_id 105).
- One row was deleted (employee_id 103).
- One row remained unchanged (employee_id 104).

Without the MERGE statement, you would have had to write at least three different statements to complete the same work.

Summary

As you can tell from the examples shown so far, the SQL language offers many alternatives that can produce the same result set. What you may have also noticed is that each of the five core statements can utilize similar constructs, like subqueries. The key is to learn which constructs are the most efficient under various circumstances. We'll cover how to do that later in the book.

If you had any trouble following the examples in this chapter, make sure to take the time to review either *Beginning Oracle SQL* or the *SQL Reference Guide* in the Oracle documentation. The rest of the book will assume you are comfortable with the basic constructs for each of the five core SQL statements: SELECT, INSERT, UPDATE, DELETE, and MERGE.