

From last time

Explain briefly how starvation may occur in process scheduling. (2 marks)

In round-robin scheduling, new processes are typically placed at the end of the ready-state queue rather than at the beginning. Suggest a good reason for this. (2 marks)

A scheduler uses a time-slice of 4.5msec, and a context switch takes 0.5msec. What percentage of CPU time is spent on executing process instructions: (a) if processes use the whole time-slice? (b) if processes only need 0.5msec CPU-bursts?

In general, how would you improve the percentage of CPU time spent on executing process instructions? (3 marks)

COMP25111: Operating Systems

Lecture 8: Process/Thread Synchronisation

John Gurd

School of Computer Science, University of Manchester

Autumn 2015

Overview & Learning Outcomes

Process Synchronisation

Semaphores

Deadlocks

Dining Philosophers

Message-passing

**Everything in this handout about process synchronisation
also applies to thread synchronisation**

Problem: Too much milk

Time	Person 1	Person 2	Person 3
8:00	sleeping	sleeping	check fridge - no milk
8:15	wakes up	sleeping	leave for lecture
8:30	check - no milk	sleeping	travelling
8:45	leave for store	wakes up	arrive at University
9:00	arrive at store	check - no milk	go to lecture
9:15	bought milk	leave for store	(lecture)
9:30	return home	arrive at store	(lecture)
9:45		bought milk	go to store
10:00		return home	arrive at store
10:15			bought milk
10:45			return home

Data inconsistency

Concurrent access to shared data

```
static int i; // shared by A & B
// process or thread A      // process or thread B
i = 0;                      i = 0;
while (i < 100) {           while (i > -100) {
    i++;                    i--;
}                           }
System.out.println("A");    System.out.println("B");
```

(assume both access & assignment are **atomic** i.e. indivisible)

Questions:

- which will finish first?
- will they ever finish?
- if one finishes, will the other also finish?
- does it help if one gets a head start?

Race Condition

Several processes manipulate shared data concurrently
& outcome depends on precise order of what happens when

Q: what are a CPU's atomic operations?

e.g.: shared variable i in memory, initial value 4

process A: $i++$;

i.e. load Reg from i ; $\text{Reg} = \text{Reg} + 1$; store Reg to i

process B: $i--$;

i.e. load Reg from i ; $\text{Reg} = \text{Reg} - 1$; store Reg to i

Question: what can the final value of i be?

Definitions

Data inconsistency: disagreement about data values

Synchronisation: using appropriate policies and mechanisms to ensure the correct operation of cooperating processes

Critical section (Critical region): section of code in which shared data is used

Mutual exclusion (mutex): at most 1 process can be in its critical section at once

i.e. if more than one process tries to enter a critical section simultaneously, only one can succeed – others must wait

Semaphores

Dijkstra, 1965: An integer variable (e.g. s)
accessed via two atomic operations (with Dutch names!):

P(S) (“try-to-reduce”, down, wait, acquire, probe, procure)

```
while (S <= 0)
    ; /*no action*/
S--;
```

V(S) (“increase”, up, signal, release, vacate)

```
S++;
```

Initialise s appropriately

= number of processes allowed in critical section at once
(usually 1)

Don't loop – yield

In practice, it is silly to busy-wait in $P()$

$P()$ adds process to a queue & gives up CPU

$V()$ takes process from queue & makes it “ready”

Example – 1 semaphore

Two processes sharing $A[100]$; initialise $S=1$;

...	...
$P(S);$	$P(S);$
$r1=A[100];$	$r2=A[100];$ critical
$r1++;$	$r2++;$
$A[100]=r1;$	$A[100]=r2;$ section
$V(S);$	$V(S);$
...	...

Example – 2 semaphores

Initialise: `S1=0; S2=0; I=0;`

<code>... /* A */</code>	<code>... /* B */</code>	<code>... /* C */</code>
<code>P (S1);</code>	<code>P (S2);</code>	<code>for (j=0; j<10; j++)</code>
<code>I=I*2;</code>	<code>I=I+5;</code>	<code>I++;</code>
<code>V (S1);</code>	<code>V (S1);</code>	<code>V (S2);</code>
<code>...</code>	<code>...</code>	<code>...</code>

Question: What is the the sequence of events?

Deadlock: everyone waiting for everyone else



Deadlock

A set of waiting processes,
where each process is waiting for something
that can only be provided by another of the processes


e.g. S1, S2 initialised to 1

// A	// B
P(S1);	P(S2);
P(S2);	P(S1);
...	...
V(S1);	V(S2);
V(S2);	V(S1);

Detection complicated

Once occurred, almost impossible for OS to solve

Prevention e.g. Deadlock (trajectory) Diagram



Benjamin D. Esham / Wikimedia Commons



The Dining Philosophers

5 philosophers, 5 plates of spaghetti, 5 forks, 1 circular table
Alternately think (alone) or eat with 2 forks (from right and left)

```
Semaphore fork[5]; // initialise to 1
```

```
philosopher (int i) {  
    while(true) {  
        System.out.println (i + " hungry");  
        P(fork[i]); P(fork[(i+1)mod 5]);  
        System.out.println (i + " eating");  
        V(fork[i]); V(fork[(i+1)mod 5]);  
        System.out.println (i + " thinking"); // delay?  
    }  
}
```

This may deadlock

Message-passing

Shared data → Semaphore(s) control access

No shared data → send copy to other processes

Wait for message instead of semaphore

– from known source(s), or from any source

Con: slower(?)

Pro: more general/flexible

– RPC

– multi-CPU

– Distributed/Network

Summary of key points

Process Synchronisation

Semaphores

Deadlocks – may occur in a variety of situations, usually fatal

Dining Philosophers – many solutions

Message-passing

Everything in this handout about process synchronisation also applies to thread synchronisation

Next: Threads in Java.

Your Questions

For next time

Explain briefly how a deadlock may occur (2 marks)

shared variables x, y, s ; initial values $S1=S2=0$ $x=y=1$ $s=0$

line	Thread A	Thread B	Explain the purpose of the semaphores in:
1.	do{	do{	of the semaphores in:
2.	V(S1)	P(S1)	– lines 2 & 3
3.	P(S2)	V(S2)	of both threads
4.	$x=x+y$	P(S1)	– line 4 of B
5.	V(S1)	$y=x-y$	& line 5 of A
6.	V(S1)	$s=s+1$	– lines 6 & 7 of A
7.	P(S2)	P(S1)	and 7 & 8 of B
8.	print s, y	V(S2)	
9.	}while($s < 7$)	}while($s < 7$)	(1 mark each)

Will A ever terminate? Justify your answer. (1 mark)

What is output by the print statement in line 8 of A? (3 marks)

Exam Questions

shared variables $x_1, x_2, x_3, x_4, x_5, x_6$; initially $S_1=S_2=S_3=0$

	Thread A	Thread B	Thread C
2.	$x_1 = 1$	$x_2 = 2$	$x_3 = 3$
3.	$V(S_1)$	$V(S_1)$	$P(S_1); P(S_1)$
4.	$P(S_2)$	$P(S_3)$	$V(S_2); V(S_3)$
5.	$x_4 = x_2 + x_3$	$x_5 = x_1 + x_3$	$x_6 = x_1 + x_2$
6.	$V(S_1)$	$V(S_1)$	$P(S_1); P(S_1)$
7.	$P(S_2)$	$P(S_3)$	$V(S_2); V(S_3)$
8.	$x_2 = x_5 + x_6$	$x_3 = x_4 + x_6$	$x_1 = x_4 + x_5$

continued...

Exam Question ctd.

Is it possible that:

- line 5 of A executes before line 2 of B?
- line 5 of C executes before line 2 of A?
- line 5 of B executes before line 5 of C?
- line 5 of A executes before line 5 of B?
- line 5 of A executes after line 5 of C?
- line 8 of A executes before line 5 of B?
- line 8 of A executes before line 5 of C?
- line 8 of C executes before line 5 of B?

What are the final values of x1, x2, x3, x4, x5, x6?

If S3 is replaced by S2 throughout, can you find a pattern of execution which gives different answers to the questions above?

Glossary

Shared data

Concurrent access

Data (in)consistency

Atomic action/operation

Race condition

Synchronisation

Critical section/region

Mutual exclusion (mutex)

Semaphore

P()

V()

Busy-wait

Deadlock

Deadlock trajectory diagram

Dining Philosophers

Message-passing

Reading

OSC/J: 6.1, 6.2, 6.5 (skim 6.3, 6.4, 6.6)

older OSC/J: 7.1, 7.2, 7.5 (skim 7.3, 7.4, 7.6)

MOS: 2.3 (opening paragraphs), 2.3.1, 2.3.2, 2.3.5 (skim 2.3, 2.4, intro of MOS2 ch.3 or MOS3 ch.6)

Dining philosophers: AOS 7.6.3, MOS2 2.4.1)

Traffic deadlock: <http://www.glommer.net/blogs/?p=189>