

Before we start

What is RAM fragmentation?

Why does it happen?

What is the cure?

What is file-store defragmentation?

Why is it needed?

What (& why) is each of these:

- disk partition
- mounting e.g. a storage device
- hard & soft (symbolic) links in a file-store

COMP25111: Operating Systems

Lecture 16: The File Manager

John Gurd

School of Computer Science, University of Manchester

Autumn 2015

Overview & Learning Outcomes

Files & File Systems

Naming Service

Storage Service

 Data Structures

 Allocation

File Manager & Virtual Memory

What is a file?

Collection of related information on secondary storage:
e.g. data, programs (.java, .c, .h, .class, .o, binary, ...)

Structure: none (sequence of bytes); or lines; or ... ?

Attributes: (name?) size, last update, owner, ...
(try `ls -la`)

Operations: create, open, read, write, close, delete, ...

Types: should OS recognise/support?

- in the name: .com .exe etc. (MSDOS)
- “magic number” at the beginning of some files (Unix)

Access:

- Sequential: processed in order, from start to end
- Direct (Random): logical records, processed in any order

File system

File Identifiers:

SFID – System – lifetime of file

UFID – User – lifetime of process

(Unix: “file descriptor”, Windows: “file handle”)

Requirements – system calls:

`open`: file-name → UFID

`read`: UFID & count → data

`write`: UFID & data →

Multiple OS Layers:

- naming service: e.g. `open`

- storage service: e.g. `read` & `write` (vector of bytes)

- disk driver: access disk sectors

File system organisation: Directories

Directory: file-name \rightarrow SFID

(SFID gives access to contents & attributes)


Originally:

1 directory per partition (1-level), or
1 per user (2-level)

Nowadays: Tree (or Forest) of directories

Stored on disk just like files but treated differently

Example Directory Tree-Structure



Directories contain files & directories

A normal file is a leaf in the tree

Decoding a path name

Split path at separators (e.g. / or \)

Absolute: left-most component = root directory

Relative: implicitly starts with current working directory

Each component from left must:

- identify a directory
- contain the next component

Final component = file or directory

Data Structures

Each process has:

- working directory (inherited on creation)
- UFIDs

File attributes (metadata, File Control Block FCB):

- file size, permissions, owner, group, dates, ...
- where to find data on disk

Open file table in memory:

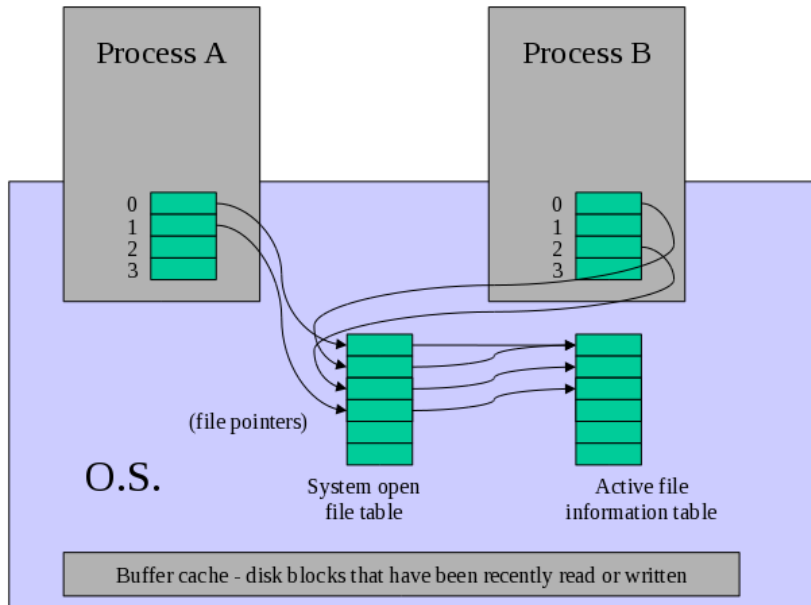
- entry = attributes, number of readers & writers.
- indexed by UFID

`open`: create entry in the file table


(last) `close`: write attributes to disk

Implementations vary e.g. maybe also a table per process

Data Structures



Disks (MOS2 fig 6.11)



Physical structure: platters, tracks, sectors, etc.

Logical structure: blocks

Free Space

e.g. bit-map, or list of block-no? (no = number)

e.g. 100GB partition = $25M * 4kB$ blocks

Q: bitmap size (blocks)?

Q: block-no size (bytes)?

Q: list size (blocks)?

list can use free blocks, bitmap needs extra disk space


Q: search $O(?)$

File structure – contiguous blocks

(e.g. CD/DVD – ROM/WORM)

file = start-block-no & block-count

e.g. file starts at 3, uses 4 blocks:



+ simple, fast

– interleaved user requests → seeks

– fragmentation

File structure – list of blocks (1)

“next” within block

e.g. file starts at 3 (next=0 indicates EOF)

1	2	3	4	5	6	7	8	9	: block numbers
		F 4	I 7			L 9		E 0	: disk blocks

– random/direct access very slow

File structure – list of blocks (2)

“next” in separate monolithic table
(e.g. MSDOS FAT – File Allocation Table)

+ table can also hold free-block info.

e.g. file starts at 3; abc starts at 2 (–1 indicates free)

1	2	3	4	5	6	7	8	9	: block numbers
	a	F	I	c		L	b	E	: disk blocks
–1	8	4	7	0	–1	9	5	0	: FAT

– need to cache table (e.g. 100MB) in memory

File structure – list of blocks (3)

“next” in separate partitioned data-structure

+ one table in RAM (proportional to file size) per open file

e.g. file table in block 1; abc table in block 6

1	2	3	4	5	6	7	8	9
3 4 7 9	a	F	I	c	2 8 5	L	b	E

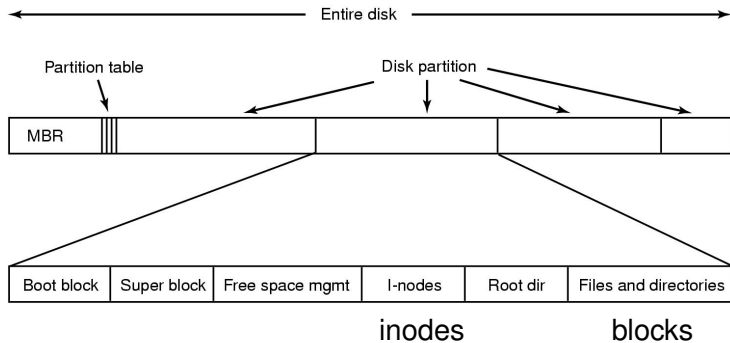
e.g. UNIX i-node/inode = file-attributes + 11 to 15 block-nos

– first 8 to 12 = first blocks of file

– last 3 = block of block-nos, block of blocks of block-nos, ...

inodes in separate disk area (earlier slide)

Disks (again)



inodes in separate area from file/directory blocks

Directory structure

File Name → directory-entry sizes

- max length → fixed-size

(e.g. MSDOS = 8+3, early UNIX = 14)

- unlimited → variable-size (including strlen)

- unlimited → fixed-size (+ “heap” for strings)

File Attributes:

- in directory entry (e.g. FAT)

- pointed at by directory entry (e.g. in inode)

Disk Address:

e.g. FAT: block-number of start of file

e.g. inode: via inode-number

Q: efficiency? – what are the commonest directory operations?

Other Issues

Concurrency: how should multiple accesses be coordinated?

- allow either 1 writer; or many readers (inflexible?)
- applications (e.g. database) define specific protocols

Performance:

- cache
- efficiency dependent on algorithms/types of data
- RAID (Redundant Array of Independent Disks): striping

Access Protection:

- R/W/X permissions (man chmod)
- access control lists (e.g. man acl)

Recovery: backups!

- consistency checking (MSDOS:chkdsk, Linux:fsck)
- partitions
- journalling
- RAID: mirroring

Virtual Memory and Storage

Virtual Memory & File managers both copy info RAM \Leftrightarrow disk

Unified VM/File Manager: memory-mapped files

File Operation	Equivalent VM Operation
open	map file into Virtual Address Space
read	access virtual address: page fault causes disk read
write	access virtual address: (eventual) page rejection causes disk write [n.b. zero fill on write to unallocated page?]
close	unmap pages

Pros: programs & libraries; sharing

Cons: different access patterns?

Summary of key points

Files & File Systems

Naming Service

Storage Service

 Data Structures

 Allocation

File Manager & Virtual Memory

Your Questions

For next time

A file system uses inodes which contain 8 block-numbers. These are for the first 7 blocks of the file and an indirect block, which just contains block-numbers for the remaining blocks in the file.

A block-number occupies 2 bytes. Each block is 4k bytes.

What is the maximum size of a file in bytes?

What is the maximum total size of directories and files in a single disk partition?

Exam Questions

Explain the algorithm used to locate the file referred to by a full path name in a hierarchical file system. (5 marks)

Using a FAT16 file system (i.e. each FAT entry occupies 16 bits) how much space would be available on a 160MB disk for directories and files, for block sizes of 2KB and of 4KB?

Explain your reasoning. (5 marks)

A disk storing a hierarchical file system will hold three forms of data: directories, file contents, and metadata. Illustrate this for a system using a File Allocation Table (FAT). Explain how this information is used and modified by a process making a new copy of an existing file on the disk. (9 marks)

Glossary

file

file attributes

sequential access

direct/random access

magic number

naming service

storage service

directory

hierarchical (tree-structured) directories

SFID, UFID, file descriptor/handle

pathname

relative v. absolute pathnames

metadata, FCB

File Allocation Table, FAT

free space

memory-mapped file

Reading

MOS2: 6.1-6.3

MOS3: 4.1-4.3

OSC/J: 10, 11.1-11.5