


COMP27112

Computer Graphics and Image Processing

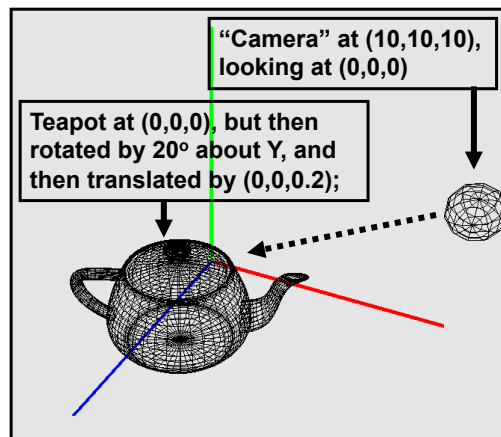
5: Viewing 1: The camera

Toby.Howard@manchester.ac.uk



Introduction

- In part 1 of our study of Viewing, we'll look at
 - Viewing in 2D
 - Clipping
 - Viewing in 3D
 - Viewing in OpenGL

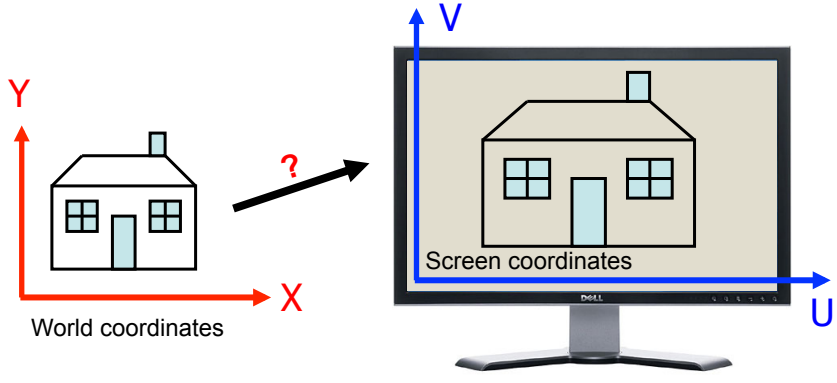


2

MANCHESTER
1824

Viewing in 2D

- We'll start with looking at viewing in 2D



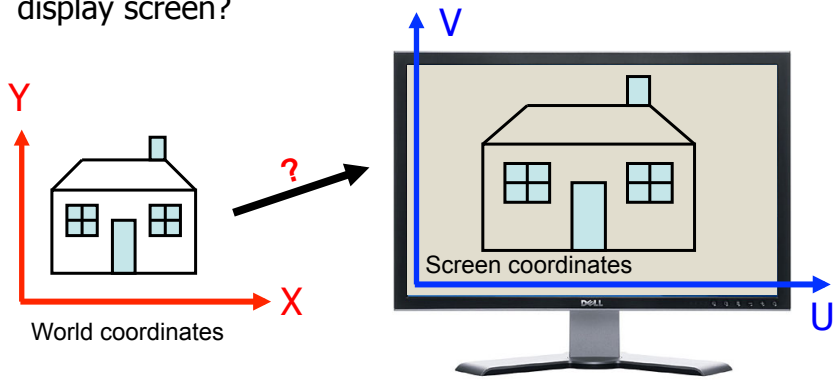
- We need to specify **what** we want to see and **where** we want to see it

3

MANCHESTER
1824

Viewing in 2D

- How do we specify the mapping from our scene to the display screen?



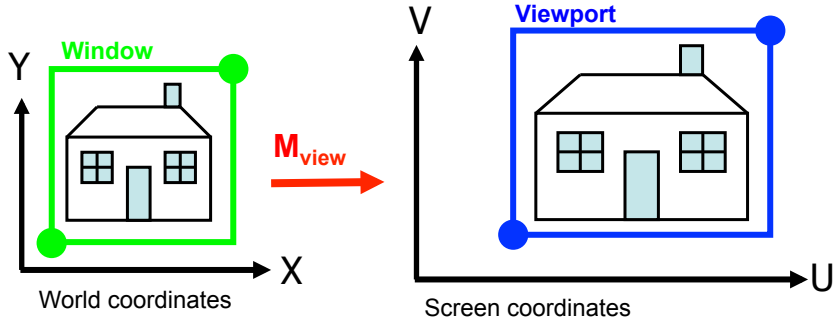
- We use the **analogy** of photographing the scene with a camera
- In this case the camera is confined to the 2D plane

4

MANCHESTER
1824

Viewing in 2D

- We specify a “window” in world coordinates, and a “viewport” in screen coordinates



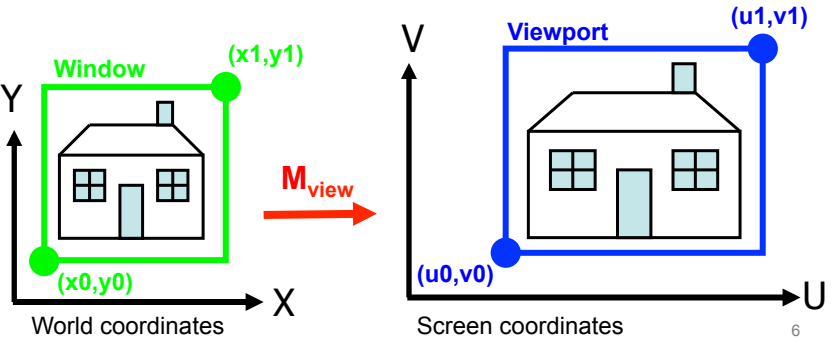
- We find the matrix M_{view} which transforms the window to the viewport.
- M_{view} is called the **viewing transformation**

5

MANCHESTER
1824

Window to viewport mapping

- We can find M_{view} easily, in 3 steps
 - M1** : Translate by $(-x_0, -y_0)$ to place the window at the origin
 - M2** : Scale the window to be the same shape as the viewport
 - M3** : Shift to the viewport position



6

MANCHESTER
1824

Window to viewport mapping

- **M1** : Translate $(-x_0, -y_0)$
- **M2** : Scale $(u_1-u_0/x_1-x_0, v_1-v_0/y_1-y_0)$
- **M3** : Translate (u_0, v_0)
- **M_{view}** = **M3** * **M2** * **M1**
- **P_{screen}** = **M_{view}** * **P_{world}**

7

MANCHESTER
1824

Window to viewport mapping

- **M_{view}** = **M3** * **M2** * **M1**

$$\begin{bmatrix} 1 & 0 & u_0 \\ 0 & 1 & v_0 \\ 0 & 0 & 1 \end{bmatrix} \star \begin{bmatrix} \frac{u_1-u_0}{x_1-x_0} & 0 & 0 \\ 0 & \frac{v_1-v_0}{y_1-y_0} & 0 \\ 0 & 0 & 1 \end{bmatrix} \star \begin{bmatrix} 1 & 0 & -x_0 \\ 0 & 1 & -y_0 \\ 0 & 0 & 1 \end{bmatrix}$$

M₃ **M₂** **M₁**

- So we can now map a world space point **P_{world}** into screen coordinates **P_{screen}** as follows. This is our "view".

$$\mathbf{P}_{\text{screen}} = \begin{bmatrix} \frac{u_1-u_0}{x_1-x_0} & 0 & -x_0 * \frac{u_1-u_0}{x_1-x_0} + u_0 \\ 0 & \frac{v_1-v_0}{y_1-y_0} & -y_0 * \frac{v_1-v_0}{y_1-y_0} + v_0 \\ 0 & 0 & 1 \end{bmatrix} \star \mathbf{P}_{\text{world}}$$


M_{view}

Note: We rarely multiply matrices by hand like this. The graphics system will multiply all matrices together for us.

8


Clipping

- Normally we will want to **CLIP** against the viewport
 - ...to remove those parts of primitives whose coordinates are outside the window
 - There are standard algorithms for clipping lines and polygons

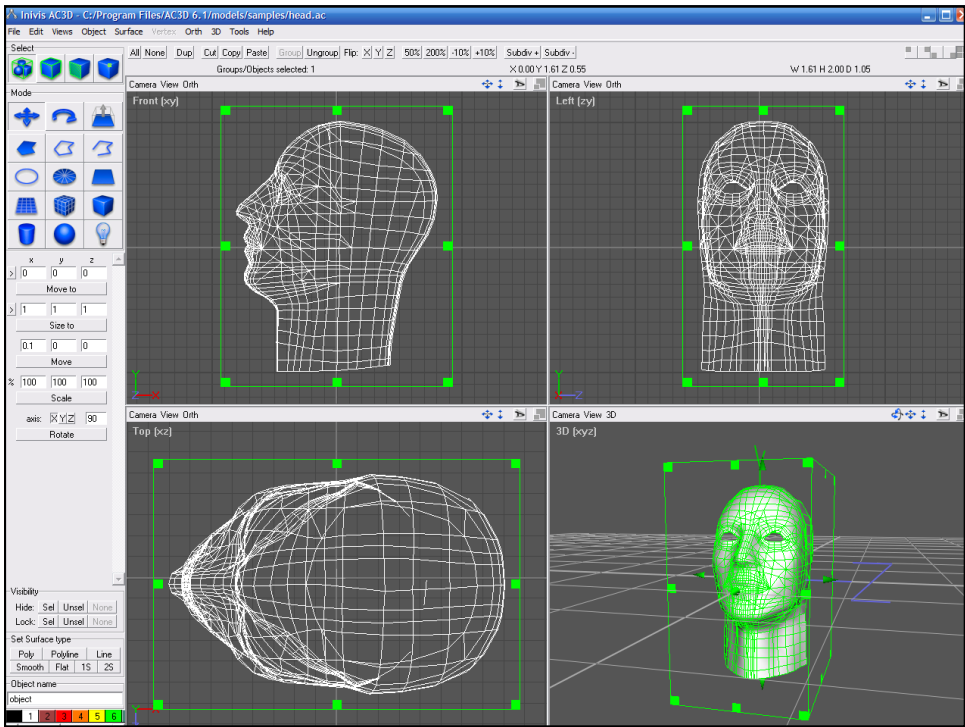


Multiple windows and viewports

- Sometimes it's useful to use multiple windows and viewports, to help arrange items on the screen



- Real example: the 4 different views in AC3D



MANCHESTER
1824

Viewing in 2D: summary


- We use the analogy of photographing our scene with a 2D “camera” which can slide in the XY plane
- We compute a **viewing transformation** M_{view}
- $P_{screen} = M_{view} * P_{world}$

The diagram illustrates the viewing transformation. On the left, a 3D scene with a house and a sun is shown in 'World coordinates' (X, Y). A green box labeled 'Window' represents the viewing frustum. An arrow labeled M_{view} points to the right, where the 2D projection of the scene is shown in 'Screen coordinates' (U, V). The projected scene is enclosed in a blue box labeled 'Viewport'.

12

Viewing in 3D

- In 2D graphics, we “view” our world by mapping from 2D world coordinates to 2D screen coordinates: easy and obvious
- In 3D graphics, in order to “view” our 3D world, we have to somehow **reduce** our 3D information to 2D information, so that it can be displayed on the 2D display: **not** so easy and not obvious
- Here we see a 2D view of an object defined in 3D. It’s been projected from 3D to 2D.
- To specify how this view is created, we again use the analogy of “taking a picture using a camera”, but this time our “camera” is like a real-world camera: It has a position and orientation in 3D space, and a particular type of lens.



13

The camera analogy

- The process of transforming a synthetic 3D model into a 2D view is analogous to using a camera in the real world to take 2D pictures of a 3D scene



14

The camera analogy



Step 1: Arrange the scene into the desired composition

15

The camera analogy



Step 2: Position and point the camera at the scene

16

The camera analogy



Step 3: Choose a camera lens (wide-angle? zoom?)

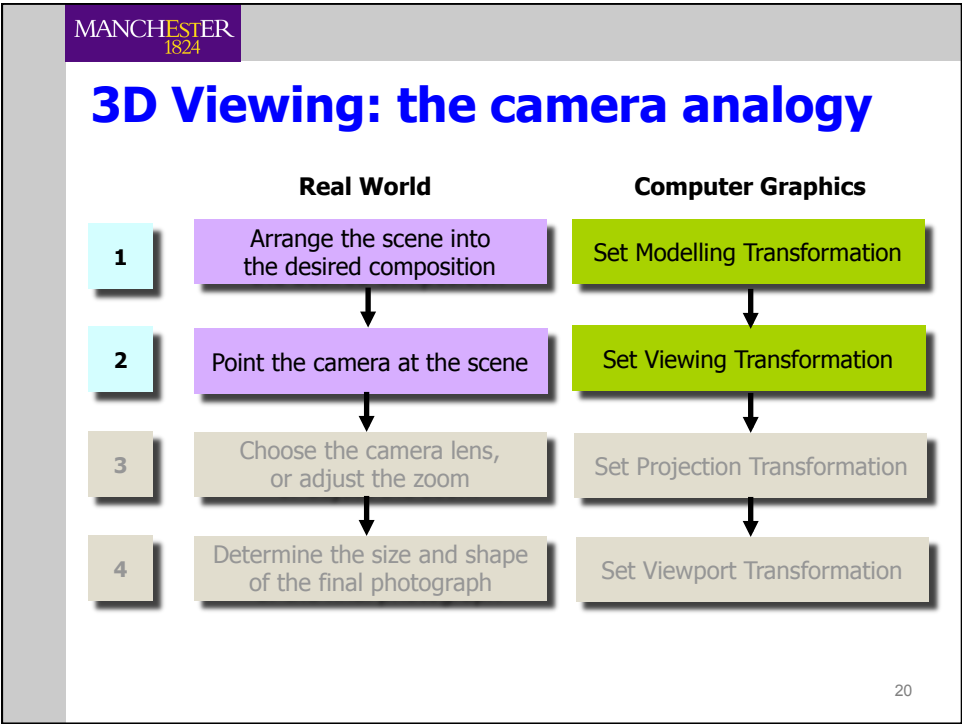
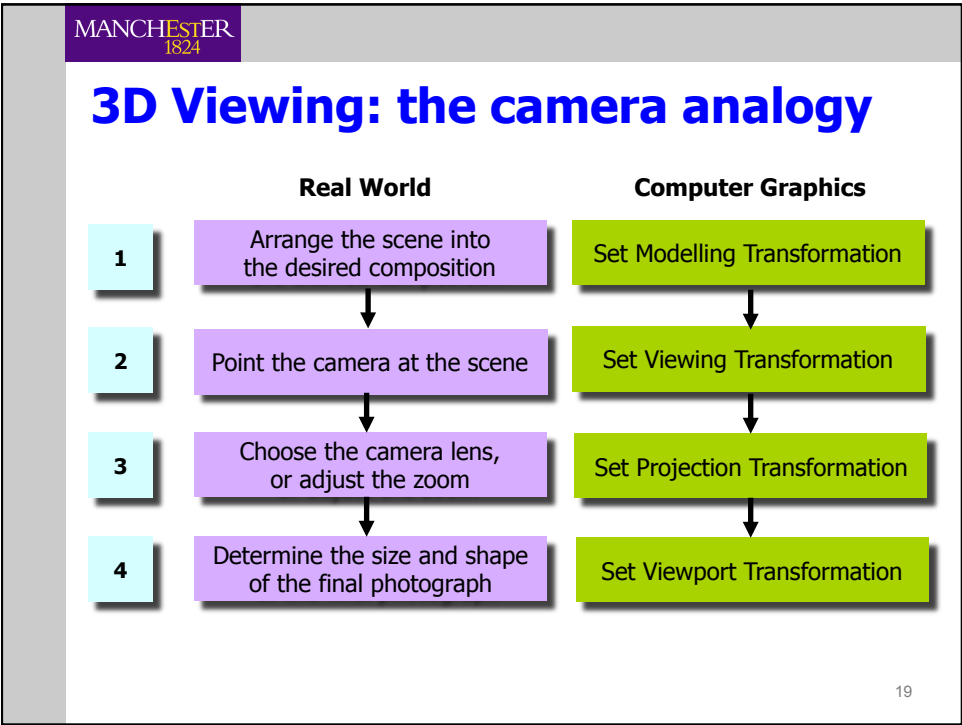
17

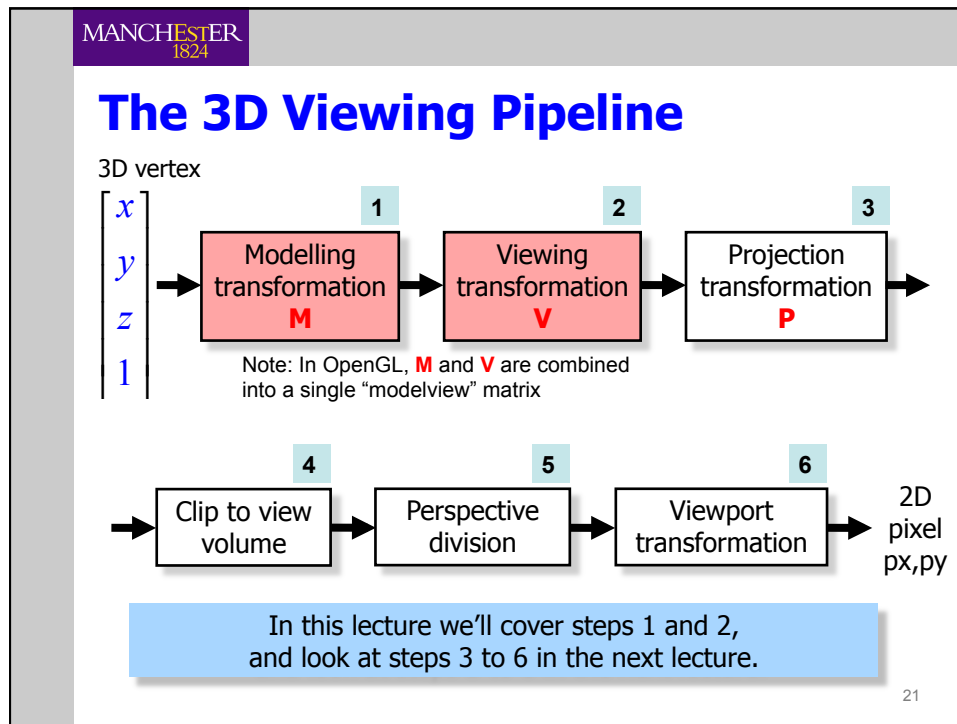
The camera analogy



Step 4: Decide the size of the final photograph

18






MANCHESTER
1824

The duality of Modelling and Viewing

- Question: to obtain these two images of the model, was the camera moved **or** was the model moved?
- Answer: it is impossible to tell, just by looking at the images. It could have been either method.
- Example: moving the model by (x,y,z) is equivalent to moving the camera by $(-x,-y,-z)$. This also applies to rotations.



22

MANCHESTER
1824

The duality of Modelling and Viewing

- Here, in 2D, we have an object **O** and a camera **C**, both sitting on the X axis
- If we keep the object fixed and move the camera by **+2**, **O** and **C** are now **4** units apart.
- If we keep the camera fixed and move the object by **-2**, **O** and **C** are now **4** units apart.
- Whether we move **O** or **C**, their **relative positions** will be the same, so the view from the camera will be the same

23

MANCHESTER
1824

The duality of Modelling and Viewing

- Now for the **KEY IDEA**, which we'll present in 2D
- Imagine we have an object **O** sitting at the origin
- We want to view **O** with a camera **C** located at position **X=3**
- But we don't actually have a camera!
- But we can **SIMULATE** the effect, by instead moving the object by **-3 in X**

24


Achieving viewing by modelling

- We've just seen that we can create the same view from a camera at a certain location and orientation, by **instead** transforming the object
- This is exactly what we do in computer graphics
- However, the idea of having a camera is very natural to us, so we **pretend** we really have a camera...
- ...and we express the view we want in terms of "camera location and orientation", but to implement this we actually compute a suitable viewing transformation which we **apply to the object**

25

Specifying the "camera"

- We specify where the "camera" is located in 3D space. This is the **eye point, E**
- We specify a **centre of interest C**, a 3D point at which the "camera" is looking
- We specify the **up** direction of the "camera", using a "view up vector" **V**
- We can then use **E, C**, and **V** to derive a transformation which, when applied to the **model**, would give the same view as if we really had this "camera".



MANCHESTER
1824

The viewing transformation

- First we use **E**, **C** and **V** to derive a coordinate system for the "camera":
 - $F = C - E$, then normalise F
 - Normalise V (up vector)
 - $\hat{S} = \hat{F} \times \hat{V}$ (cross product)
 - $\hat{U} = \hat{S} \times \hat{F}$ (cross product)
- We now have a coordinate system for the "camera": with axes $\hat{S} \hat{U} \hat{F}$

MANCHESTER
1824

The viewing transformation

- Next we derive a transformation **M** that maps the "camera" coordinate system into world coordinates, in two steps:
 - We translate the origin of the camera system to the origin of the world system
 - We rotate the camera axes to be coincident with the world axes, with \hat{F} aligned with **-Z**

$$M = \begin{bmatrix} \hat{S}_x & \hat{S}_y & \hat{S}_z & 0 \\ \hat{U}_x & \hat{U}_y & \hat{U}_z & 0 \\ -\hat{F}_x & -\hat{F}_y & -\hat{F}_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 0 & -E_x \\ 0 & 1 & 0 & -E_y \\ 0 & 0 & 1 & -E_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

rotation translation
(we omit the derivation)

- M** is the **viewing transformation**

MANCHESTER
1824

The viewing transformation: summary

- The viewing transformation M maps the $\hat{S}\hat{U}\hat{F}$ "camera" axes system to the world coordinates axes system:

world coordinate system

camera coordinate system

Translation and rotation by M

KEY IDEA: If we apply M to **objects**, we will get the **same view** as if we had a real camera

29

MANCHESTER
1824

The viewing transformation in OpenGL

- This OpenGL function takes (E, C, V) and computes the viewing transformation we have just seen

```
void gluLookAt (GLdouble eyex,
                 GLdouble eyez,
                 GLdouble centx,
                 GLdouble centy,
                 GLdouble centz,
                 GLdouble upx,
                 GLdouble upy,
                 GLdouble upz);
```

OpenGL

- As we have seen with other OpenGL transformation functions, `gluLookAt()` creates a temporary matrix T , and then multiplies the modelview matrix by T : $M_{\text{modelview}} = M_{\text{modelview}} * T$

30

The viewing transformation in OpenGL

- The viewing transformation specifies the location and orientation of the “camera” (by in fact transforming the model)
- We incorporate this transformation into the modelview matrix as follows:

```
glMatrixMode(GL_MODELVIEW);  
glLoadIdentity(); // M= identity matrix (I)  
gluLookAt(...stuff...) // M is now I * VIEW
```

- Because we want the viewing transformation to take place **AFTER** any true modelling transformations, we need to “pre-load” the modelview matrix with the viewing transformation...
- ...And then all subsequent modelling transformations will get multiplied into the modelview matrix

31

Modelling and viewing together

```
// First set the viewing transformation  
glMatrixMode(GL_MODELVIEW);  
glLoadIdentity(); // M= identity matrix (I)  
gluLookAt(...stuff...) // M is now I * VIEW  
  
// Now draw a transformed teapot  
glTranslatef(tx, ty, tz);  
// OpenGL computes temp translation matrix T,  
// then sets M= M x T, so now M is (VIEW x T)  
glRotatef(theta, 0.0, 1.0, 0.0);  
// OpenGL computes temp rotation matrix R,  
// then sets M= M x R, so M is now (VIEW x T x R)  
glutWireTeapot(1.0);
```

- So all points **P** will be transformed first by **R**, then **T**, then **VIEW**

32

Example: Setting a view in OpenGL

- Here's a real fragment showing the use of a view transformation and a modelling transformation together
- Note that we also need to set the **projection**, but we'll cover that in the next lecture, so ignore it for now.

```
glMatrixMode(GL_MODELVIEW); // select modelview matrix

glLoadIdentity(); // initialise it


// set the projection (see next lecture)
gluPerspective(...stuff...);
// set the view transformation
gluLookAt(10,10,10, 0,0,0, 0,1,0);

// move/rotate the model however we want
glTranslatef(0.0, 0.0, 0.2);
glRotatef(20.0, 0.0, 1.0, 0.0);

glutWireTeapot(3.0); // draw it
```

See the next slide for a visualisation of this.

33



34

Demonstrations

- Nate Robbins' demonstrations will help you to visualise viewing.
- `/opt/info/courses/OpenGL/tutor` (Linux)

